

AXI Memory Mapped to PCI Express (PCIe) Gen2 v2.9

LogiCORE IP Product Guide

Vivado Design Suite

PG055 November 24, 2023

AMD Adaptive Computing is creating an environment where employees, customers, and partners feel welcome and included. To that end, we're removing non-inclusive language from our products and related collateral. We've launched an internal initiative to remove language that could exclude people or reinforce historical biases, including terms embedded in our software and IPs. You may still find examples of non-inclusive language in our older products as we work to make these changes and align with evolving industry standards. Follow this [link](#) for more information.



Table of Contents

IP Facts

Chapter 1: Overview

Feature Summary	6
Unsupported Features	7
Licensing and Ordering	7

Chapter 2: Product Specification

Standards	9
Performance and Resource Utilization	9
Port Descriptions	10
Bridge Parameters	13
Parameter Dependencies	20
Memory Map	24

Chapter 3: Designing with the Core

General Design Guidelines	40
Clocking	40
Resets	41
Shared Logic	42
Clocking Interface	48
AXI Transactions for PCIe	49
Transaction Ordering for PCIe	50
BAR and Address Translation	51
Interrupts	56
Malformed TLP	58
Abnormal Conditions	58
Root Port	62

Chapter 4: Design Flow Steps

Customizing and Generating the Core	65
Constraining the Core	79
Simulation	81

Chapter 5: Example Design

Overview	84
Simulation Design Overview	84
Implementation Design Overview	86
Example Design Elements	87
Example Design Output Structure	87

Chapter 6: Test Bench

Root Port Model Test Bench for Endpoint	89
Endpoint Model Test Bench for Root Port	91

Appendix A: Upgrading

Migrating to the Vivado Design Suite	93
Upgrading in the Vivado Design Suite	93

Appendix B: Debugging

Finding Help with AMD Adaptive Computing Solutions	95
Debug Tools	96
Simulation Debug	98
Hardware Debug	100
Interface Debug	110

Appendix C: Additional Resources and Legal Notices

Finding Additional Documentation	111
Support Resources	112
References	112
Revision History	113
Please Read: Important Legal Notices	116

Introduction

The AMD AXI Memory Mapped to PCI Express[®] core is an interface between AXI4 and PCI Express.

Features

- Zynq[™] 7000 SoC, Virtex[™] 7, Kintex[™] 7, and Artix[™] 7 FPGA Integrated Blocks for PCI Express⁽³⁾
- Maximum Payload Size (MPS) up to 256 bytes
- Multiple Vector Messaged Signaled Interrupts (MSIs)
- Legacy interrupt support
- Memory-mapped AXI4 access to PCIe[®] space
- PCIe access to memory-mapped AXI4 space
- Tracks and manages Transaction Layer Packets (TLPs) completion processing
- Detects and indicates error conditions with interrupts
- Optimal AXI4 pipeline support for enhanced performance
- Compliant with Advanced RISC Machine (Arm[®]) Advanced Microcontroller Bus Architecture 4 (AMBA[®]) AXI4 specification
- Supports up to three PCIe 32-bit or 64-bit PCIe Base Address Registers (BARs) as Endpoint

Supports a single PCIe 32-bit or 64-bit BAR as Root Port

LogiCORE [™] IP Facts Table	
Core Specifics	
Supported Device Family	AMD Zynq [™] 7000 SoC, AMD Virtex [™] 7, AMD Kintex [™] 7, and AMD Artix [™] 7 ⁽³⁾
Supported User Interfaces	AXI4
Resources	Performance and Resource Utilization web page
Provided with Core	
Design Files	VHDL and Verilog
Example Design	Verilog
Test Bench	Verilog
Constraints File	XDC
Simulation Model	Not Provided
Supported S/W Driver ⁽²⁾	Standalone and Linux
Tested Design Flows ⁽¹⁾	
Design Entry	Vivado [™] Design Suite
Simulation	For supported simulators, see the Vivado Design Suite User Guide: Release Notes, Installation, and Licensing .
Synthesis	Vivado synthesis
Support	
Release Notes and Known Issues	Master Answer Record: 54646
All Vivado IP Change Logs	Master Vivado IP Change Logs: 72775
Support web page	

Notes:

1. For the supported versions of third-party tools, see the [Vivado Design Suite User Guide: Release Notes, Installation, and Licensing](#).
2. Standalone driver details can be found in the Vitis directory (<install_directory>/Vitis/<release>/data/embeddedsw/doc/xilinx_drivers.htm). Linux OS and driver support information is available from the [Wiki page](#).
3. Except for XC7VX485T, XC7V585T, and XC7V2000T, Virtex 7 devices are not supported. XA7A12T, XC7A12T, XA7A15T, XC7A15T, XA7A25T, and XC7A25T devices are not supported due to limited available logic cells within them.

Overview

The AXI Memory Mapped to PCI Express core is designed for the Vivado™ IP integrator in the Vivado Design Suite. The AXI Memory Mapped to PCI Express core provides an interface between an AXI4 customer user interface and PCI Express using the AMD Integrated Block for PCI Express. The AXI Memory Mapped to PCI Express core provides the translation level between the AXI4 embedded system to the PCI Express system. The AXI Memory Mapped to PCI Express core translates the AXI4 memory read or writes to PCIe Transaction Layer Packets (TLP) packets and translates PCIe memory read and write request TLP packets to AXI4 interface commands.

The architecture of the AXI Memory Mapped to PCI Express is shown in [Figure 1-1](#).

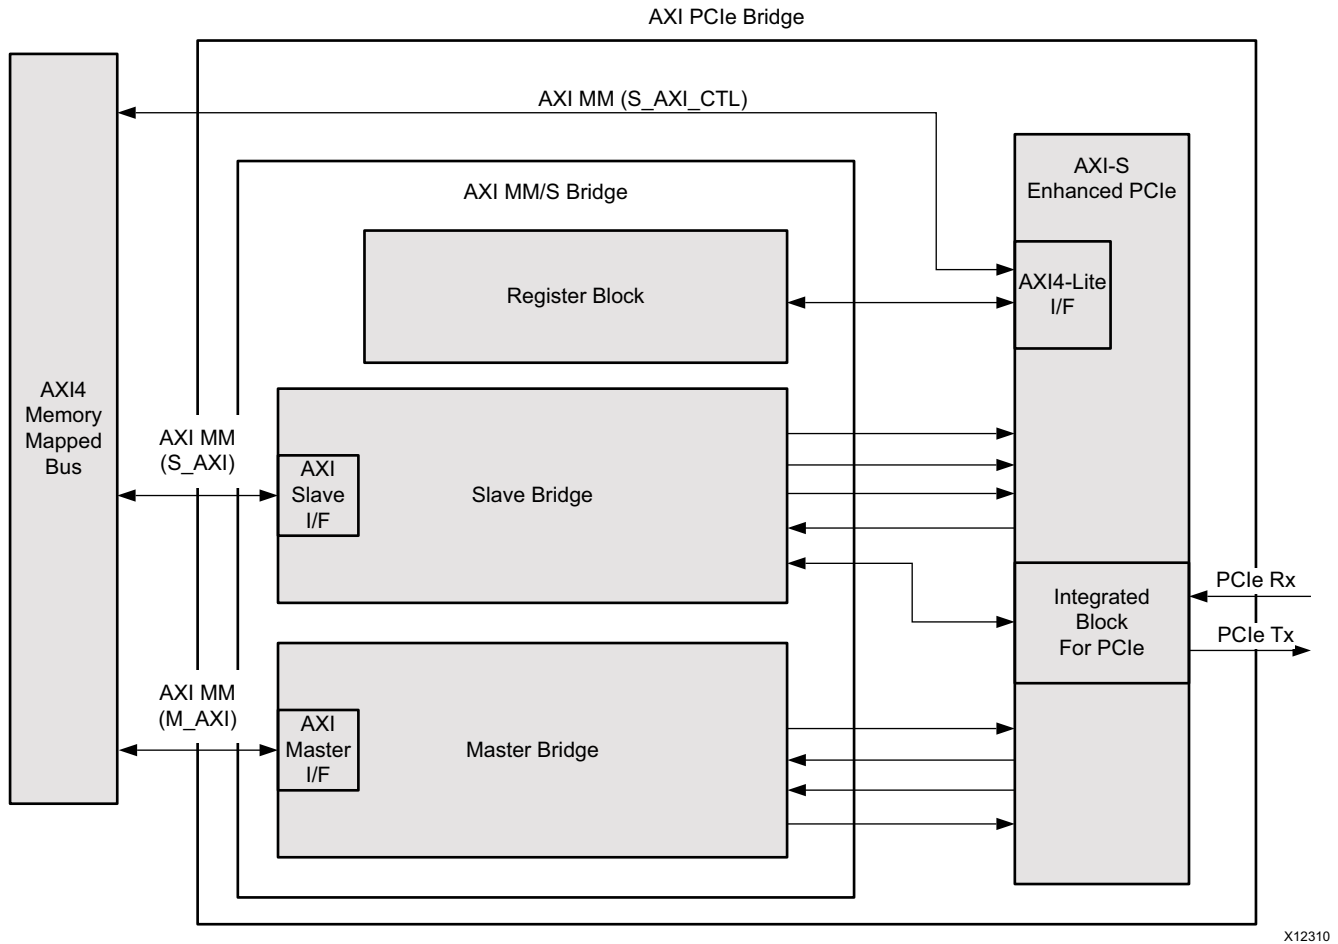


Figure 1-1: High-Level AXI Memory Mapped to PCI Express Architecture

Feature Summary

The AXI Memory Mapped to PCI Express core is an interface between the AXI4 and PCI Express. It contains the memory mapped AXI4 to AXI4-Stream Bridge and the AXI4-Stream Enhanced Interface Block for PCIe. The memory-mapped AXI4 to AXI4-Stream Bridge contains a register block and two functional half bridges, referred to as the Slave Bridge and Master Bridge. The slave bridge connects to the AXI4 Interconnect as a slave device to handle any issued AXI4 master read or write requests. The master bridge connects to the AXI4 Interconnect as a master to process the PCIe generated read or write TLPs. The core uses a set of interrupts to detect and flag error conditions.

The AXI Memory Mapped to PCI Express core supports both Root Port and Endpoint configurations.

- When configured as an Endpoint, the AXI Memory Mapped to PCI Express core supports up to three 32-bit or 64-bit PCIe Base Address Registers (BARs).

- When configured as a Root Port, the core supports a single 32-bit or 64-bit PCIe BAR.

The AXI Memory Mapped to PCI Express core is compliant with the PCI Express Base Specification v2.0 [\[Ref 8\]](#) and with the AMBA[®] AXI Protocol Specification [\[Ref 7\]](#). PCIe compliance is run with the base PCI Express IP for each silicon family. You can check the list of qualified devices at [Integrators List](#).

Unsupported Features

The following features are not supported in the AXI Memory Mapped to PCI Express core.

- Tandem PROM and Tandem PCIe
- Advanced Error Reporting (AER)
- MSI-X

The PCIe transactions that can be generated are those that are compatible with the AXI4 specification. lists the supported PCIe transaction types.

Licensing and Ordering

This AMD LogiCORE™ IP module is provided at no additional cost with the AMD Vivado™ Design Suite under the terms of the [AMD End User License](#).

For more information, visit the [AXI Memory Mapped to PCI Express \(PCIe\) Gen2](#) product page.

Information about other AMD LogiCORE IP modules is available at the [AMD Intellectual Property](#) page. For information on pricing and availability of other AMD LogiCORE IP modules and tools, contact your [local AMD sales representative](#).

Product Specification

Figure 2-1 shows the architecture of the AXI Memory Mapped to PCI Express® core.

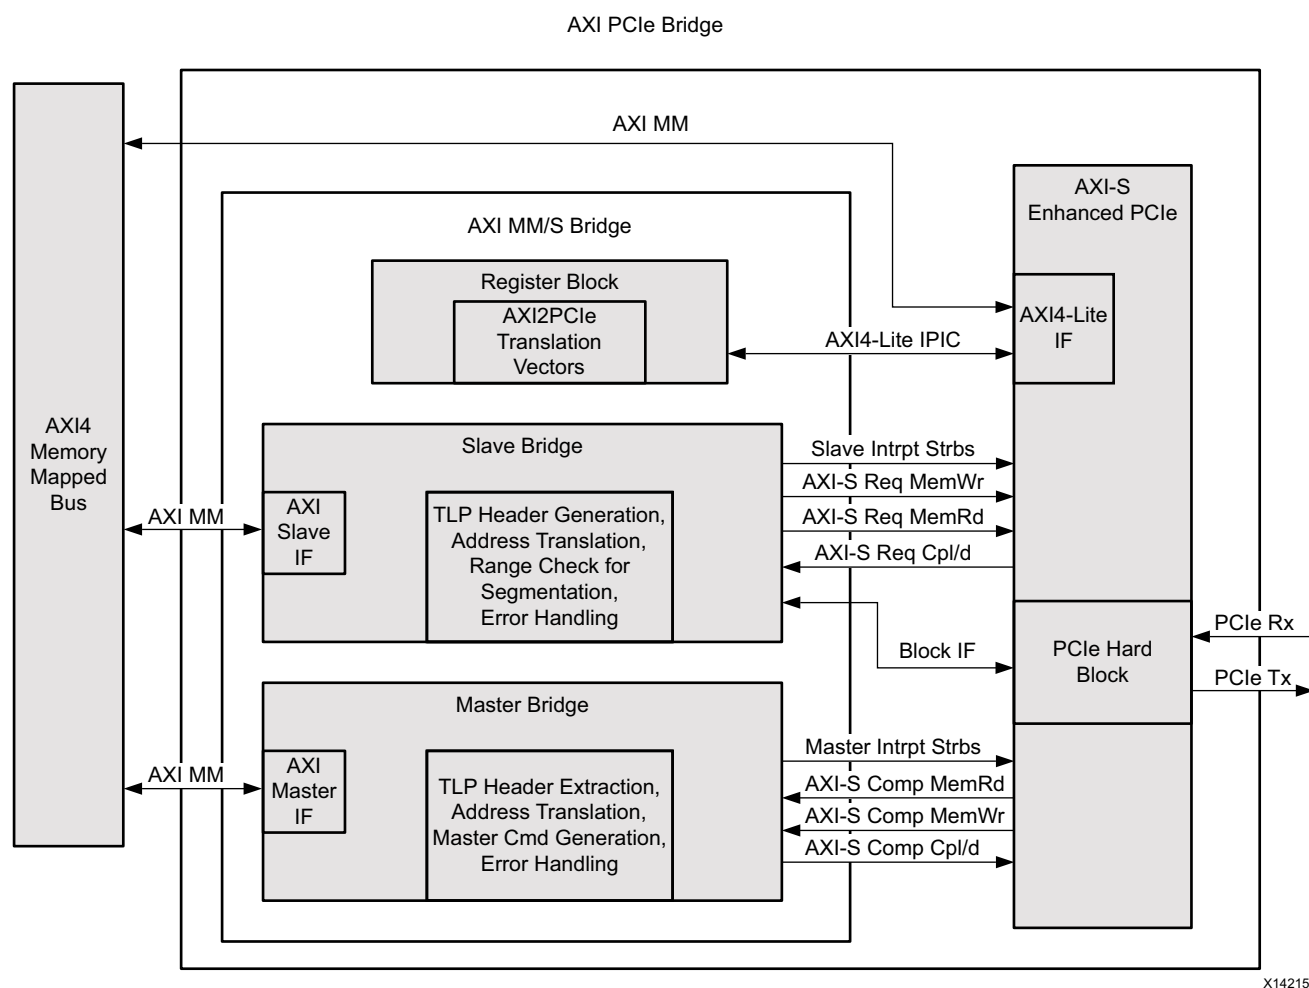


Figure 2-1: AXI Memory Mapped to PCI Express Architecture

The Register block contains registers used in the AXI Memory Mapped to PCI Express core for dynamically mapping the AXI4 memory mapped (MM) address range provided using the AXIBAR parameters to an address for PCIe® range.

The slave bridge provides termination of memory-mapped AXI4 transactions from an AXI master device (such as a processor). The slave bridge provides a way to translate addresses

that are mapped within the AXI4 memory mapped address domain to the domain addresses for PCIe. When a remote AXI master initiates a write transaction to the slave bridge, the write address and qualifiers are captured and write data is queued in a first in first out (FIFO). These are then converted into one or more MemWr TLPs, depending on the configured Max Payload Size setting, which are passed to the integrated block for PCI Express. The Slave Bridge can support up to two active AXI4 memory mapped write transactions.

When a remote AXI master initiates a read transaction to the slave bridge, the read address and qualifiers are captured and a MemRd request TLP is passed to the core and a completion timeout timer is started. Completions received through the core are correlated with pending read requests and read data is returned to the AXI master. The slave bridge is capable of handling up to eight memory mapped AXI4 read requests with pending completions.

The master bridge processes both PCIe MemWr and MemRd request TLPs received from the integrated block for PCI Express and provides a means to translate addresses that are mapped within the address for PCIe domain to the memory mapped AXI4 address domain. Each PCIe MemWr request TLP header is used to create an address and qualifiers for the memory mapped AXI4 bus and the associated write data is passed to the addressed memory mapped AXI4 Slave. The master bridge can support up to four active PCIe MemWr request TLPs.

Each PCIe MemRd request TLP header is used to create an address and qualifiers for the memory-mapped AXI4 bus. Read data is collected from the addressed memory mapped AXI4 Slave and used to generate completion TLPs which are then passed to the integrated block for PCI Express. The master bridge can handle up to four read requests with pending completions for improved AXI4 pipelining performance.

The instantiated AXI4-Stream Enhanced PCIe block contains submodules including the Requester/Completer interfaces to the AXI bridge and the Register block. The Register block contains the status, control, interrupt registers, and the AXI4-Lite interface.

Standards

The AXI Memory Mapped to PCI Express core is compliant with the AMBA[®] AXI Protocol Specification [\[Ref 7\]](#) and the PCI Express Base Specification v2.0 [\[Ref 8\]](#).

Performance and Resource Utilization

For full details about performance and resource utilization, visit the [Performance and Resource Utilization web page](#).

Port Descriptions

The interface signals for the AXI Memory Mapped to PCI Express are described in [Table 2-1](#).

Table 2-1: Top-Level Interface Signals

Signal Name	I/O	Description
Global Signals		
refclk	I	PCIe Reference Clock
axi_aresetn	I	Global reset signal for the IP. Resets GT, PCIe, and AXI interfaces.
axi_aclk_out	O	PCIe derived clock output for axi_aclk.
axi_ctl_aclk_out	O	PCIe derived clock output for axi_ctl_aclk
mmcm_lock	O	Indicates axi_aclk_out from the axi_enhanced_pcie block is stable
interrupt_out	O	Interrupt signal. This signal is asserted for as long as there exists at least one bit asserted in the Interrupt Decode register and is not masked in the Interrupt Mask register.
user_link_up	O	This signal is asserted when the PCI Express core is linked up with a link partner PCIe device.
AXI Slave Interface		
s_axi_awid[c_s_axi_id_width-1:0]	I	Slave write address ID
s_axi_awaddr[c_s_axi_addr_width-1:0]	I	Slave write address
s_axi_awregion[3:0]	I	Slave write region decode
s_axi_awlen[7:0]	I	Slave write burst length
s_axi_awsz[2:0]	I	Slave write burst size
s_axi_awburst[1:0]	I	Slave write burst type
s_axi_awvalid	I	Slave address write valid
s_axi_awready	O	Slave address write ready
s_axi_wdata[c_s_axi_data_width-1:0]	I	Slave write data
s_axi_wstrb[c_s_axi_data_width/8-1:0]	I	Slave write strobe
s_axi_wlast	I	Slave write last
s_axi_wvalid	I	Slave write valid
s_axi_wready	O	Slave write ready
s_axi_bid[c_s_axi_id_width-1:0]	O	Slave response ID
s_axi_bresp[1:0]	O	Slave write response
s_axi_bvalid	O	Slave write response valid
s_axi_bready	I	Slave response ready

Table 2-1: Top-Level Interface Signals (Cont'd)

Signal Name	I/O	Description
s_axi_arid[c_s_axi_id_width-1:0]	I	Slave read address ID
s_axi_araddr[c_s_axi_addr_width-1:0]	I	Slave read address
s_axi_arregion[3:0]	I	Slave read region decode
s_axi_arlen[7:0]	I	Slave read burst length
s_axi_arsize[2:0]	I	Slave read burst size
s_axi_arburst[1:0]	I	Slave read burst type
s_axi_arvalid	I	Slave read address valid
s_axi_arready	O	Slave read address ready
s_axi_rid[c_s_axi_id_width-1:0]	O	Slave read ID tag
s_axi_rdata[c_s_axi_data_width-1:0]	O	Slave read data
s_axi_rresp[1:0]	O	Slave read response
s_axi_rlast	O	Slave read last
s_axi_rvalid	O	Slave read valid
s_axi_rready	I	Slave read ready
AXI Master Interface		
m_axi_awaddr[c_m_axi_addr_width-1:0]	O	Master write address
m_axi_awlen[7:0]	O	Master write burst length
m_axi_awsz[2:0]	O	Master write burst size
m_axi_awburst[1:0]	O	Master write burst type
m_axi_awprot[2:0]	O	Master write protection type
m_axi_awvalid	O	Master write address valid
m_axi_awready	I	Master write address ready
m_axi_wdata[c_m_axi_data_width-1:0]	O	Master write data
m_axi_wstrb[c_m_axi_data_width/8-1:0]	O	Master write strobe
m_axi_wlast	O	Master write last
m_axi_wvalid	O	Master write valid
m_axi_wready	I	Master write ready
m_axi_bresp[1:0]	I	Master write response
m_axi_bvalid	I	Master write response valid
m_axi_bready	O	Master response ready
m_axi_araddr[c_m_axi_addr_width-1:0]	O	Master read address
m_axi_arlen[7:0]	O	Master read burst length
m_axi_arsize[2:0] ⁽¹⁾	O	Master read burst size
m_axi_arburst[1:0]	O	Master read burst type
m_axi_arprot[2:0]	O	Master read protection type

Table 2-1: Top-Level Interface Signals (Cont'd)

Signal Name	I/O	Description
m_axi_arvalid	O	Master read address valid
m_axi_arready	I	Master read address ready. This signal only responds when Bus Master Enable bit is set in the Command register within PCI Configuration Space.
m_axi_rdata[c_m_axi_data_width-1:0]	I	Master read data
m_axi_rresp[1:0]	I	Master read response
m_axi_rlast	I	Master read last
m_axi_rvalid	I	Master read valid
m_axi_rready	O	Master read ready
AXI4-Lite Control Interface		
s_axi_ctl_awaddr[31:0]	I	Slave write address
s_axi_ctl_awvalid	I	Slave write address valid
s_axi_ctl_awready	O	Slave write address ready
s_axi_ctl_wdata[31:0]	I	Slave write data
s_axi_ctl_wstrb[3:0]	I	Slave write strobe
s_axi_ctl_wvalid	I	Slave write valid
s_axi_ctl_wready	O	Slave write ready
s_axi_ctl_bresp[1:0]	O	Slave write response
s_axi_ctl_bvalid	O	Slave write response valid
s_axi_ctl_bready	I	Slave response ready
s_axi_ctl_araddr[31:0]	I	Slave read address
s_axi_ctl_arvalid	I	Slave read address valid
s_axi_ctl_arready	O	Slave read address ready
s_axi_ctl_rdata[31:0]	O	Slave read data
s_axi_ctl_rresp[1:0]	O	Slave read response
s_axi_ctl_rvalid	O	Slave read valid
s_axi_ctl_rready	I	Slave read ready
MSI Signals		
intx_msi_request	I	Legacy interrupt input (see c_interrupt_pin) when msi_enable = 0. Initiates a MSI write request when msi_enable = 1. Intx_msi_request is asserted for one clock period.
intx_msi_grant	O	Indicates legacy interrupt/MSI grant signal. The intx_msi_grant signal is asserted for one clock period when the interrupt is accepted by the PCIe core.
msi_enable	O	Indicates when MSI is enabled.

Table 2-1: Top-Level Interface Signals (Cont'd)

Signal Name	I/O	Description
msi_vector_num [4:0]	I	Indicates MSI vector to send when writing a MSI write request.
msi_vector_width [2:0]	O	Indicates the size of the MSI field (the number of MSI vectors allocated to the device).
PCIe Interface		
pci_exp_rxp[C_NO_OF_LANES-1:0][[]]	I	PCIe RX serial interface
pci_exp_rxn[C_NO_OF_LANES-1:0][[]]	I	PCIe RX serial interface
pci_exp_txp[C_NO_OF_LANES-1:0][[]]	O	PCIe TX serial interface
pci_exp_txn[C_NO_OF_LANES-1:0][[]]	O	PCIe TX serial interface

Notes:

- When a read request is received with a length that is not 1DW and is shorter than the Master AXI data width, m_axi_arsize always indicates that the requested size is equal to the Master AXI data width. The core drops the extra data when a completion packet is formed and sent back to the requester.

Bridge Parameters

Because many features in the AXI Memory Mapped to PCI Express core design can be parameterized, you can uniquely tailor the implementation of the core using only the resources required for the desired functionality. This approach also achieves the best possible performance with the lowest resource usage.

The parameters defined for the AXI Memory Mapped to PCI Express are shown in [Table 2-2](#).

Table 2-2: Top-Level Parameters

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
Bridge Parameters					
	C_PCIE_BLK_LOCN	PCIe integrated block location within FPGA	0: X0Y0 1: X0Y1 2: X0Y2 3: X1Y0 4: X1Y1	0	String
	C_XLNX_REF_BOARD	Target FPGA Board	NONE KC705_REVA KC705_REVB KC705_REVC VC707	NONE	String
G1	C_FAMILY	Target FPGA Family	kintex7, virtex7, artix7, zynq		String

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
G2	C_INCLUDE_RC	Configures the AXI bridge for PCIe to be a Root Port or an Endpoint	0: Endpoint 1: Root Port (applies only for 7 series, and Zynq 7000 SoC devices)	0	Integer
G3	C_COMP_TIMEOUT	Selects the slave bridge completion timeout counter value	0: 50 μ s 1: 50 ms	0	Integer
G4	C_INCLUDE_BAROFFSET_REG	Include the registers for high-order bits to be substituted in translation in slave bridge	0: Exclude 1: Include	0	Integer
G5	C_SUPPORTS_NARROW_BURST	Instantiates internal logic to support narrow burst transfers. Only enable when AXI master bridge generates narrow burst traffic.	0: Not supported 1: Supported	0	Integer
G6	C_AXIBAR_NUM	Number of AXI address apertures that can be accessed	1: BAR_0 enabled 2: BAR_0, BAR_1 enabled 3: BAR_0, BAR_1, BAR_2 enabled 4: BAR_0 through BAR_3 enabled 5: BAR_0 through BAR_4 enabled 6: BAR_0 through BAR_5 enabled	6	Integer
G7	C_AXIBAR_0	AXI BAR_0 aperture low address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0xFFFF_FFFF	std_logic_vector
G8	C_AXIBAR_HIGHADDR_0	AXI BAR_0 aperture high address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0x0000_0000	std_logic_vector
G9	C_AXIBAR_AS_0	AXI BAR_0 address size	0: 32 bit 1: 64 bit	0	Integer
G10	C_AXIBAR2PCIEBAR_0	Initial address translation from an AXI BAR_0 address to a PCI Express address	Valid address for PCIe ⁽²⁾	0xFFFF_FFFF	std_logic_vector
G11	C_AXIBAR_1	AXI BAR_1 aperture low address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0xFFFF_FFFF	std_logic_vector

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
G12	C_AXIBAR_HIGHADDR_1	AXI BAR_1 aperture high address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0x0000_0000	std_logic_vector
G13	C_AXIBAR_AS_1	AXI BAR_1 address size	0: 32 bit 1: 64 bit	0	Integer
G14	C_AXIBAR2PCIEBAR_1	Initial address translation from an AXI BAR_1 address to a PCI Express address	Valid address for PCIe ⁽²⁾	0xFFFF_FFFF	std_logic_vector
G15	C_AXIBAR_2	AXI BAR_2 aperture low address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0xFFFF_FFFF	std_logic_vector
G16	C_AXIBAR_HIGHADDR_2	AXI BAR_2 aperture high address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0x0000_0000	std_logic_vector
G17	C_AXIBAR_AS_2	AXI BAR_2 address size	0: 32 bit 1: 64 bit	0	Integer
G18	C_AXIBAR2PCIEBAR_2	Initial address translation from an AXI BAR_2 address to a PCI Express address	Valid address for PCIe ⁽²⁾	0xFFFF_FFFF	std_logic_vector
G19	C_AXIBAR_3	AXI BAR_3 aperture low address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0xFFFF_FFFF	std_logic_vector
G20	C_AXIBAR_HIGHADDR_3	AXI BAR_3 aperture high address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0x0000_0000	std_logic_vector
G21	C_AXIBAR_AS_3	AXI BAR_3 address size	0: 32 bit 1: 64 bit	0	Integer
G22	C_AXIBAR2PCIEBAR_3	Initial address translation from an AXI BAR_3 address to a PCI Express address	Valid address for PCIe ⁽²⁾	0xFFFF_FFFF	std_logic_vector
G23	C_AXIBAR_4	AXI BAR_4 aperture low address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0xFFFF_FFFF	std_logic_vector
G24	C_AXIBAR_HIGHADDR_4	AXI BAR_4 aperture high address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0x0000_0000	std_logic_vector
G25	C_AXIBAR_AS_4	AXI BAR_4 address size	0: 32 bit 1: 64 bit	0	Integer
G26	C_AXIBAR2PCIEBAR_4	Initial address translation from an AXI BAR_4 address to a PCI Express address	Valid address for PCIe ⁽²⁾	0xFFFF_FFFF	std_logic_vector

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
G27	C_AXIBAR_5	AXI BAR_5 aperture low address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0xFFFF_FFFF	std_logic_vector
G28	C_AXIBAR_HIGHADDR_5	AXI BAR_5 aperture high address	Valid AXI address ⁽¹⁾⁽³⁾⁽⁴⁾	0x0000_0000	std_logic_vector
G29	C_AXIBAR_AS_5	AXI BAR_5 address size	0: 32 bit 1: 64 bit	0	Integer
G30	C_AXIBAR2PCIEBAR_5	Initial address translation from an AXI BAR_5 address to a PCI Express address	Valid address for PCIe ⁽²⁾	0xFFFF_FFFF	std_logic_vector
G31	C_PCIEBAR_NUM	Number of address for PCIe apertures that can be accessed	1: BAR_0 enabled 2: BAR_0, BAR_1 enabled 3: BAR_0, BAR_1, BAR_2 enabled	3	Integer
G32	C_PCIEBAR_AS	Configures PCIEBAR aperture width to be 32 bits wide or 64 bits wide	0: Generates three 32-bit PCIEBAR address apertures. 32-bit BAR example: PCIEBAR_0 is 32 bits PCIEBAR_1 is 32 bits PCIEBAR_2 is 32 bits 1: Generates three 64 bit PCIEBAR address apertures. 64-bit BAR example: PCIEBAR_0 and PCIEBAR_1 concatenate to comprise 64-bit PCIEBAR_0. PCIEBAR_2 and PCIEBAR_3 concatenate to comprise 64-bit PCIEBAR_1. PCIEBAR_4 and PCIEBAR_5 concatenate to comprise 64-bit PCIEBAR_2	1	Integer
G33	C_PCIEBAR_LEN_0	Specifies the size of the PCIe BAR	13-31	16	Integer
G34	C_PCIEBAR2AXIBAR_0	Initial address translation from an AXI BAR_0 address to a PCI Express address	Valid AXI address	0x0000_0000	std_logic_vector

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
	C_PCIEBAR2AXIBAR_0_SEC	Defines the AXIBAR memory space (PCIe BAR_0) (accessible from PCIe) to be either secure or non-secure memory mapped.	0: Denotes a non-secure memory space 1: Marks the AXI memory space as secure	0	Integer
G35	C_PCIEBAR_LEN_1	Specifies the size of the PCIe BAR.	13-31	16	Integer
G36	C_PCIEBAR2AXIBAR_1	Initial address translation from an AXI BAR_1 address to a PCI Express address	Valid AXI address	0x0000_0000	std_logic_vector
G37	C_PCIEBAR_LEN_2	Specifies the size of the PCIe BAR.	13-31	16	Integer
G38	C_PCIEBAR2AXIBAR_2	Initial address translation from an AXI BAR_2 address to a PCI Express address.	Valid AXI address	0x0000_0000	std_logic_vector
	C_PCIEBAR2AXIBAR_2_SEC	Defines the AXIBAR memory space (PCIe BAR_2) (accessible from PCIe) to be either secure or non-secure memory mapped.	0: Denotes a non-secure memory space 1: Marks the AXI memory space as secure	0	Integer
AXI4-Lite Parameters					
G39	C_BASEADDR	Device base address Note: When configured as an RP, the minimum alignment granularity must be 256 MB. Bit [27:0] are used for Bus Number, Device Number, Function number.	Valid AXI address	0xFFFF_FFFF	std_logic_vector
G40	C_HIGHADDR	Device high address	Valid AXI address	0x0000_0000	std_logic_vector
	C_S_AXI_CTL_PROTOCOL	AXI4-Lite port connection definition to AXI Interconnect in the Vivado IP integrator.	AXI4LITE	AXI4LITE	String

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
Core for PCIe Configuration Parameters					
G41	C_NO_OF_LANES	Number of PCIe Lanes	1, 2, 4, 8: 7 series FPGAs	1	Integer
G42	C_DEVICE_ID	Device ID	16-bit vector	0x0000	std_logic_vector
G43	C_VENDOR_ID	Vendor ID	16-bit vector	0x0000	std_logic_vector
G44	C_CLASS_CODE	Class Code	24-bit vector	0x00_0000	std_logic_vector
G45	C_REV_ID	Rev ID	8-bit vector	0x00	std_logic_vector
G46	C_SUBSYSTEM_ID	Subsystem ID	16-bit vector	0x0000	std_logic_vector
G47	C_SUBSYSTEM_VENDOR_ID	Subsystem Vendor ID	16-bit vector	0x0000	std_logic_vector
	C_PCIE_USE_MODE	Specifies PCIe use mode for underlying serial transceiver wrapper usage/configuration (specific only to 7 series). This parameter ignored for Zynq 7000 SoC devices (set to 3.0).	See Table 2-4 . 1.0: For Kintex 7 325T IES (initial ES) silicon 1.1: For Virtex 7 485T IES (initial ES) silicon 3.0: For GES (general ES) silicon	1.0	String
G48	C_PCIE_CAP_SLOT_IMPLEMENTED	PCIE Capabilities Register Slot Implemented	0: No add-in card slot 1: Downstream port is connected to add-in card slot (valid only for Root Complex)	0	Integer
G49	C_REF_CLK_FREQ	REFCLK input Frequency	0: 100 MHz 1: 125 MHz 2: 250 MHz - 7 series FPGAs only	0	Integer
	C_NUM_MSI_REQ	Specifies the size of the MSI request vector for selecting the number of requested message values.	0-5	0	Integer
Memory Mapped AXI4 Parameters					
G50	C_M_AXI_DATA_WIDTH	AXI Master Bus Data width	64: 7 series FPGAs only 128: 7 series FPGAs only	64	Integer

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
G51	C_M_AXI_ADDR_WIDTH	AXI Master Bus Address width	32	32	Integer
G52	C_S_AXI_ID_WIDTH	AXI Slave Bus ID width	4	4	Integer
G53	C_S_AXI_DATA_WIDTH	AXI Slave Bus Data width	64: 7 series FPGAs only 128: 7 series FPGAs only	64	Integer
G54	C_S_AXI_ADDR_WIDTH	AXI Slave Bus Address width	32	32	Integer
G55	C_MAX_LINK_SPEED	Maximum PCIe link speed supported	0: 2.5 GT/s - 7 series 1: 5.0 GT/s - 7 series	0	Integer
G56	C_INTERRUPT_PIN	Legacy INTX pin support/select	0: No INTX support (setting for Root Port) 1: INTA selected (only allowable when core in Endpoint configuration)	0	Integer
AXI4 Slave Interconnect Parameters⁽⁶⁾					
G57	NUM_WRITE_OUTSTANDING	AXI Interconnect Slave Port Write Pipeline Depth	1: Only one active AXI AWADDR can be accepted in the AXI slave bridge for PCIe 2: Maximum of two active AXI AWADDR values can be stored in AXI slave bridge for PCIe	2	Integer
G58	NUM_READ_OUTSTANDING	AXI Interconnect Slave Port Read Pipeline Depth	1: Only one active AXI ARADDR can be accepted in AXI slave bridge PCIe. 2, 4, 8: Size of pipeline for active AXI ARADDR values to be stored in AXI slave bridge PCIe A value of 8 is not allowed for 128-bit core (Gen2 7 series) configurations. The maximum setting of this parameter value is 4.	8	Integer
AXI4 Master Interconnect Parameters					
G59	NUM_WRITE_OUTSTANDING	AXI Interconnect master bridge write address issue depth	1, 2, 4: Number of actively issued AXI AWADDR values on the AXI Interconnect to the target slave device(s).	4	Integer

Table 2-2: Top-Level Parameters (Cont'd)

Generic	Parameter Name	Description	Allowable Values	Default Value	VHDL Type
G60	NUM_READ_OUTSTANDING	AXI Interconnect master bridge read address issue depth	1, 2, 4: Number of actively issued AXI ARADDR values on the AXI Interconnect to the target slave device(s).	4	Integer

Notes:

1. This is a 32-bit address.
2. The width of this should match the address size (C_AXIBAR_AS) for this BAR.
3. The range specified must comprise a complete, contiguous power of two range, such that the range = 2^n and the n least significant bits of the Base Address are zero. The address value is a 32-bit AXI address.
4. The difference between C_AXIBAR_n and C_AXIBAR_HIGHADDR_n must be less than or equal to 0x7FFF_FFFF and greater than or equal to 0x0000_1FFF.
5. It is recommended that you do not edit these default values on the AXI Memory Mapped to PCI Express IP unless you need to reduce the resource utilization. Doing so impacts the AXI bridge performance.
6. These are the user parameters of the AXI4 Interconnect. By default, the slave bridge handles up to two AXI4 write requests and eight AXI4 read requests. The master bridge handles up to four PCIe write/read requests.

Reference Clock for PCIe Frequency Value

The `refclk` input used by the serial transceiver for PCIe must be 100 MHz, 125 MHz, or 250 MHz for 7 series and Zynq 7000 SoC device configurations. The C_REF_CLK_FREQ parameter is used to set this value, as defined in [Table 2-2](#). The `refclk` input must be fed in from a clock source that is external to the chip.

Parameter Dependencies

[Table 2-3](#) lists the parameter dependencies. The parameters are defined in [Table 2-2](#).

Table 2-3: Parameter Dependencies

Generic	Parameter	Affects	Depends	Description
Bridge Parameters				
G1	C_FAMILY	G2, G41, G49, G55		
G2	C_INCLUDE_RC		G1	
	C_MSI_DECODE_ENABLE		G2	TRUE: Allows the bridge to decode incoming MSI packet. FALSE: Ignore incoming MSI packet and treats it as a regular Memory Write packet.
G3	C_COMP_TIMEOUT			

Table 2-3: Parameter Dependencies (Cont'd)

Generic	Parameter	Affects	Depends	Description
G4	C_INCLUDE_BAROFFSET_REG	G10, G14, G18, G22, G26, G30	G6	If G4 = 0, then G10, G14, G18, G22, G26 and G30 have no meaning. The number of registers included is set by G6.
G5	C_SUPPORTS_NARROW_BURST			
G6	C_AXIBAR_NUM	G4, G7 - G30		If G6 = 1, then G7 - G10 are enabled. If G6 = 2, then G7 - G14 are enabled. If G6 = 3, then G7 - G18 are enabled. If G6 = 4, then G7 - G22 are enabled. If G6 = 5, then G7 - G26 are enabled. If G6 = 6, then G7 - G30 are enabled.
G7	C_AXIBAR_0	G8	G6, G8	G7 and G8 define the range in AXI memory space that is responded to by this device (AXIBAR)
G8	C_AXIBAR_HIGHADDR_0	G7	G6, G7	G7 and G8 define the range in AXI memory space that is responded to by this device (AXIBAR)
G9	C_AXIBAR_AS_0		G6	
G10	C_AXIBAR2PCIEBAR_0		G4, G6	Meaningful when G4 = 1.
G11	C_AXIBAR_1	G12	G12	G11 and G12 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G12	C_AXIBAR_HIGHADDR_1	G11	G6, G11	G11 and G12 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G13	C_AXIBAR_AS_1		G6	
G14	C_AXIBAR2PCIEBAR_1		G4, G6	Meaningful when G4 = 1.
G15	C_AXIBAR_2	G16	G16	G15 and G16 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G16	C_AXIBAR_HIGHADDR_2	G15	G6, G15	G15 and G16 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G17	C_AXIBAR_AS_2		G6	
G18	C_AXIBAR2PCIEBAR_2		G4, G6	Meaningful when G4 = 1.
G19	C_AXIBAR_3	G20	G20	G19 and G20 define the range in AXI-memory space that is responded to by this device (AXIBAR)

Table 2-3: Parameter Dependencies (Cont'd)

Generic	Parameter	Affects	Depends	Description
G20	C_AXIBAR_HIGHADDR_3	G19	G6, G19	G19 and G20 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G21	C_AXIBAR_AS_3		G6	
G22	C_AXIBAR2PCIEBAR_3		G4, G6	Meaningful when G4 = 1.
G23	C_AXIBAR_4	G24	G24	G23 and G24 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G24	C_AXIBAR_HIGHADDR_4	G23	G6, G23	G23 and G24 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G25	C_AXIBAR_AS_4		G6	
G26	C_AXIBAR2PCIEBAR_4		G4, G6	Meaningful if G4 = 1.
G27	C_AXIBAR_5	G28	G28	G27 and G28 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G28	C_AXIBAR_HIGHADDR_5	G27	G6, G27	G27 and G28 define the range in AXI-memory space that is responded to by this device (AXIBAR)
G29	C_AXIBAR_AS_5		G6	
G30	C_AXIBAR2PCIEBAR_5		G4, G6	Meaningful if G4 = 1.
G31	C_PCIEBAR_NUM	G33-G38		If G31 = 1, then G32, G33 are enabled. If G31 = 2, then G32 - G36 are enabled. If G31 = 3, then G32 - G38 are enabled
G32	C_PCIEBAR_AS			
G33	C_PCIEBAR_LEN_0	G34	G31	
G34	C_PCIEBAR2AXIBAR_0		G31, G33	Only the high-order bits above the length defined by G33 are meaningful.
G35	C_PCIEBAR_LEN_1	G36	G31	
G36	C_PCIEBAR2AXIBAR_1		G31, G35	Only the high-order bits above the length defined by G35 are meaningful.
G37	C_PCIEBAR_LEN_2	G38	G31	
G38	C_PCIEBAR2AXIBAR_2		G31, G37	Only the high-order bits above the length defined by G37 are meaningful.

Table 2-3: Parameter Dependencies (Cont'd)

Generic	Parameter	Affects	Depends	Description						
Core for PCIe Configuration Parameters										
G41	C_NO_OF_LANES		G1, G50, G53	<table><tr><th>Parameter Setting</th><th>Result</th></tr><tr><td>G1 = Kintex7 (for FBG484 max X4), Artix7 (CPG236 max x2), zc7030 and zc7015 (max x4). For all other devices and family max G41 is 8. G50 = G53 = 64</td><td>G41 = 1, 2, or 4 (Gen1) or G41 = 1 or 2 (Gen2)</td></tr><tr><td>G1 = Kintex 7 G50 = G53 = 128</td><td>G41 = 1, 2, 4, or 8 (Gen1), or 1, 2, or 4 (Gen2)</td></tr></table>	Parameter Setting	Result	G1 = Kintex7 (for FBG484 max X4), Artix7 (CPG236 max x2), zc7030 and zc7015 (max x4). For all other devices and family max G41 is 8. G50 = G53 = 64	G41 = 1, 2, or 4 (Gen1) or G41 = 1 or 2 (Gen2)	G1 = Kintex 7 G50 = G53 = 128	G41 = 1, 2, 4, or 8 (Gen1), or 1, 2, or 4 (Gen2)
				Parameter Setting	Result					
				G1 = Kintex7 (for FBG484 max X4), Artix7 (CPG236 max x2), zc7030 and zc7015 (max x4). For all other devices and family max G41 is 8. G50 = G53 = 64	G41 = 1, 2, or 4 (Gen1) or G41 = 1 or 2 (Gen2)					
G1 = Kintex 7 G50 = G53 = 128	G41 = 1, 2, 4, or 8 (Gen1), or 1, 2, or 4 (Gen2)									
G42	C_DEVICE_ID									
G43	C_VENDOR_ID									
G44	C_CLASS_CODE									
G45	C_REV_ID									
G46	C_SUBSYSTEM_ID									
G47	C_SUBSYSTEM_VENDOR_ID									
G48	C_PCIE_CAP_SLOT_IMPLEMENTED		G2	If G2 = 0, G48 is not meaningful						
G49	C_REF_CLK_FREQ		G1							
Memory-Mapped AXI4 Bus Parameters										
G50	C_M_AXI_DATA_WIDTH	G53	G1, G41, G53	G50 must be equal to G53						
G51	C_M_AXI_ADDR_WIDTH	G54	G54	G51 must be equal to G54						
G52	C_S_AXI_ID_WIDTH									
G53	C_S_AXI_DATA_WIDTH	G50	G1, G41, G50	G53 must be equal to G50						
G54	C_S_AXI_ADDR_WIDTH	G51	G51	G54 must be equal to G51						
G55	C_MAX_LINK_SPEED		G1							
G56	C_INTERRUPT_PIN									

Table 2-4 summarizes the relationship between the IP design parameters, C_FAMILY and C_PCIE_USE_MODE. The C_PCIE_USE_MODE is used to specify the 7 series (and derivative FPGA technology) serial transceiver wrappers to use based on the silicon version. Initial Engineering Silicon (IES) as well as General Engineering Silicon (GES) must be specified.

Table 2-4: Silicon Version Specification

C_FAMILY	C_PCIE_USE_MODE
Kintex 7	1.1 = for Kintex 7 325T IES (initial silicon) 3.0 = for GES (general silicon)
Virtex 7	1.1 = for Virtex 7 485T IES (initial silicon) 3.0 = for GES and Production (general silicon and Production silicon)
Artix 7	1.0 = for IES (initial silicon) as well as GES and Production (General silicon and Production silicon) to use latest serial transceiver wrappers (only allowable value)
Zynq	Not applicable. (set internally = 3.0)

Memory Map

The memory map shown in [Table 2-5](#) shows the address mapping for the AXI Memory Mapped to PCI Express core. These registers are described in more detail in the following section. All registers are accessed through the AXI4-Lite Control Interface and are offset from C_BASEADDR. During a reset, all registers return to default values.

Table 2-5: Register Memory Map

Accessibility	Offset	Contents	Location
RO - EP, R/W - RC	0x000 - 0x124	PCIe Configuration Space Header	Part of integrated PCIe configuration space.
RO	0x128	Vendor-Specific Enhanced Capability (VSEC) Capability	VSEC of integrated PCIe configuration space.
RO	0x12C	VSEC Header	
RO	0x130	Bridge Info	AXI bridge defined memory-mapped register space.
RO - EP, R/W - RC	0x134	Bridge Status and Control	
R/W	0x138	Interrupt Decode	
R/W	0x13C	Interrupt Mask	
RO - EP, R/W - RC	0x140	Bus Location	
RO	0x144	Physical-Side Interface (PHY) Status/Control	
RO - EP, R/W - RC	0x148	Root Port Status/Control	
RO - EP, R/W - RC	0x14C	Root Port MSI Base 1	
RO - EP, R/W - RC	0x150	Root Port MSI Base 2	
RO - EP, R/W - RC	0x154	Root Port Error FIFO Read	
RO - EP, R/W - RC	0x158	Root Port Interrupt FIFO Read 1	
RO - EP, R/W - RC	0x15C	Root Port Interrupt FIFO Read 2	
RO	0x160 - 0x1FF	Reserved (zeros returned on read)	
RO	0x200	VSEC Capability 2	

Table 2-5: Register Memory Map (Cont'd)

Accessibility	Offset	Contents	Location
RO	0x204	VSEC Header 2	
R/W	0x208 - 0x234	AXI Base Address Translation Configuration Registers	AXI bridge defined memory-mapped space.
RO	0x238 - 0xFFFF	Reserved (zeros returned on read)	

PCIe Configuration Space Header

The PCIe Configuration Space Header is a memory aperture for accessing the core for PCIe configuration space. For 7 series devices, this area is read-only when configured as an Endpoint. Writes are permitted for some registers when a 7 series device is configured as a Root Port. Special access modes can be enabled using the PHY Status/Control register. All reserved or undefined memory-mapped addresses must return zero and writes have no effect.

VSEC Capability Register (Offset 0x128)

The VSEC Capability register (described in [Table 2-6](#)) allows the memory space of the core to appear as though it is a part of the underlying core configuration space. The VSEC is inserted immediately following the last enhanced capability structure in the underlying block. VSEC is defined in §7.18 of the *PCI Express Base Specification, v1.1* (§7.19 of v2.0) [\[Ref 8\]](#).

Table 2-6: VSEC Capability Register

Bits	Name	Core Access	Reset Value	Description
15:0	VSEC Capability ID	RO	0x000B	PCI-SIG [®] defined ID identifying this enhanced capability as a vendor-specific capability. Hardcoded to 0x000B.
19:16	Capability Version	RO	0x1	Version of this capability structure. Hardcoded to 0x1.
31:20	Next Capability Offset	RO	0x200	Offset to next capability. Hardcoded to 0x0200.

VSEC Header Register (Offset 0x12C)

The VSEC Header register (described in [Table 2-7](#)) provides a unique (within a given vendor) identifier for the layout and contents of the VSEC structure, as well as its revision and length.

VSEC Header register is part of the PCI Express Hard Block which contains Loopback Control registers. For more information about Loopback Control Registers, see the AMD Defined Vendor Specific Capability section in the *7 Series FPGAs Integrated Block for PCI Express Product Guide* (PG054) [\[Ref 2\]](#).

Table 2-7: VSEC Header Register

Bits	Name	Core Access	Reset Value	Description
15:0	VSEC ID	RO	0x0001	ID value uniquely identifying the nature and format of this VSEC structure.
19:16	VSEC REV	RO	0	Version of this capability structure. Hardcoded to 0h.
31:20	VSEC Length	RO	0x038	Length of the entire VSEC capability structure, in bytes, including the VSEC capability register. Hardcoded to 0x038 (56 decimal).

Bridge Info Register (Offset 0x130)

The Bridge Info register (described in [Table 2-8](#)) provides general configuration information about the AXI4-Stream Bridge. Information in this register is static and does not change during operation.

Table 2-8: Bridge Info Register

Bits	Name	Core Access	Reset Value	Description
0	Gen2 Capable	RO	0	If set, underlying integrated block supports PCIe Gen2 speed.
1	Root Port Present	RO	0	Indicates the underlying integrated block is a Root Port when this bit is set. If set, Root Port registers are present in this interface.
2	Up Config Capable	RO		Indicates the underlying integrated block is upconfig capable when this bit is set.
15:3	Reserved	RO	0	Reserved
18:16	ECAM Size	RO	0	Size of Enhanced Configuration Access Mechanism (ECAM) Bus Number field, in number of bits. If ECAM window is present, value is between 1 and 8. If not present, value is 0. Total address bits dedicated to ECAM window is 20+(ECAM Size). The size of the ECAM is determined by the parameter settings of C_BASEADDR and C_HIGHADDR.
31:19	Reserved	RO	0	Reserved

Bridge Status and Control Register (Offset 0x134)

The Bridge Status and Control register (described in [Table 2-9](#)) provides information about the current state of the AXI4-Stream Bridge. It also provides control over how reads and writes to the Core Configuration Access aperture are handled.

Table 2-9: Bridge Status and Control Register

Bits	Name	Core Access	Reset Value	Description
0	ECAM Busy	RO	0	Indicates an ECAM access is in progress (waiting for completion). This bit is tied to 0.
7:1	Reserved	RO	0	Reserved
8	Global Disable	RW	0	When set, disables interrupt line from being asserted. Does not prevent bits in Interrupt Decode register from being set.
15:9	Reserved	RO	0	Reserved
16	RW1C as RW	RW	0	When set, allows writing to core registers which are normally RW1C.
17	RO as RW	RW	0	When set, allows writing to certain registers which are normally RO. (Only supported for 7-series and Zynq 7000 SoC device cores.)
31:18	Reserved	RO	0	Reserved

Interrupt Decode Register (Offset 0x138)

The Interrupt Decode register (described in [Table 2-10](#)) provides a single location where the host processor interrupt service routine can determine what is causing the interrupt to be asserted and how to clear the interrupt. Writing a 1'b1 to any bit of the Interrupt Decode register clears that bit except for the Correctable, Non-Fatal, and Fatal bits.

Follow this sequence to clear the Correctable, Non-Fatal, and Fatal bits:

1. Clear the Root Port Error FIFO (0x154) by performing first a read, followed by write-back of the same register.
2. Read Root Port Status/Control Register (0x148) bit 16, and ensure that the Error FIFO is empty.
Note: If the error FIFO is still not empty, repeat [step 1](#) and [step 2](#) until the Error FIFO is empty.
3. Write to the Interrupt Decode Register (0x138) with 1 to the appropriate error bit to clear it.



IMPORTANT: An asserted bit in the Interrupt Decode register does not cause the interrupt line to assert unless the corresponding bit in the Interrupt Mask register is also set.

Table 2-10: Interrupt Decode Register

Bits	Name	Core Access	Reset Value	Description
0	Link Down	RW1C	0	Indicates that Link-Up on the PCI Express link was lost. Not asserted unless link-up had previously been seen.
1	ECRC Error	RW1C	0	Indicates received packet failed ECRC check. (Only applicable to 7 series and Zynq 7000 SoC device cores.)

Table 2-10: Interrupt Decode Register (Cont'd)

Bits	Name	Core Access	Reset Value	Description
2	Streaming Error	RW1C	0	Indicates a gap was encountered in a streamed packet on the TX interface (RW, RR, or CC).
3	Hot Reset	RW1C	0	Indicates a Hot Reset was detected (Only applicable to Endpoint cores).
4	Reserved	RO	0	Reserved
7:5	Cfg Completion Status	RW1C	0	Indicates config completion status.
8	Cfg Timeout	RW1C	0	Indicates timeout on an ECAM access. (Only applicable to Root Port cores.)
9	Correctable	RW1C	0	Indicates a correctable error message was received. Requester ID of error message should be read from the Root Port FIFO. (Only applicable to Root Port cores.)
10	Non-Fatal	RW1C	0	Indicates a non-fatal error message was received. Requester ID of error message should be read from the Root Port FIFO. (Only applicable to Root Port cores.)
11	Fatal	RW1C	0	Indicates a fatal error message was received. Requester ID of error message should be read from the Root Port FIFO. (Only applicable to Root Port cores.)
15:12	Reserved	RO	0	Reserved
16	INTx Interrupt Received	RW1C	0	Indicates an INTx interrupt was received. Interrupt details should be read from the Root Port FIFO. (Only applicable to Root Port cores.)
17	MSI Interrupt Received	RW1C	0	Indicates an MSI(x) interrupt was received. Interrupt details should be read from the Root Port FIFO. (Only applicable to Root Port cores.)
19:18	Reserved	RO	0	Reserved
20	Slave Unsupported Request	RW1C	0	Indicates that a completion TLP was received with a status of 0b001 - Unsupported Request.
21	Slave Unexpected Completion	RW1C	0	Indicates that a completion TLP was received that was unexpected.
22	Slave Completion Timeout	RW1C	0	Indicates that the expected completion TLP(s) for a read request for PCIe was not returned within the time period selected by the C_COMP_TIMEOUT parameter.
23	Slave Error Poison	RW1C	0	Indicates the EP bit was set in a completion TLP.
24	Slave Completer Abort	RW1C	0	Indicates that a completion TLP was received with a status of 0b100 - Completer Abort.

Table 2-10: Interrupt Decode Register (Cont'd)

Bits	Name	Core Access	Reset Value	Description
25	Slave Illegal Burst	RW1C	0	Indicates that a burst type other than INCR was requested by the AXI master.
26	Master DECERR	RW1C	0	Indicates a Decoder Error (DECERR) response was received.
27	Master SLVERR	RW1C	0	Indicates a Slave Error (SLVERR) response was received.
28	Master Error Poison	RW1C	0	Indicates an EP bit was set in a MemWR TLP for PCIe.
31:29	Reserved	RO	0	Reserved

Interrupt Mask Register (Offset 0x13C)

The Interrupt Mask register controls whether each individual interrupt source can cause the interrupt line to be asserted. A one in any location allows the interrupt source to assert the interrupt line. The Interrupt Mask register initializes to all zeros. Therefore, by default no interrupt is generated for any event. Table 2-11 describes the Interrupt Mask register bits and values.

Table 2-11: Interrupt Mask Register

Bits	Name	Core Access	Reset Value	Description
0	Link Down	RW	0	Enables interrupts for Link Down events when bit is set.
1	ECRC Error	RW	0	Enables interrupts for ECRC Error events when bit is set.
2	Streaming Error	RW	0	Enables interrupts for Streaming Error events when bit is set.
3	Hot Reset	RW	0	Enables interrupts for Hot Reset events when bit is set. (Only writable for EP configurations, otherwise = 0)
4	Reserved	RO	0	Reserved
7:5	Cfg Completion Status	RW	0	Enables interrupts for config completion status. (Only writable for Root Port Configurations, otherwise = 0)
8	Cfg Timeout	RO	0	Enables interrupts for Config (Cfg) Timeout events when bit is set. (Only writable for Root Port Configurations, otherwise = 0)
9	Correctable	RO	0	Enables interrupts for Correctable Error events when bit is set. (Only writable for Root Port Configurations, otherwise = 0)
10	Non-Fatal	RO	0	Enables interrupts for Non-Fatal Error events when bit is set. (Only writable for Root Port Configurations, otherwise = 0)
11	Fatal	RO	0	Enables interrupts for Fatal Error events when bit is set. (Only writable for Root Port Configurations, otherwise = 0)
15:12	Reserved	RO	0	Reserved
16	INTx Interrupt Received	RO	0	Enables interrupts for INTx Interrupt events when bit is set. (Only writable for Root Port Configurations, otherwise = 0)

Table 2-11: Interrupt Mask Register (Cont'd)

Bits	Name	Core Access	Reset Value	Description
17	MSI Interrupt Received	RO	0	Enables interrupts for MSI Interrupt events when bit is set. (Only writable for Root Port Configurations, otherwise =0)
19:18	Reserved	RO	0	Reserved
20	Slave Unsupported Request	RW	0	Enables the Slave Unsupported Request interrupt when bit is set.
21	Slave Unexpected Completion	RW	0	Enables the Slave Unexpected Completion interrupt when bit is set.
22	Slave Completion Timeout	RW	0	Enables the Slave Completion Timeout interrupt when bit is set.
23	Slave Error Poison	RW	0	Enables the Slave Error Poison interrupt when bit is set.
24	Slave Completer Abort	RW	0	Enables the Slave Completer Abort interrupt when bit is set.
25	Slave Illegal Burst	RW	0	Enables the Slave Illegal Burst interrupt when bit is set.
26	Master DECERR	RW	0	Enables the Master DECERR interrupt when bit is set.
27	Master SLVERR	RW	0	Enables the Master SLVERR interrupt when bit is set.
28	Master Error Poison	RW	0	Enables the Master Error Poison interrupt when bit is set.
31:29	Reserved	RO	0	Reserved

Bus Location Register (Offset 0x140)

The Bus Location register reports the Bus, Device, and Function number, and the Port number for the PCIe port ([Table 2-12](#)).

Table 2-12: Bus Location Register

Bits	Name	Core Access	Reset Value	Description
2:0	Function Number	RO	0	Function number of the port for PCIe. Hard-wired to 0.
7:3	Device Number	RO	0	Device number of port for PCIe. For Endpoint, this register is RO and is set by the Root Port.
15:8	Bus Number	RO	0	Bus number of port for PCIe. For Endpoint, this register is RO and is set by the external Root Port.
23:16	Port Number	RW	0	Sets the Port number field of the Link Capabilities register. EP: Read Only on all devices except for Spartan-6 FPGA. RP: Read and writeable.
31:24	Reserved	RO	0	Reserved

PHY Status/Control Register (Offset 0x144)

The PHY Status/Control register (described in [Table 2-13](#)) provides the status of the current PHY state, as well as control of speed and rate switching for Gen2-capable cores.

Table 2-13: PHY Status/Control Register

Bits	Name	Core Access	Reset Value	Description
0	Link Rate	RO	0	Reports the current link rate. <ul style="list-style-type: none"> • 0b = 2.5 GT/s • 1b = 5.0 GT/s
2:1	Link Width	RO	0	Reports the current link width. 00b = x1, 01b = x2, 10b = x4, 11b = x8.
8:3	LTSSM State	RO	0	Reports the current Link Training and Status State Machine (LTSSM) state. Encoding is specific to the underlying integrated block.
10:9	Lane Reversal	RO	0	Reports the current lane reversal mode. <ul style="list-style-type: none"> • 00b: No reversal • 01b: Lanes 1:0 reversed • 10b: Lanes 3:0 reversed • 11b: Lanes 7:0 reversed
11	Link Up	RO	0	Reports the current PHY Link-up state. <ul style="list-style-type: none"> • 1b: Link up • 0b: Link down Link up indicates the core has achieved link up status, meaning the LTSSM is in the L0 state and the core can send/receive data packets.
15:12	Reserved	RO	0	Reserved
17:16	Directed Link Width	RW	0	Specifies completer link width for a directed link change operation. Only acted on when Directed Link Change specifies a width change. <ul style="list-style-type: none"> • 00b: x1 • 01b: x2 • 10b: x4 • 11b: x8
18	Directed Link Speed	RW	0	Specifies completer link speed for a directed link change operation. Only acted on when Directed Link Change specifies a speed change. <ul style="list-style-type: none"> • 0b: 2.5 GT/s • 1b: 5.0 GT/s

Table 2-13: PHY Status/Control Register (Cont'd)

Bits	Name	Core Access	Reset Value	Description
19	Directed Link Autonomous	RW	0	Specifies link reliability or autonomous for directed link change operation. <ul style="list-style-type: none"> 0b: Link reliability 1b: Autonomous
21:20	Directed Link Change	RW	0	Directs LTSSM to initiate a link width and/or speed change. <ul style="list-style-type: none"> 00b: No change 01b: Force link width 10b: Force link speed 11b: Force link width and speed
31:22	Reserved	RO	0	Reserved

Root Port Status/Control Register (Offset 0x148)

The Root Port Status/Control register provides access to Root Port specific status and control. This register is only implemented for Root Port cores. For non-Root Port cores, reads return 0 and writes are ignored (described in [Table 2-14](#)).

Table 2-14: Root Port Status/Control Register

Bits	Name	Core Access	Reset Value	Description
0	Bridge Enable	RW	0	When set, allows the reads and writes to the AXIBARs to be presented on the PCIe bus. Root Port software needs to write 1 to this bit when enumeration is done. AXI Enhanced PCIe Bridge clears this location when link up to link down transition occurs. Default is set to 0.
15:1	Reserved	RO	0	Reserved.
16	Error FIFO Not Empty	RO	0	Indicates that the Root Port Error FIFO has data to read.
17	Error FIFO Overflow	RW1C	0	Indicates that the Root Port Error FIFO overflowed and an error message was dropped. Writing a 1 clears the overflow status.
18	Interrupt FIFO Not Empty	RO	0	Indicates that the Root Port Interrupt FIFO has data to read.
19	Interrupt FIFO Overflow	RW1C	0	Indicates that the Root Port Interrupt FIFO overflowed and an interrupt message was dropped. Writing a 1 clears the overflow status
31:20	Reserved	RO	0	Reserved.

Root Port MSI Base Register 1 (Offset 0x14C)

The Root Port MSI Base register contains the upper 32-bits of the 64-bit MSI address (described in [Table 2-15](#)).

For EP configurations, read returns zero.

Table 2-15: Root Port MSI Base Register 1

Bits	Name	Core Access	Reset Value	Description
31:0	MSI Base	RW	0	4Kb-aligned address for MSI interrupts. In case of 32-bit MSI, it returns 0 but captures the upper 32-bits of the MSI address in case of 64-bit MSI.

Root Port MSI Base Register 2 (Offset 0x150)

The Root Port MSI Base register 2 (described in [Table 2-16](#)) sets the address window in Root Port cores used for MSI interrupts. MemWr TLPs to this address are interpreted as MSI interrupts. MSI TLPs are interpreted based on the address programmed in this register. For EP configurations, a read returns zero. However, the AXI Memory Mapped to PCI Express core does not support MSI-X.

Table 2-16: Root Port MSI Base Register 2

Bits	Name	Core Access	Reset Value	Description
11:0	Reserved	RO	0	Reserved
31:12	MSI Base	RW	0	4 Kb-aligned address for MSI interrupts.

Root Port Error FIFO Read Register (Offset 0x154)

Reads from this location return queued error (Correctable/Non-fatal/Fatal) messages. Data from each read follows the format shown in [Table 2-17](#). For EP configurations, read returns zero.

Reads are non-destructive. Removing the message from the FIFO requires a write. The write value is ignored.

Table 2-17: Root Port Error FIFO Read Register

Bits	Name	Core Access	Reset Value	Description
15:0	Requester ID	RWC	0	Requester ID belonging to the requester of the error message.
17:16	Error Type	RWC	0	Indicates the type of the error. 00b: Correctable 01b: Non-Fatal 10b: Fatal 11b: Reserved

Table 2-17: Root Port Error FIFO Read Register

Bits	Name	Core Access	Reset Value	Description
18	Error Valid	RWC	0	Indicates whether read succeeded. 1b: Success 0b: No message to read
31:19	Reserved	RO	0	Reserved

Root Port Interrupt FIFO Read Register 1 (Offset 0x158)

Reads from this location return queued interrupt messages. Data from each read follows the format shown in Table 2-18. For MSI interrupts, the message payload is presented in the Root Port Interrupt FIFO Read 2 register. The interrupt-handling flow should be to read this register first, immediately followed by the Root Port Interrupt FIFO Read 2 register. For non-Root Port cores, reads return zero.

Note: Reads are non-destructive. Removing the message from the FIFO requires a write to either this register or the Root Port Interrupt FIFO Read 2 register. The write value is ignored.

Table 2-18: Root Port Interrupt FIFO Read Register 1

Bits	Name	Core Access	Reset Value	Description
15:0	Requester ID	RWC	0	Requester ID belonging to the requester of the error message.
26:16	MSI Address	RWC	0	For MSI interrupts, contains address bits 12:2 from the TLP address field.
28:27	Interrupt Line	RWC	0	Indicates interrupt line used. 00b: INTA 01b: INTB 10b: INTC 11b: INTD For MSI, this field is set to 00b and should be ignored.
29	Interrupt Assert	RWC	0	Indicates assert or deassert for INTx. 1b: Assert 0b: Deassert For MSI, this field is set to 0b and should be ignored.
30	MSI Interrupt	RWC	0	Indicates whether interrupt is MSI or INTx. 1b = MSI, 0b = INTx.
31	Interrupt Valid	RWC	0	Indicates whether read succeeded. 1b: Success 0b: No interrupt to read

Root Port Interrupt FIFO Read Register 2 (Offset 0x15C)

Reads from this location return queued interrupt messages. Data from each read follows the format shown in Table 2-19. For MSI interrupts, the message payload is presented in this register, while the header information is presented in the Root Port Interrupt FIFO Read 1

register. The interrupt-handling flow should be to read the Root Port Interrupt FIFO Read 1 register first, immediately followed by this register. For non-Root Port cores, reads return 0. For INTx interrupts, reads return zero.

Note: Reads are non-destructive. Removing the message from the FIFO requires a write to either this register or the Root Port Interrupt FIFO Read 1 register (write value is ignored).

Table 2-19: Root Port Interrupt FIFO Read Register 2

Bits	Name	Core Access	Reset Value	Description
15:0	Message Data	RWC	0	Payload for MSI messages.
31:16	Reserved	RO	0	Reserved

VSEC Capability Register 2 (Offset 0x200)

The VSEC capability register (described in [Table 2-20](#)) allows the memory space for the core to appear as though it is a part of the underlying integrated block PCIe configuration space. The VSEC is inserted immediately following the last enhanced capability structure in the underlying block. VSEC is defined in §7.18 of the PCI Express Base Specification, v1.1 (§7.19 of v2.0) [\[Ref 8\]](#).

This register is only included if C_INCLUDE_BAR_OFFSET_REG = 1.

Table 2-20: VSEC Capability Register 2

Bits	Name	Core Access	Reset Value	Description
15:0	VSEC Capability ID	RO	0x000B	PCI-SIG defined ID identifying this Enhanced Capability as a Vendor-Specific capability. Hardcoded to 0x000B.
19:16	Capability Version	RO	0x1	Version of this capability structure. Hardcoded to 0x1.
31:20	Next Capability Offset	RO	0x000	Offset to next capability.

VSEC Header Register 2 (Offset 0x204)

The VSEC Header Register 2 (described in [Table 2-21](#)) provides a unique (within a given vendor) identifier for the layout and contents of the VSEC structure, as well as its revision and length. VSEC Header Register 2 is part of the AXI Memory Mapped to PCI Express core that contains AXI Base Address Translation Configuration Registers which start immediately after VSEC Header Register 2 (Offset 0x208).

This register is only included if C_INCLUDE_BAR_OFFSET_REG = 1.

Table 2-21: VSEC Header Register 2

Bits	Name	Core Access	Reset Value	Description
15:0	VSEC ID	RO	0x0002	ID value uniquely identifying the nature and format of this VSEC structure.
19:16	VSEC REV	RO	0x0	Version of this capability structure. Hardcoded to 0x0.
31:20	VSEC Length	RO	0x038	Length of the entire VSEC Capability structure, in bytes, including the VSEC Capability register. Hardcoded to 0x038 (56 decimal).

AXI Base Address Translation Configuration Registers (Offset 0x208 - 0x234)

The AXI Base Address Translation Configuration Registers and their offsets are shown in [Table 2-22](#) and the register bits are described in [Table 2-23](#). This set of registers can be used in two configurations based on the top-level parameter C_AXIBAR_AS_n. When the BAR is set to a 32-bit address space, then the translation vector should be placed into the AXIBAR2PCIEBAR_nL register where n is the BAR number. When the BAR is set to a 64-bit address space, then the most significant 32 bits are written into the AXIBAR2PCIEBAR_nU and the least significant 32 bits are written into AXIBAR2PCIEBAR_nL. These registers are only included if C_INCLUDE_BAR_OFFSET_REG = 1.

Table 2-22: AXI Base Address Translation Configuration Registers

Offset	Bits	Register Mnemonic
0x208	31-0	AXIBAR2PCIEBAR_0U
0x20C	31-0	AXIBAR2PCIEBAR_0L
0x210	31-0	AXIBAR2PCIEBAR_1U
0x214	31-0	AXIBAR2PCIEBAR_1L
0x218	31-0	AXIBAR2PCIEBAR_2U
0x21C	31-0	AXIBAR2PCIEBAR_2L
0x220	31-0	AXIBAR2PCIEBAR_3U
0x224	31-0	AXIBAR2PCIEBAR_3L
0x228	31-0	AXIBAR2PCIEBAR_4U
0x22C	31-0	AXIBAR2PCIEBAR_4L
0x230	31-0	AXIBAR2PCIEBAR_5U
0x234	31-0	AXIBAR2PCIEBAR_5L

Table 2-23: AXI Base Address Translation Configuration Register Bit Definitions

Bits	Name	Core Access	Reset Value	Description
31-0	Lower Address	R/W	C_AXIBAR2PCIEBAR_0(31 to 0)	To create the address for PCIe—this is the value substituted for the least significant 32 bits of the AXI address.
31-0	Upper Address	R/W	if (C_AXIBAR2PCIEBAR_0 = 64 bits), then reset value = C_AXIBAR2PCIEBAR_0(63 to 32) if (C_AXIBAR2PCIEBAR_0 = 32 bits), then reset value = 0x00000000	To create the address for PCIe—this is the value substituted for the most significant 32 bits of the AXI address.
31-0	Lower Address	R/W	C_AXIBAR2PCIEBAR_1(31 to 0)	To create the address for PCIe—this is the value substituted for the least significant 32 bits of the AXI address.
31-0	Upper Address	R/W	if (C_AXIBAR2PCIEBAR_1 = 64 bits), then reset value = C_AXIBAR2PCIEBAR_1(63 to 32) if (C_AXIBAR2PCIEBAR_1 = 32 bits), then reset value = 0x00000000	To create the address for PCIe— this is the value substituted for the most significant 32 bits of the AXI address.
31-0	Lower Address	R/W	C_AXIBAR2PCIEBAR_2(31 to 0)	To create the address for PCIe—this is the value substituted for the least significant 32 bits of the AXI address.
31-0	Upper Address	R/W	if (C_AXIBAR2PCIEBAR_2 = 64 bits), then reset value = C_AXIBAR2PCIEBAR_2(63 to 32) if (C_AXIBAR2PCIEBAR_2 = 32 bits), then reset value = 0x00000000	To create the address for PCIe—this is the value substituted for the most significant 32 bits of the AXI address.
31-0	Lower Address	R/W	C_AXIBAR2PCIEBAR_3(31 to 0)	To create the address for PCIe—this is the value substituted for the least significant 32 bits of the AXI address.
31-0	Upper Address	R/W	if (C_AXIBAR2PCIEBAR_3 = 64 bits) then reset value = C_AXIBAR2PCIEBAR_3(63 to 32) if (C_AXIBAR2PCIEBAR_3 = 32 bits) then reset value = 0x00000000	To create the address for PCIe—this is the value substituted for the most significant 32 bits of the AXI address.
31-0	Lower Address	R/W	C_AXIBAR2PCIEBAR_4(31 to 0)	To create the address for PCIe—this is the value substituted for the least significant 32 bits of the AXI address.
31-0	Upper Address	R/W	if (C_AXIBAR2PCIEBAR_4 = 64 bits), then reset value = C_AXIBAR2PCIEBAR_4(63 to 32) if (C_AXIBAR2PCIEBAR_4 = 32 bits), then reset value = 0x00000000	To create the address for PCIe—this is the value substituted for the most significant 32 bits of the AXI address.

Table 2-23: AXI Base Address Translation Configuration Register Bit Definitions (Cont'd)

Bits	Name	Core Access	Reset Value	Description
31-0	Lower Address	R/W	C_AXIBAR2PCIEBAR_5(31 to 0)	To create the address for PCIe—this is the value substituted for the least significant 32 bits of the AXI address.
31-0	Upper Address	R/W	if (C_AXIBAR2PCIEBAR_5 = 64 bits), then reset value = C_AXIBAR2PCIEBAR_5(63 to 32) if (C_AXIBAR2PCIEBAR_5 = 32 bits), then reset value = 0x00000000	To create the address for PCIe—this is the value substituted for the most significant 32 bits of the AXI address.

Enhanced Configuration Access

When the AXI Memory Mapped to PCI Express core is configured as a Root Port, configuration traffic is generated by using the PCI Express Enhanced Configuration Access Mechanism (ECAM). ECAM functionality is available only when the core is configured as a Root Port. Reads and writes to a certain memory aperture are translated to configuration reads and writes, as specified in the PCI Express Base Specification (v1.1 and v2.1), §7.2.2 [Ref 8].

Depending on the core configuration, the ECAM memory aperture is 2^{21} – 2^{28} (byte) addresses. The address breakdown is defined in Table 2-24. The ECAM window begins at memory map base address and extends to $2^{(20+ECAM_SIZE)} - 1$. ECAM_SIZE is calculated from the C_BASEADDR and C_HIGHADDR parameters. The number N of low-order bits of the two parameters that do not match, specifies the 2^N byte address range of the ECAM space. If C_INCLUDE_RC = 0, then ECAM_SIZE = 0.

When an ECAM access is attempted to the primary bus number, which defaults as bus 0 from reset, then access to the type 1 PCI™ Configuration Header of the integrated block in the Enhanced Interface for PCIe is performed. When an ECAM access is attempted to the secondary bus number, then type 0 configuration transactions are generated. When an ECAM access is attempted to a bus number that is in the range defined by the secondary bus number and subordinate bus number (not including the secondary bus number), then type 1 configuration transactions are generated. The primary, secondary, and subordinate bus numbers are written by Root Port software to the type 1 PCI Configuration Header of the Enhanced Interface for PCIe in the beginning of the enumeration procedure.

When an ECAM access is attempted to a bus number that is out of the bus_number and subordinate bus number, the bridge does not generate a configuration request and signal SLVERR response on the AXI4-Lite bus. When the AXI Memory Mapped to PCI Express is configured for EP (C_INCLUDE_RC = 0), the underlying Integrated Block configuration space and the core memory map are available at the beginning of the memory space. The memory space looks like a simple PCI Express configuration space. When the AXI Memory Mapped to PCI Express is configured for RC (C_INCLUDE_RC = 1), the same is true, but it also looks like an ECAM access to primary bus, Device 0, Function 0.

When the AXI Memory Mapped to PCI Express core is configured as a Root Port, the reads and writes of the local ECAM are Bus 0. Because the FPGA only has a single Integrated Block for PCIe core, all local ECAM operations to Bus 0 return the ECAM data for Device 0, Function 0.

Configuration write accesses across the PCI Express bus are non-posted writes and block the AXI4-Lite interface while they are in progress. Because of this, system software is not able to service an interrupt if one were to occur. However, interrupts due to abnormal terminations of configuration transactions can generate interrupts. ECAM read transactions block subsequent Requester read TLPs until the configuration read completions packet is returned to allow unique identification of the completion packet.

Table 2-24: ECAM Addressing

Bits	Name	Description
1:0	Byte Address	Ignored for this implementation. The S_AXI_CTL_WSTRB[3:0] signals define byte enables for ECAM accesses.
7:2	Register Number	Register within the configuration space to access.
11:8	Extended Register Number	Along with Register Number, allows access to PCI Express Extended Configuration Space.
14:12	Function Number	Function Number to complete.
19:15	Device Number	Device Number to complete.
(20+n-1):20	Bus Number	Bus Number, $1 \leq n \leq 8$. n is the number of bits available for Bus Number as derived from core parameters C_INCLUDE_RC, C_BASEADDR, and C_HIGHADDR.

Unsupported Memory Space

Advanced Error Reporting (AER) is not supported in the AXI Memory Mapped to PCI Express core. The AER register space is not accessible in the AXI Memory Mapped to PCI Express memory mapped space.

Designing with the Core

This chapter includes guidelines and additional information to make designing with the core easier.

General Design Guidelines

The AMD Vivado™ Design Suite has been optimized to provide a starting point for designing with the AXI Memory Mapped to PCI Express core.

Clocking

Figure 3-1 shows the clocking diagram for the core. The main memory mapped AXI4 bus clock `axi_aclk` is driven by `axi_aclk_out`.

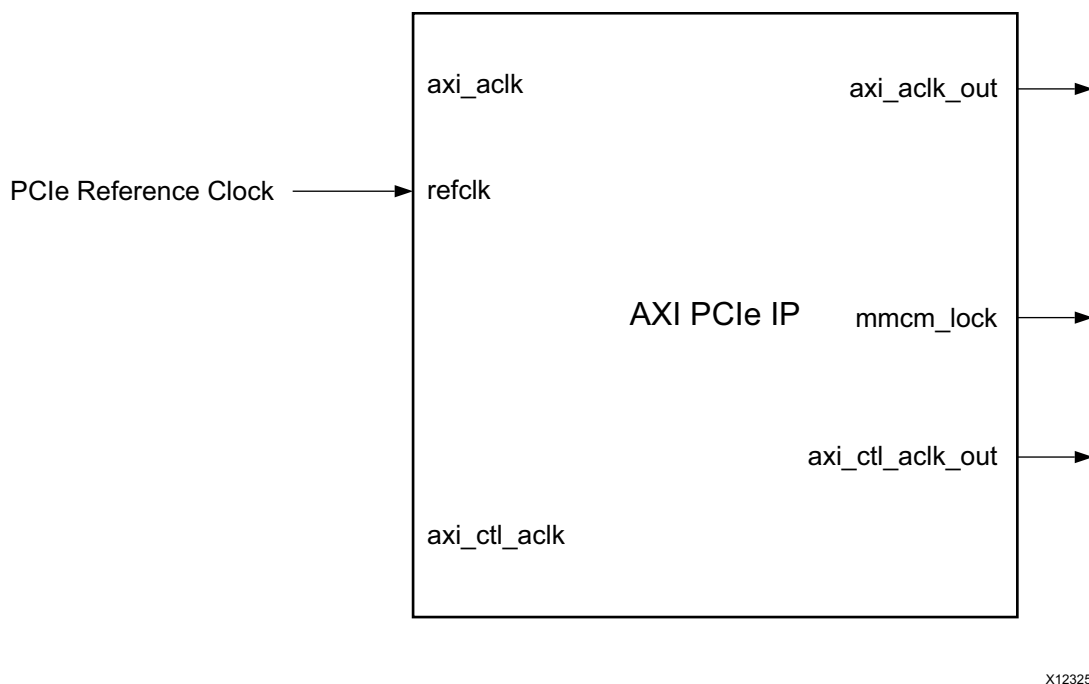


Figure 3-1: Clocking Diagram



IMPORTANT: *axi_aclk_out and axi_ctl_aclk_out are internally connected to axi_aclk and axi_ctl_aclk, respectively, and they do not need to be connected in the design.*

The `refclk` input is used to generate the internal clocks used by the core and the output clocks. This clock must be provided at the reference clock frequency selected in the Vivado Integrated Design Environment (IDE) during IP generation.

The AXI4-Lite interconnect clock `axi_ctl_aclk` is driven by `axi_ctl_aclk_out`. The `axi_ctl_aclk_out` clock is rising edge aligned and an integer division of the `axi_aclk_out` clock.

The output clock frequency is either 62.5 MHz or 125 MHz depending on the PCIe link speed and link width and device selection.

Resets

The AXI Memory Mapped to PCI Express core is designed to be used with the Processing System Reset module for generation of the `axi_areset` input. When using the Vivado IP integrator to build a system, it is best to connect the `perstn` pin of the host connector for PCIe to the `Aux_Reset_In` port of the Processing System Reset module. The bridge does not use `perstn` directly. Also, the `mmcm_lock` output must be connected to the `dcm_locked` input of the Processing System Reset module to make sure that `axi_aresetn` is held active for 16 clocks after `mmcm_lock` becomes active. See [Figure 3-2](#).

Note: Be sure to set the correct polarity on the `aux_reset_in` signal of the `proc_sys_reset` IP block. when `PERSTN` is active-Low, set the parameter as follows:

```
PARAMETER C_AUX_RESET_HIGH = 0
```

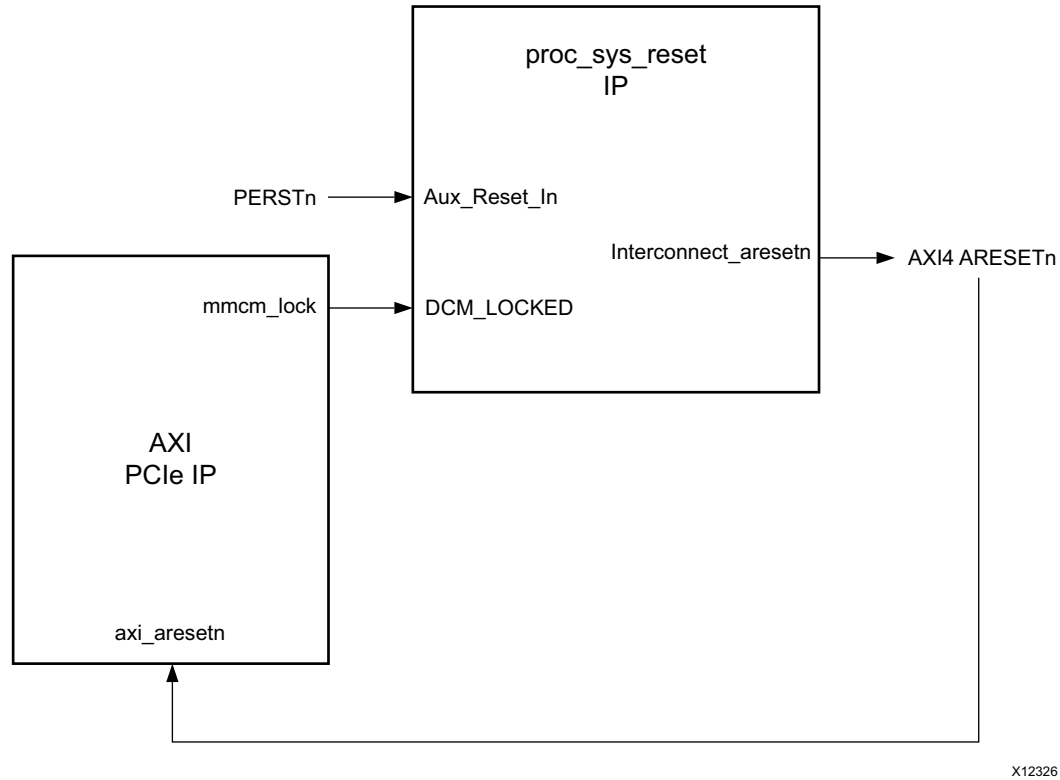


Figure 3-2: System Reset Connection

Shared Logic

This new feature allows you to share common blocks across multiple instantiations. It minimizes the amount of required HDL modifications. You can use your own system-level clocking or reset circuit. You can modify some of these blocks due to system requirements (for example, swapping a BUFG for a BUFGH). You can instantiate multiple cores and share 'common/shared logic' from one core among all instantiated cores. This is applicable for both Endpoint mode and Root Port mode.

In the Vivado Design Suite, the shared logic options are available in the **Shared Logic** page when customizing the core.

There are four types of logic sharing:

- [Shared Clocking](#)
- [Shared GT_COMMON](#)
- [Shared GT_COMMON and Clocking](#)
- [Internal Shared GT_COMMON and Clocking](#)

Shared Clocking

To use the share clocking feature, select **Include Shared Logic (Clocking) in example design** option in the Shared Logic tab (Figure 3-3).

When this feature is selected, the mixed-mode clock manager (MMCM) instance is removed from the pipe wrappers and is moved into the support wrapper of the example design. It also brings out additional ports to the top level to enable sharing of the clocks.

You also have the option to modify and use the unused outputs of the MMCM.

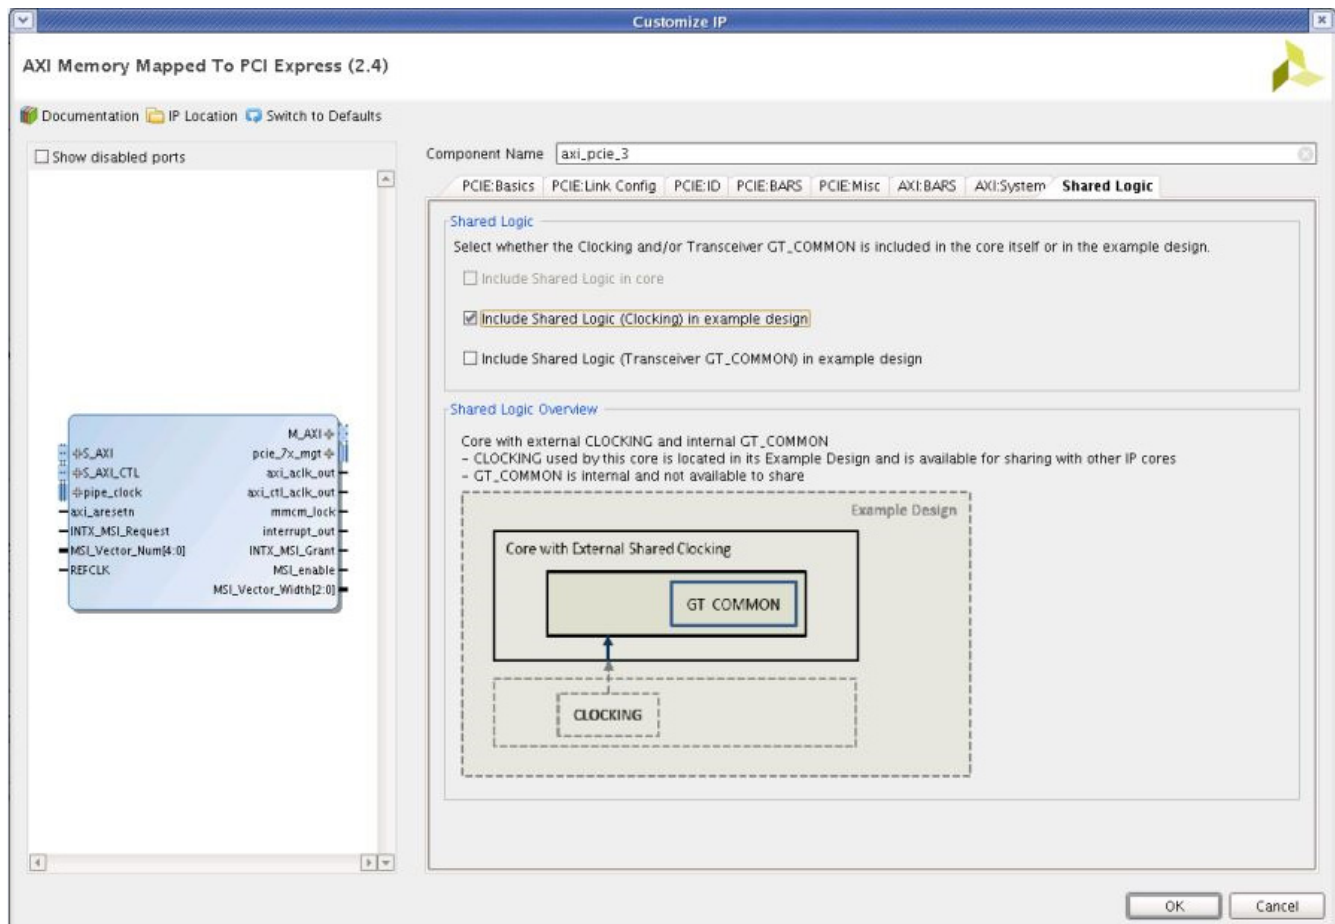


Figure 3-3: Shared Clocking

The MMCM generates the following clocks for PCIe solution wrapper:

- **clk_125mhz**: 125 MHz clock.
- **clk_250mhz**: 250 MHz clock.
- **userclk**: 62.5 MHz/125 MHz/250 MHz clock, depending on selected PCIe core lane width, link speed, and AXI interface width.
- **userclk2**: 250 MHz/500 MHz clock, depending on selected PCIe core link speed.

- **oobclk**: 50 MHz clock.

The other cores/logic present in the user design can use any of the MMCM outputs listed above.

The MMCM instantiated in the PCIe example design has two unconnected outputs: `clkout5`, and `clkout6`. You can use those outputs to generate other desired clock frequencies by selecting the appropriate `CLKOUT5_DIVIDE` and `CLKOUT6_DIVIDE` parameters for MMCM.



TIP: *Sharing the MMCM between PCIe and other cores in your design saves FPGA resources and eases output clock path routing.*

Limitations

- Reference clock input to MMCM is restricted to 100 MHz in most use cases.
 - There is an option for selecting a reference clock of 125 MHz or 250 MHz, which is not a common use case.
- The MMCM reset is tied to a static value in the top module. The MMCM can be reset as required by the system design. Note that MMCM reset can be asserted only after reference clock is recovered and is stable. Also, MMCM reset is indirectly tied to the PCIe core reset and asserting MMCM reset will reset the PCIe core.
- The `userclk1` and `userclk2` outputs are selected based on the **PCIe Lane Width**, **Link Speed**, and **AXI width** selections (for details, see [Customizing the Core in Chapter 4](#)). Sharing cores must comply with these requirements.

Shared GT_COMMON

A quad phase-locked loop (QPLL) in GT_COMMON can serve a quad of GT_CHANNEL instances. If the PCIe core is configured as X1 or X2 and is using a QPLL, the remaining GT_CHANNEL instances can be used by other cores by sharing the same QPLL and GT_COMMON.

To share GT_COMMON instances, select **Include Shared Logic (Transceiver GT_COMMON) in example design** option in the Shared Logic tab ([Figure 3-4](#)).

When this feature is selected, the GT_COMMON instance is removed from the pipe wrappers and is moved into the support wrapper of the example design. It also brings out additional ports to the top level to enable sharing of the GT_COMMON.

Shared logic feature for GT_COMMON helps save FPGA resources and also eases dedicated clock routing within the single GT Quad.

Shared GT_COMMON Use Cases with GTX and GTP

Table 3-1: Shared GT_COMMON Use Cases

GT – PCIe max Link Speed	Device – PCIe Max Link Speed	Shared GT_COMMON
GTX	Kintex7, Virtex7 (485T) – PCIe Gen2	PCIe Pipe Wrappers instantiate GT_COMMON. Shared IP uses QPLL because PCIe uses CPLL inside GT_CHANNEL
GTP	Artix7 – PCIe Gen2	GTP_COMMON has 2 QPLLs. PCIe Pipe Wrappers use only one QPLL. The remaining one can be used by shared IP core.

Limitations

- The reset logic in the pipe wrapper resets the QPLL when the PCIe Block performs a rate change. When sharing is enabled, the core/logic which is sharing the QPLL must be able to handle and recover from this reset.
- The settings of the GT_COMMON should not be changed as they are optimized for the PCIe core.

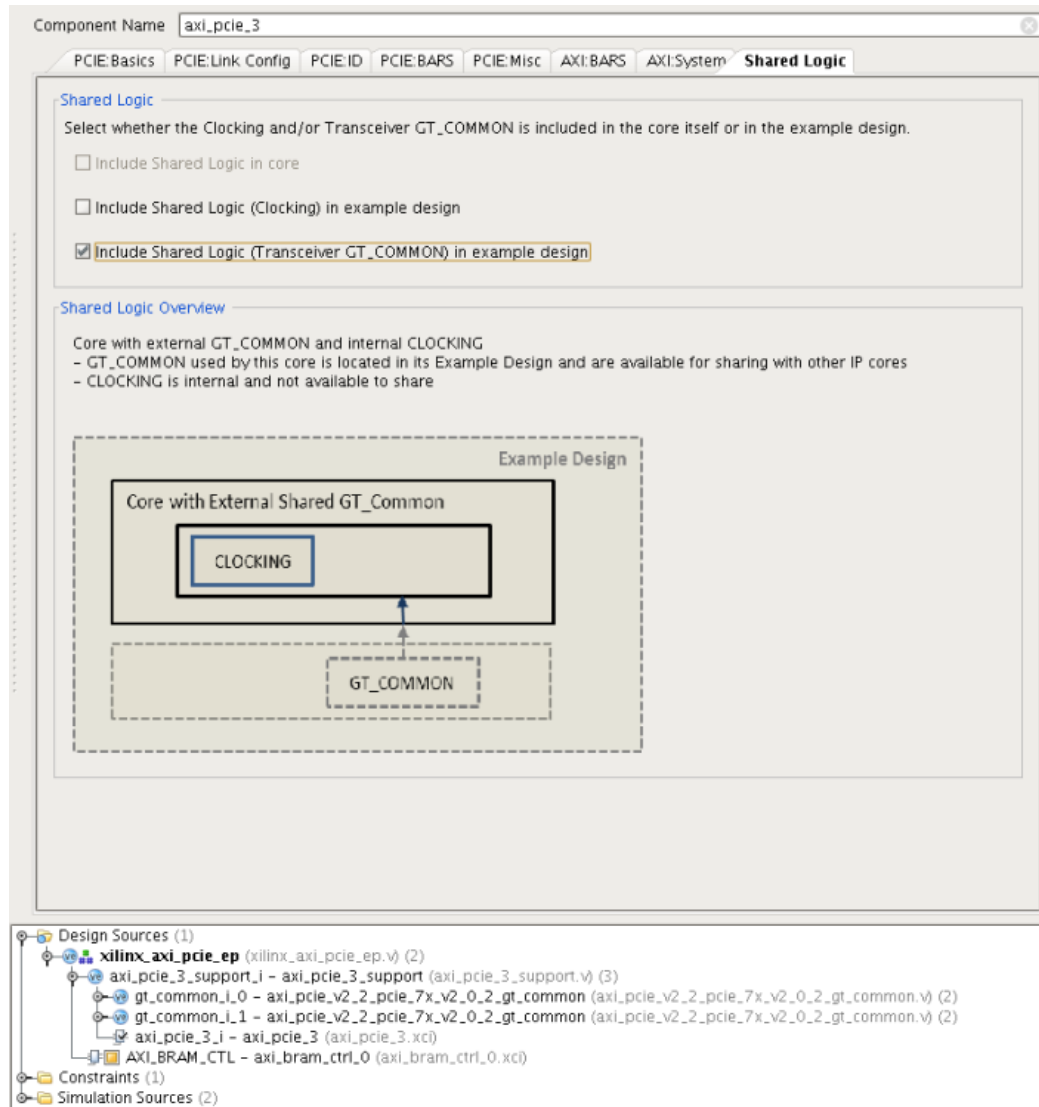


Figure 3-4: Shared GT_COMMON

Shared GT_COMMON and Clocking

You can share both GT_COMMON and Clocking instances when you select **Include Shared Logic (Clocking) in example design** and **Include Shared Logic (Transceiver GT_COMMON) in example design** in the Shared Logic page (see [Figure 3-5](#)).

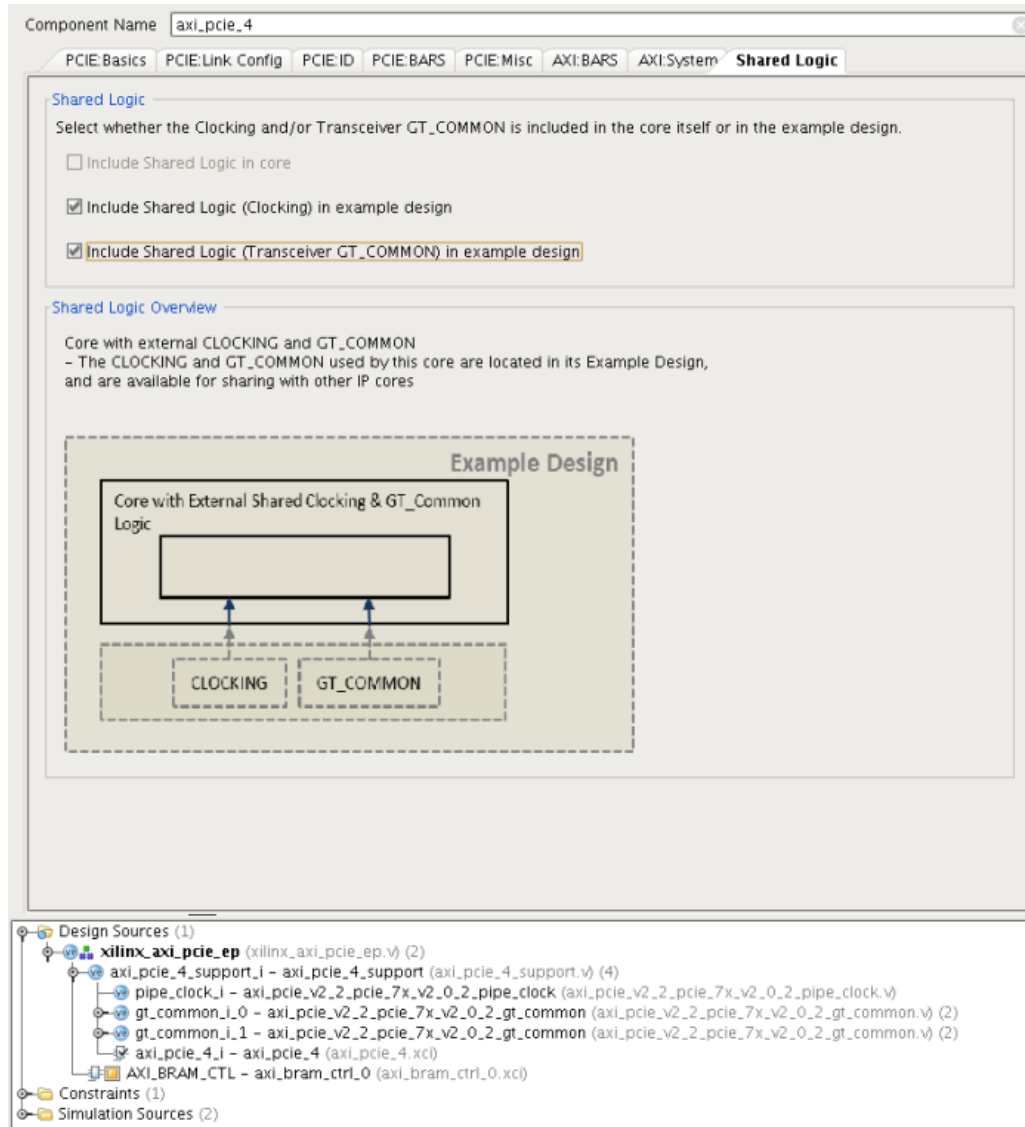


Figure 3-5: Shared GT_COMMON and Clocking

Internal Shared GT_COMMON and Clocking

This feature allows sharing of GT_COMMON and Clocks while these modules are still internal to the core (not brought up to the support wrapper). It can be enabled when you select **Include Shared Logic in Core** in the Shared Logic page (see [Figure 3-6](#)).

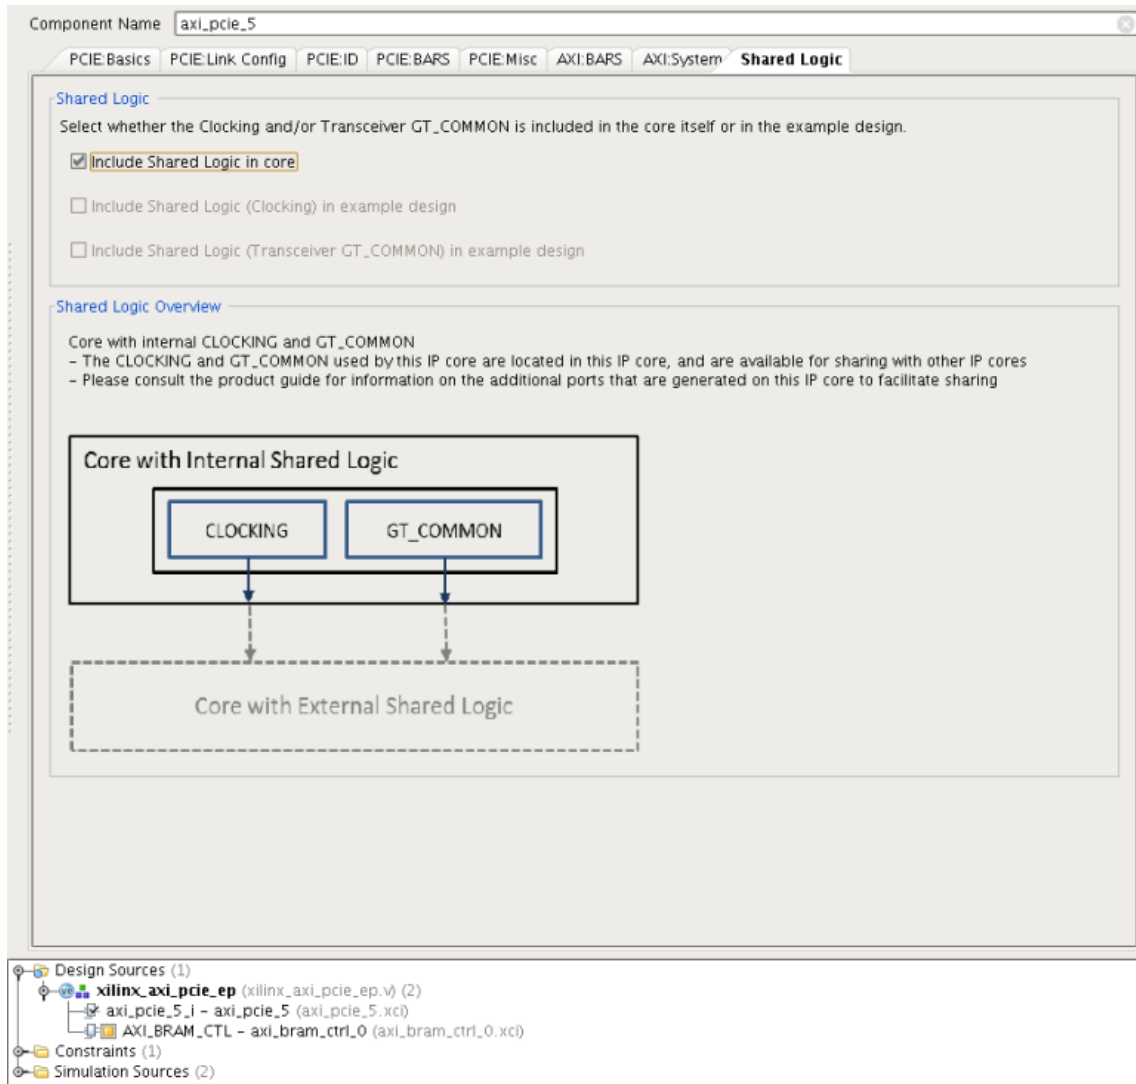


Figure 3-6: Internal Shared Logic

Clocking Interface

Table 3-2 defines the clocking interface signals.

Table 3-2: Clocking Interface Signals

Name	Direction	Description
clk_pclk	Input	Parallel clock used to synchronize data transfers across the parallel interface of the transceiver.
clk_rxusrclk	Input	Clock for the internal RX PCS datapath of the transceiver.
clk_fab_refclk	Input	Reserved.

Table 3-2: Clocking Interface Signals (Cont'd)

Name	Direction	Description
clk_dclk	Input	GT DRP clock input. This clock is required for proper core operation even when GT DRP interface is not used by the user.
clk_userclk1 ⁽¹⁾	Input	Clock used internally for PCIe block logic.
clk_userclk2 ⁽¹⁾	Input	User clock used for the majority portion of the AXI-PCIe Bridge, AXI4 (axi_aclk_out), and AXI4-Lite (axi_ctl_aclk_out) interfaces.
clk_oobclk_in	Input	Free running clock used for OOB detect circuitry in the Transceivers
clk_mmcm_lock	Input	Indicates if the MMCM is locked onto the source clk. This signal is used only when external shared clocking logic is enabled.
clk_txoutclk	Output	Recommended clock output to the FPGA logic. Frequency of this clock will be equal to the sys_clk frequency.
clk_rxoutclk	Output	Reserved.
clk_pclk_sel	Output	Parallel clock select. This signal toggles when PCIe link up-trains to Gen2 speed.
clk_gen3	Output	Reserved.
pipe_mmcm_rst_n	Input	MMCM reset port. This port could be used by the upper layer to reset MMCM if error recovery is required. If the system detects the deassertion of MMCM lock, AMD recommends that you reset the MMCM. The recommended approach is to reset the MMCM after the MMCM input clock recovers (if MMCM reset occurs before the input reference clock recovers, the MMCM might never relock). After MMCM is reset, wait for MMCM to lock and then reset the PIPE Wrapper as normally done. Currently this port is tied High.

Notes:

1. clk_userclk1 and clk_userclk2 are the user clocks for PCIe and AXI related logic and are different than TX/RXUSRCLK and TX/RXUSRCLK2 referred in 7-series FPGAs GTX/GTH Transceivers User Guide (UG476) [Ref 3]. TX/RXUSRCLK and TX/RXUSRCLK2 referred there are 'user clocks' from the Transceivers perspective which are clk_pclk (Parallel clock).

The Clocking architecture is described in detail in the Use Model chapter of the *7 Series FPGAs GTX/GTH Transceivers User Guide (UG476)* [Ref 3].

AXI Transactions for PCIe

Table 3-3 and Table 3-4 are the translation tables for AXI4-Stream and memory-mapped transactions.

Table 3-3: AXI4 Memory-Mapped Transactions to AXI4-Stream PCIe TLPs

AXI4 Memory-Mapped Transaction	AXI4-Stream PCIe TLPs
INCR Burst Read of 32-bit address AXIBAR	MemRd 32 (3DW)
INCR Burst Write to 32-bit address AXIBAR	MemWr 32 (3DW)
INCR Burst Read of 32-bit address AXIBAR	MemRd 64 (4DW)
INCR Burst Write to 32-bit address AXIBAR	MemWr 64 (4DW)

Table 3-4: AXI4-Stream PCIe TLPs to AXI4 Memory Mapped Transactions

AXI4-Stream PCIe TLPs	AXI4 Memory-Mapped Transaction
MemRd 32 (3DW) of PCIEBAR	INCR Burst Read with 32-bit address
MemWr 32 (3DW) to PCIEBAR	INCR Burst Write with 32-bit address
MemRd 64 (4DW) of PCIEBAR	INCR Burst Read with 32-bit address
MemWr 64 (4DW) to PCIEBAR	INCR Burst Write with 32-bit address

Note: `s_axi_wstrb` can be used to facilitate data alignment to an address boundary. `s_axi_wstrb` may equal to 0 in the beginning of valid data cycle and will appropriately calculate an offset to the given address. However, the valid data identified by `s_axi_wstrb` must be continuous from the first byte enable to the last byte enable.

Transaction Ordering for PCIe

The AXI Memory Mapped to PCI Express core conforms to strict PCIe transaction ordering rules. See the PCIe v2.1 Specification [Ref 8] for the complete rule set. The following behaviors are implemented in the AXI Memory Mapped to PCI Express core to enforce the PCIe transaction ordering rules on the highly-parallel AXI bus of the bridge. The rules are enforced without regard to the Relaxed Ordering attribute bit within the TLP header:

- The `bresp` to the remote (requesting) AXI4 master device for a write to a remote PCIe device is not issued until the MemWr TLP transmission is guaranteed to be sent on the PCIe link before any subsequent TX-transfers.
- A remote AXI master read of a remote PCIe device is not permitted to pass any previous or simultaneous AXI master writes to a remote PCIe device that occurs previously or at the same time. Timing is based off the AXI `arvalid` signal timing relative to the AXI `awvalid`. Any AXI write transaction in which `awvalid` was asserted before or at the same time as the `arvalid` for a read from PCIe is asserted causes the MemRd TLP(s) to be held until the pipelined or simultaneous MemWr TLP(s) have been sent.
- A remote PCIe device read of a remote AXI slave is not permitted to pass any previous remote PCIe device writes to a remote AXI slave received by the AXI Memory Mapped to PCI Express core. The AXI read address phase is held until the previous AXI write transactions have completed and `bresp` has been received for the AXI write transactions.

- Read completion data received from a remote PCIe device are not permitted to pass any remote PCIe device writes to a remote AXI slave received by the AXI Memory Mapped to PCI Express core prior to the read completion data. The `bresp` for the AXI write(s) must be received before the completion data is presented on the AXI read data channel.
- Read data from a remote AXI slave is not permitted to pass any remote AXI master writes to a remote PCIe device initiated on the AXI bus prior to or simultaneously with the read data being returned on the AXI bus. Timing is based on the AXI `awvalid` signal timing relative to the AXI `rvalid` assertion. Any AXI write transaction in which `awvalid` was asserted before or simultaneously with the `rvalid` being asserted up to and including the last data beat, causes the Completion TLP(s) to be held until the pipelined or simultaneous MemWr TLP(s) have been sent.



IMPORTANT: *The transaction ordering rules for PCIe might have an impact on data throughput in heavy bidirectional traffic.*

BAR and Address Translation

BAR Addressing

`C_AXIBAR_n` and `C_AXIBAR_HIGHADDR_n` are used to calculate the size of the AXI BAR *n* and during address translation to PCIe address.

- `C_AXIBAR_n` provides the low address where AXI BAR *n* starts and will be regarded as address offset 0x0 when the address is translated.
- `C_AXIBAR_HIGHADDR_n` is the high address of the last valid byte address of AXI BAR *n*. (For more details on how the address gets translated, see [Address Translation](#).)

The difference between the two parameters are your AXI BAR *n* size. These parameters must be set accordingly such that AXI BAR *n* size is a power of two and must have at least 8 KB and a maximum of 2 GB.

When a packet is sent to the core (outgoing PCIe packets), the packet must have an address that is in the range of `C_AXIBAR_n` and `C_AXIBAR_HIGHADDR_n`. Any packet that is received by the core that has an address outside of this range will be responded to with a SLVERR (Slave Error). When IP Integrator is used, these parameters are derived from the Address Editor tab within IP Integrator. The Address Editor sets the AXI Interconnect as well as the core so the address range matches, and the packet is routed to the core only when the packet has an address within the valid range.

Address Translation

The address space for PCIe is different than AXI address space. To access one address space from another address space requires an address translation process. On the AXI side, the bridge supports mapping to PCIe on up to six 32-bit or 64-bit AXI base address registers (BARs). The generics used to configure the BARs follow.

`C_AXIBAR_NUM`, `C_AXIBAR_n`, `C_AXIBAR_HIGHADDR_n`, `C_AXIBAR2PCIEBAR_n` and `C_AXIBAR_AS_n`

where n represents an AXIBAR number from 0 to 5. The bridge for PCIe supports mapping on up to three 64-bit BARs for PCIe. The generics used to configure the BARs are:

`C_PCIEBAR_NUM`, `C_PCIEBAR2AXIBAR_n` and `C_PCIEBAR_LEN_n`

where n represents a particular BAR number for PCIe from 0 to 2.

Note: The `C_INCLUDE_BAROFFSET_REG` generic allows for dynamic address translation. When this parameter is set to one, the `AXIBAR2PCIEBAR_n` translation vectors can be changed by using software.

Four examples follow:

- [Example 1 \(32-bit PCIe Address Mapping\)](#) demonstrates how to set up three 32-bit AXI BARs and translate the AXI address to a 32-bit address for PCIe.
- [Example 2 \(64-bit PCIe Address Mapping\)](#) demonstrates how to set up three 64-bit AXI BARs and translate the AXI address to a 64-bit address for PCIe.
- [Example 3](#) demonstrates how to set up two 64-bit PCIe BARs and translate the address for PCIe to an AXI address.
- [Example 4](#) demonstrates how to set up a combination of two 32-bit AXI BARs and two 64 bit AXI BARs, and translate the AXI address to an address for PCIe.

Example 1 (32-bit PCIe Address Mapping)

This example shows the generic settings to set up three independent 32-bit AXI BARs and address translation of AXI addresses to a remote 32-bit address space for PCIe. This setting of AXI BARs does not depend on the BARs for PCIe within the AXI Memory Mapped to PCI Express core.

In this example, where `C_AXIBAR_NUM=3`, the following assignments for each range are made:

```
C_AXIBAR_AS_0=0
C_AXIBAR_0=0x12340000
C_AXI_HIGHADDR_0=0x1234FFFF
C_AXIBAR2PCIEBAR_0=0x5671XXXX (Bits 15-0 do not matter as the lower 16-bits hold the
actual lower 16-bits of the PCIe address)
```

```

C_AXIBAR_AS_1=0
C_AXIBAR_1=0xABCDE000
C_AXI_HIGHADDR_1=0xABCDFFFF
C_AXIBAR2PCIEBAR_1=0xFEDC0XXX (Bits 12-0 do not matter as the lower 13-bits hold the
actual lower 13-bits of the PCIe address)

C_AXIBAR_AS_2=0
C_AXIBAR_2=0xFE000000
C_AXI_HIGHADDR_2=0xFFFFFFFF
C_AXIBAR2PCIEBAR_2=0x40XXXXXX (Bits 24-0 do not care)

```

- Accessing the Bridge AXIBAR_0 with address 0x12340ABC on the AXI bus yields 0x56710ABC on the bus for PCIe.

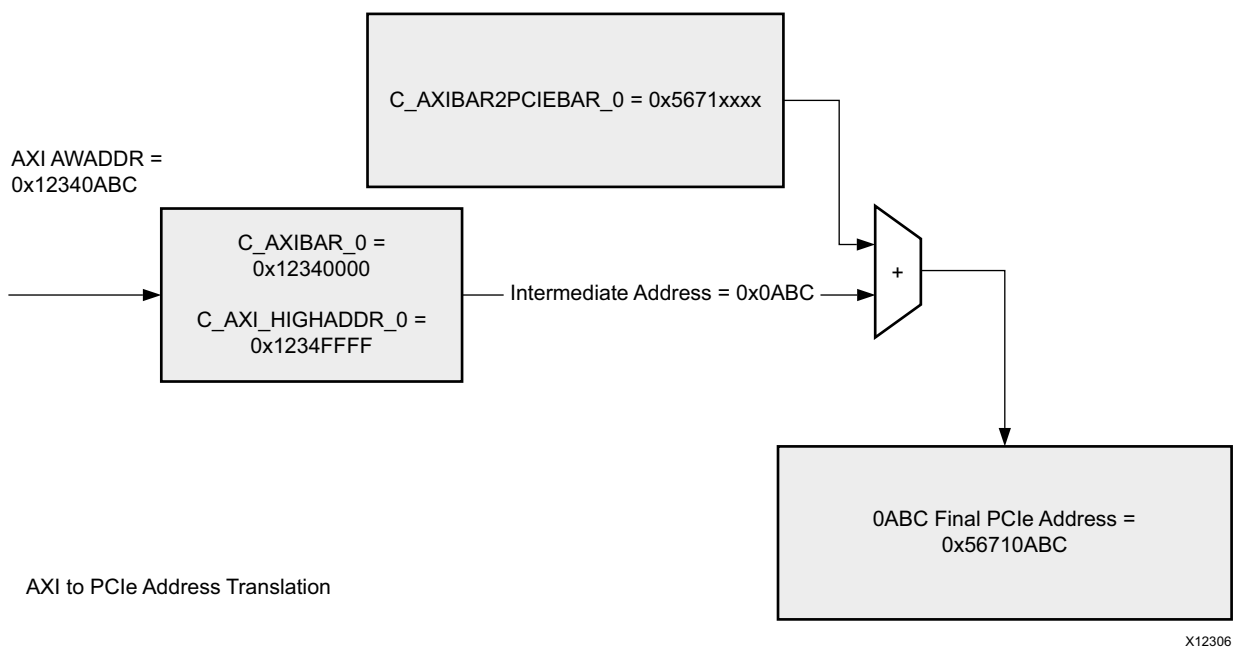


Figure 3-7: AXI to PCIe Address Translation

- Accessing the Bridge AXIBAR_1 with address 0xABCD123 on the AXI bus yields 0xFEDC1123 on the bus for PCIe.
- Accessing the Bridge AXIBAR_2 with address 0xFFFFEDCBA on the AXI bus yields 0x41FEDCBA on the bus for PCIe.

Example 2 (64-bit PCIe Address Mapping)

This example shows the generic settings to set up to three independent 64-bit AXI BARs and address translation of AXI addresses to a remote 64-bit address space for PCIe. This setting of AXI BARs does not depend on the BARs for PCIe within the Bridge.

In this example, where C_AXIBAR_NUM=3, the following assignments for each range are made:

```
C_AXIBAR_AS_0=1
C_AXIBAR_0=0x12340000
C_AXI_HIGHADDR_0=0x1234FFFF
C_AXIBAR2PCIEBAR_0=0x500000005671XXXX (Bits 15-0 do not matter)

C_AXIBAR_AS_1=1
C_AXIBAR_1=0xABCDE000
C_AXI_HIGHADDR_1=0xABCDFFFF
C_AXIBAR2PCIEBAR_1=0x60000000FEDC0XXX (Bits 12-0 do not matter)

C_AXIBAR_AS_2=1
C_AXIBAR_2=0xFE000000
C_AXI_HIGHADDR_2=0xFFFFFFFF
C_AXIBAR2PCIEBAR_2=0x7000000040XXXXXX (Bits 24-0 do not matter)
```

- Accessing the Bridge AXIBAR_0 with address 0x12340ABC on the AXI bus yields 0x5000000056710ABC on the bus for PCIe.
- Accessing the Bridge AXIBAR_1 with address 0xABCDF123 on the AXI bus yields 0x60000000FEDC1123 on the bus for PCIe.
- Accessing the Bridge AXIBAR_2 with address 0xFFFE DCBA on the AXI bus yields 0x7000000041FEDCBA on the bus for PCIe.

Example 3

This example shows the generic settings to set up two independent BARs for PCIe and address translation of addresses for PCIe to a remote AXI address space. This setting of BARs for PCIe does not depend on the AXI BARs within the bridge.

In this example, where C_PCIEBAR_NUM=2, the following range assignments are made:

```
BAR 0 is set to 0x20000000_ABCD8000 by the Root Port
C_PCIEBAR_LEN_0=15
C_PCIEBAR2AXIBAR_0=0x1234_0XXX (Bits 14-0 do not matter)

BAR 1 is set to 0xA000000012000000 by Root Port
C_PCIEBAR_LEN_1=25
C_PCIEBAR2AXIBAR_1=0xFEXXXXXX (Bits 24-0 do not matter)
```

- Accessing the Bridge PCIEBAR_0 with address 0x20000000_ABCDFFF4 on the bus for PCIe yields 0x1234_7FF4 on the AXI bus.

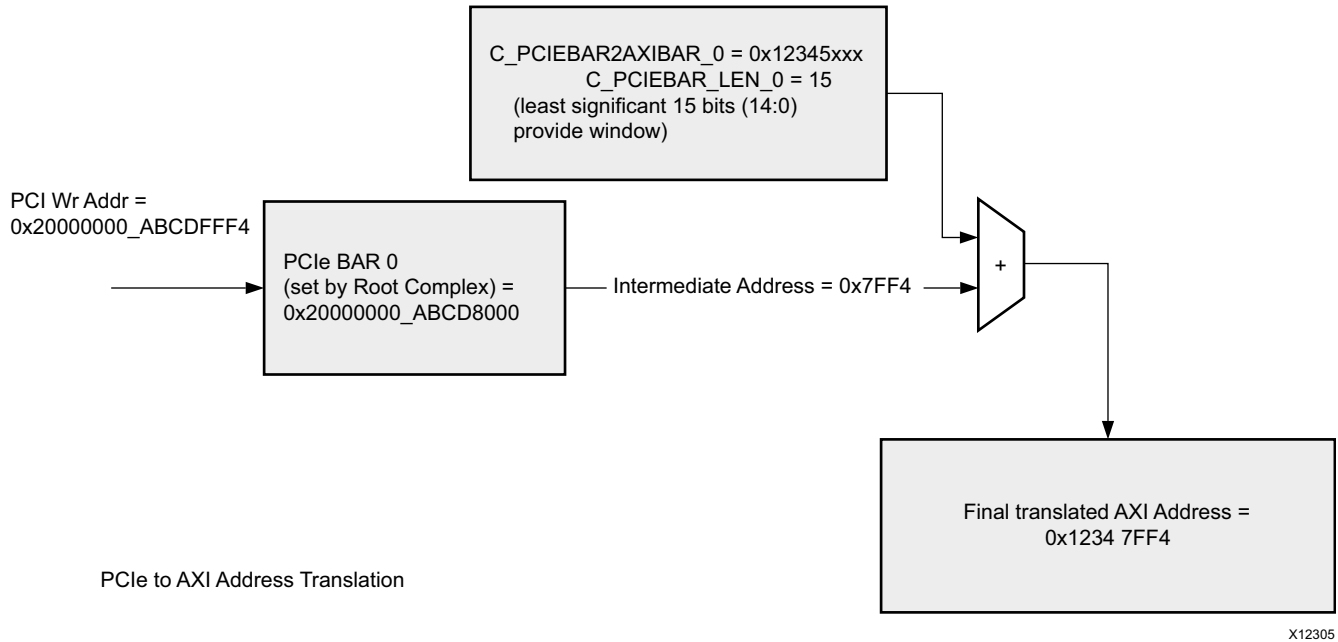


Figure 3-8: PCIe to AXI Translation

- Accessing Bridge PCIEBAR_1 with address 0xA00000001235FEDC on the bus for PCIe yields 0xFE35FEDC on the AXI bus.

Example 4

This example shows the generic settings to set up a combination of two independent 32-bit AXI BARs and two independent 64-bit BARs and address translation of AXI addresses to a remote address space for PCIe. This setting of AXI BARs do not depend on the BARs for PCIe within the Bridge.

In this example, where C_AXIBAR_NUM=4, the following assignments for each range are made:

```
C_AXIBAR_AS_0=0
C_AXIBAR_0=0x12340000
C_AXI_HIGHADDR_0=0x1234FFFF
C_AXIBAR2PCIEBAR_0=0x5671XXXX (Bits 15-0 do not matter)

C_AXIBAR_AS_1=1
C_AXIBAR_1=0xABCDE000
C_AXI_HIGHADDR_1=0xABCDDFFF
C_AXIBAR2PCIEBAR_1=0x50000000FEDC0XXX (Bits 12-0 do not matter)

C_AXIBAR_AS_2=0
C_AXIBAR_2=0xFE000000
C_AXI_HIGHADDR_2=0xFFFFFFFF
C_AXIBAR2PCIEBAR_2=0x40XXXXXX (Bits 24-0 do not matter)

C_AXIBAR_AS_3=1
C_AXIBAR_3=0x00000000
```

```
C_AXI_HIGHADDR_3=0x0000007F
C_AXIBAR2PCIEBAR_3=0x600000008765438X (Bits 6-0 do not matter)
```

- Accessing the Bridge AXIBAR_0 with address 0x12340ABC on the AXI bus yields 0x56710ABC on the bus for PCIe.
- Accessing the Bridge AXIBAR_1 with address 0xABCD123 on the AXI bus yields 0x50000000FEDC1123 on the bus for PCIe.
- Accessing the Bridge AXIBAR_2 with address 0xFFEDCBA on the AXI bus yields 0x41FEDCBA on the bus for PCIe.
- Accessing the AXI M S PCIe Bridge AXIBAR_3 with address 0x00000071 on the AXI bus yields 0x60000000876543F1 on the bus for PCIe.

Addressing Checks

When setting the following parameters for PCIe address mapping, C_PCIEBAR2AXIBAR_n and C_PCIEBAR_LEN_n, be sure these are set to allow for the 32-bit addressing space on the AXI system. For example, the following setting is illegal and results in an invalid AXI address.

```
C_PCIEBAR2AXIBAR_0 = 0xFFFF_0000
C_PCIEBAR_LEN_0 = 23
```

Also, check for a larger value on C_PCIEBAR_LEN_n compared to the value assigned to parameter, C_PCIEBAR2AXIBAR_n. For example, the following parameter settings.

```
C_PCIEBAR2AXIBAR_0 = 0xFFFF_E000
C_PCIEBAR_LEN_0 = 20
```

To keep the AXIBAR upper address bits as 0xFFFF_E000 (to reference bits [31:13]), the C_PCIEBAR_LEN_0 parameter must be set to 13.

Interrupts

This section describes the interrupt pins which include Local, MSI and Legacy Interrupts.

Local Interrupts

The `interrupt_out` pin can be configured to send interrupts based on the settings of the Interrupt Mask register. The `interrupt_out` pin signals interrupts to devices attached to the memory mapped AXI4 side of the bridge. The MSI interrupt defined in the Interrupt Mask & Interrupt Decode registers is used to indicate the receipt of a Message Signaled Interrupt only when the bridge is operating in Root Port mode (C_INCLUDE_RC=1).

MSI Interrupt

When the `msi_enable` output pin indicates that the bridge has Endpoint MSI functionality enabled (`msi_enable = 1`), the `intx_msi_request` input pin is defined as MSI_Request and can be used to trigger a Message Signaled Interrupt through a special MemWr TLP to an external Root Port for PCIe on the PCIe side of the Bridge. The `intx_msi_request` input pin is positive-edge detected and synchronous to `axi_aclk_out`. The address and data contained in this MemWr TLP are determined by an external Root Port for PCIe configuration of registers within the integrated block for PCI Express. The `intx_msi_request` pin input is valid only when the bridge is operating in Endpoint mode (`C_INCLUDE_RC=0`).

Additional MSI capability now supports multiple vectors on the Endpoint configuration of the AXI Memory Mapped to PCI Express core. Using the handshaking described here, an additional input value specifies the vector number to send with the MSI MemWr TLP upstream to the Root Port. This is specified on the input signal, `msi_vector_num`. This signal is (4:0), and represents up to (32), the allowable MSI messages that can be sent from the Endpoint (and what is enabled after configuration).

The bridge ignores any bits set on the `msi_vector_num` input signal, if they are not allocated in the Message Control Register.

The Endpoint requests the number of message as specified in the design parameter (of the AXI Memory Mapped to PCI Express), `C_NUM_MSI_REQ`. Following specification requirements, this parameter can be set up to 5. `C_NUM_MSI_REQ` represents the number of MSI vectors requested. For example, `C_NUM_MSI_REQ = 5` represents a request of $2^5 = 32$ MSI vectors. This parameter value, `C_NUM_MSI_REQ` is assigned to the Message Control Register field, Multiple Message Capable, bits (3:1).

After configuration, the number of allocated MSI vectors is specified in the design output port, `msi_vector_width`. This signal with width, (2:0), can only be values up to 5 (101), representing 32 allocated MSI vectors for the Endpoint. Output values of 6 and 7; 110 and 111 are reserved. The `msi_vector_width` output signal is a direct correlation from the value in the Multiple Message Enable field bits (6:4) of the Message Control Register as shown in [Table 3-5](#).

Table 3-5: MSI Vectors Enabled in Message Control Register

Value	Number of Messages Requested	Output Signal, MSI_Vector_Width (2:0)
000	1	000
001	2	001
010	4	010
011	8	011
100	16	100
101	32	101

Table 3-5: MSI Vectors Enabled in Message Control Register (Cont'd)

Value	Number of Messages Requested	Output Signal, MSI_Vector_Width (2:0)
110	Reserved	N/A
111	Reserved	N/A

Additional IP is required in the Endpoint PCIe system to create the priority scheme for the MSI vectors on the PCIe interface.

Legacy Interrupts

The bridge supports legacy interrupts for PCI if selected by the `C_INTERRUPT_PIN` parameter. (Can only be set to 1 when `C_INCLUDE_RC = 0`.) A value of 1 selects INTA, as defined in Table 2-2. If a legacy interrupt for PCI support is selected and the `msi_enable` output pin indicates that the bridge has endpoint MSI functionality disabled (`MSI_enable = 0`), the `intx_msi_request` pin is defined as INTX. When the INTX pin goes High, an assert INTA message is sent. When the INTX pin goes Low, a deassert INTA message is sent. The interrupts needs to be synchronized to `axi_aclk_out`. These messages are defined in the PCI 2.1 specification. The `intx_msi_request` pin input is valid only when the bridge is operating in Endpoint mode (`C_INCLUDE_RC=0`).

Malformed TLP

The integrated block for PCI Express detects a malformed TLP. For the IP configured as an Endpoint core, a malformed TLP results in a fatal error message being sent upstream if error reporting is enabled in the Device Control Register.

For the IP configured as a Root Port, when a malformed TLP is received from the Endpoint, this can fall under one of several types of violations as per the PCIe specification. For example, if a Received TLP has the Error Poison bit set, this is discarded by the MM/S master bridge, and the MEP (Master Error Poison) bit is set in the Interrupt Decode register.

Abnormal Conditions

This section describes how the Slave side (Table 3-6) and Master side (Table 3-7) of the AXI Memory Mapped to PCI Express core handle abnormal conditions.

Slave Bridge Abnormal Conditions

Slave bridge abnormal conditions are classified as: Illegal Burst Type and Completion TLP Errors. The following sections describe the manner in which the Bridge handles these errors.

Illegal Burst Type

The slave bridge monitors AXI read and write burst type inputs to ensure that only the INCR (incrementing burst) type is requested. Any other value on these inputs is treated as an error condition and the Slave Illegal Burst (SIB) interrupt is asserted. In the case of a read request, the Bridge asserts SLVERR for all data beats and arbitrary data is placed on the `s_axi_rdata` bus. In the case of a write request, the Bridge asserts SLVERR for the write response and all write data is discarded.

Completion TLP Errors

Any request to the bus for PCIe (except for posted Memory write) requires a completion TLP to complete the associated AXI request. The Slave side of the Bridge checks the received completion TLPs for errors and checks for completion TLPs that are never returned (Completion Timeout). Each of the completion TLP error types are discussed in the subsequent sections.

Unexpected Completion

When the slave bridge receives a completion TLP, it matches the header RequesterID and Tag to the outstanding RequesterID and Tag. A match failure indicates the TLP is an Unexpected Completion which results in the completion TLP being discarded and a Slave Unexpected Completion (SUC) interrupt strobe being asserted. Normal operation then continues.

Unsupported Request

A device for PCIe might not be capable of satisfying a specific read request. For example, the read request targets an unsupported address for PCIe causing the completer to return a completion TLP with a completion status of 0b001 - Unsupported Request. The completer can also return a completion TLP with a completion status that is reserved according to the 2.1 PCIe Specification, which must be treated as an unsupported request status. When the slave bridge receives an unsupported request response, the Slave Unsupported Request (SUR) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Completion Timeout

A Completion Timeout occurs when a completion (Cpl) or completion with data (CpID) TLP is not returned after an AXI to PCIe read request. Completions must complete within the `C_COMP_TIMEOUT` parameter selected value from the time the MemRd for PCIe request is issued. When a completion timeout occurs, a Slave Completion Timeout (SCT) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Poison Bit Received on Completion Packet

An Error Poison occurs when the completion TLP EP bit is set, indicating that there is poisoned data in the payload. When the slave bridge detects the poisoned packet, the Slave Error Poison (SEP) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Completer Abort

A Completer Abort occurs when the completion TLP completion status is 0b100 - Completer Abort. This indicates that the completer has encountered a state in which it was unable to complete the transaction. When the slave bridge receives a completer abort response, the Slave Completer Abort (SCA) interrupt is asserted and the SLVERR response is asserted with arbitrary data on the memory mapped AXI4 bus.

Table 3-6: Slave Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
Read	Illegal burst type	SIB interrupt is asserted. SLVERR response given with arbitrary read data.
Write	Illegal burst type	SIB interrupt is asserted. Write data is discarded. SLVERR response given.
Read	Unexpected completion	SUC interrupt is asserted. Completion is discarded.
Read	Unsupported Request status returned	SUR interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Completion timeout	SCT interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Poison bit in completion	Completion data is discarded. SEP interrupt is asserted. SLVERR response given with arbitrary read data.
Read	Completer Abort (CA) status returned	SCA interrupt is asserted. SLVERR response given with arbitrary read data.

Master Bridge Abnormal Conditions

The following sections describe the manner in which the master bridge handles abnormal conditions.

AXI DECERR Response

When the master bridge receives a DECERR response from the AXI bus, the request is discarded and the Master DECERR (MDE) interrupt is asserted. If the request was

non-posted, a completion packet with the Completion Status = Unsupported Request (UR) is returned on the bus for PCIe.

AXI SLVERR Response

When the master bridge receives a SLVERR response from the addressed AXI slave, the request is discarded and the Master SLVERR (MSE) interrupt is asserted. If the request was non-posted, a completion packet with the Completion Status = Completer Abort (CA) is returned on the bus for PCIe.

Max Payload Size for PCIe, Max Read Request Size or 4K Page Violated

It is the responsibility of the requester to ensure that the outbound request adhere to the Max Payload Size, Max Read Request Size, and 4 Kb Page Violation rules. If the master bridge receives a request that violates one of these rules, the bridge processes the invalid request as a valid request, which can return a completion that violates one of these conditions or can result in the loss of data. The master bridge does not return a malformed TLP completion to signal this violation.

Completion Packets

When the MAX_READ_REQUEST_SIZE is greater than the MAX_PAYLOAD_SIZE, a read request for PCIe can ask for more data than the master bridge can insert into a single completion packet. When this situation occurs, multiple completion packets are generated up to MAX_PAYLOAD_SIZE, with the Read Completion Boundary (RCB) observed.

Poison Bit

When the poison bit is set in a transaction layer packet (TLP) header, the payload following the header is corrupted. When the master bridge receives a memory request TLP with the poison bit set, it discards the TLP and asserts the Master Error Poison (MEP) interrupt strobe.

Zero Length Requests

When the master bridge receives a read request with the Length = 0x1, FirstBE = 0x00, and LastBE = 0x00, it responds by sending a completion with Status = Successful Completion. When the master bridge receives a write request with the Length = 0x1, FirstBE = 0x00, and LastBE = 0x00 there is no effect.

Table 3-7: Master Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
Read	DECERR response	MDE interrupt strobe asserted. Completion returned with Unsupported Request status.
Write	DECERR response	MDE interrupt strobe asserted.

Table 3-7: Master Bridge Response to Abnormal Conditions

Transfer Type	Abnormal Condition	Bridge Response
Read	SLVERR response	MSE interrupt strobe asserted. Completion returned with Completer Abort status.
Write	SLVERR response	MSE interrupt strobe asserted.
Write	Poison bit set in request	MEP interrupt strobe asserted. Data is discarded.
Read	DECERR response	MDE interrupt strobe asserted. Completion returned with Unsupported Request status.
Write	DECERR response	MDE interrupt strobe asserted.

Link Down Behavior

The normal operation of the AXI Memory Mapped to PCI Express core is dependent on the integrated block for PCIe establishing and maintaining the point-to-point link with an external device for PCIe. If the link has been lost, it must be re-established to return to normal operation.

When a Hot Reset is received by the AXI Memory Mapped to PCI Express core, the link goes down and the PCI Configuration Space must be reconfigured.

Initiated AXI4 write transactions that have not yet completed on the AXI4 bus when the link goes down have a SLVERR response given and the write data is discarded. Initiated AXI4 read transactions that have not yet completed on the AXI4 bus when the link goes down have a SLVERR response given, with arbitrary read data returned.

Any MemWr TLPs for PCIe that have been received, but the associated AXI4 write transaction has not started when the link goes down, are discarded. If the associated AXI4 write transaction is in the process of being transferred, it completes as normal. Any MemRd TLPs for PCIe that have been received, but have not returned completion TLPs by the time the link goes down, complete on the AXI4 bus, but do not return completion TLPs on the PCIe bus.

Root Port

When configured to support Root Port functionality, the AXI Memory Mapped to PCI Express core fully supports Root Port operation as supported by the underlying block. There are a few details that need special consideration. The following subsections contain information and design considerations about Root Port support.

Power Limit Message TLP

The AXI Memory Mapped to PCI Express core automatically sends a Power Limit Message TLP when the Master Enable bit of the Command Register is set. The software must set the Requester ID register before setting the Master Enable bit to ensure that the desired Requester ID is used in the Message TLP.

Root Port Configuration Read

When an ECAM access is performed to the primary bus number, self-configuration of the integrated block for PCIe is performed. A PCIe configuration transaction is not performed and is not presented on the link. When an ECAM access is performed to the bus number that is equal to the secondary bus value in the Enhanced PCIe type 1 configuration header, then type 0 configuration transactions are generated.

When an ECAM access is attempted to a bus number that is in the range defined by the secondary bus number and subordinate bus number range (not including secondary bus number), then type 1 configuration transactions are generated. The primary, secondary and subordinate bus numbers are written and updated by Root Port software to the type 1 PCI Configuration Header of the AXI Memory Mapped to PCI Express core in the enumeration procedure.

When an ECAM access is attempted to a bus number that is out of the range defined by the secondary bus_number and subordinate bus number, the bridge does not generate a configuration request and signal a SLVERR response on the AXI4-Lite bus.

When a Unsupported Request (UR) response is received for a configuration read request, all ones are returned on the AXI4-Lite bus to signify that a device does not exist at the requested device address. It is the responsibility of the software to ensure configuration write requests are not performed to device addresses that do not exist. However, the AXI Memory Mapped to PCI Express core asserts SLVERR response on the AXI4-Lite bus when a configuration write request is performed on device addresses that do not exist or a UR response is received.

Root Port BAR

During core customization in the Vivado Design Suite, select the **Hide RP BAR** option to force the core to return zeros when User Logic or CPU is probing the RP BAR0 register and to mask any write to this register to prevent any address from being allocated for Root Port BAR. The base address then defaults to 0x0000_0000.

When **BAR 64 bit Enabled** is selected, the maximum BAR size for Root Port configuration is 4 GB. This option allows all 32-bits of AXI addresses to be addressable from the PCIe link. Inherently, AXI-PCIe BAR translation is then disabled when 4 GB is selected.

BAR 0 cannot be disabled. If a system must have no BAR for Root Port and accepts all incoming packets, select **Hide RP BAR**, select **BAR 64 bit Enabled**, and set the Bar 0 size to **4 Gigabytes**.

The Root Port BAR customization options in the Vivado Design Suite are found in [PCIe Base Address Registers in Chapter 4](#).

Unsupported Request to Upstream Traffic

To receive upstream traffic from a connected device, the Root Ports PCIe BAR_0 must be configured. If BAR_0 is not configured, the Root Port responds to requests with a Completion returned with Unsupported Request status.

Configuration Transaction Timeout

Configuration transactions are non-posted transactions. The AXI Memory Mapped to PCI Express core has a timer for timeout termination of configuration transactions that have not completed on the PCIe link. SLVERR is returned when a configuration timeout occurs. Timeout of configuration transactions are flagged by an interrupt as well.

Abnormal Configuration Transaction Termination Responses

Responses on AXI4-Lite to abnormal terminations to configuration transactions are shown in [Table 3-8](#).

Table 3-8: Responses of AXI Memory Mapped to PCI Express to Abnormal Configuration Terminations

Transfer Type	Abnormal Condition	Bridge Response
Config Read or Write	Bus number not in the range of primary bus number through subordinate bus number.	SLVERR response is asserted.
Config Read or Write	Valid bus number and completion timeout occurs.	SLVERR response is asserted.
Config Read or Write	Completion timeout.	SLVERR response is asserted.
Config Write	Bus number in the range of secondary bus number through subordinate bus number and UR is returned.	SLVERR response is asserted.

Design Flow Steps

This chapter describes customizing and generating the core, constraining the core, and the simulation, synthesis and implementation steps that are specific to this IP core. More detailed information about the Vivado™ design flows and the Vivado IP integrator can be found in the following Vivado Design Suite user guides:

- *Vivado Design Suite User Guide: Designing IP Subsystems using IP Integrator* (UG994) [\[Ref 15\]](#)
- *Vivado Design Suite User Guide: Designing with IP* (UG896) [\[Ref 12\]](#)
- *Vivado Design Suite User Guide: Getting Started* (UG910) [\[Ref 11\]](#)
- *Vivado Design Suite User Guide: Logic Simulation* (UG900) [\[Ref 13\]](#)

Customizing and Generating the Core

This section includes information about using the Vivado Design Suite to customize and generate the core.

Note: If you are customizing and generating the core in the Vivado IP integrator, see the *Vivado Design Suite User Guide: Designing IP Subsystems using IP Integrator* (UG994) [\[Ref 15\]](#) for detailed information. IP integrator might auto-compute certain configuration values when validating or generating the design. To check whether the values do change, see the description of the parameter in this chapter. To view the parameter value you can run the `validate_bd_design` command in the Tcl Console.

You can customize the IP for use in your design by specifying values for the various parameters associated with the IP core using the following steps:

1. Select the IP from the IP catalog.
2. Double-click the selected IP, or select the **Customize IP** command from the toolbar or right-click menu.

For details, see the *Vivado Design Suite User Guide: Designing with IP* (UG896) [Ref 12] and the *Vivado Design Suite User Guide: Getting Started* (UG910) [Ref 11].

Note: Figures in this chapter are illustrations of the Vivado Integrated Design Environment (IDE). This layout might vary from the current version.

Customizing the Core

The AXI Memory Mapped to PCI Express core customization parameters are described in the following sections.

Basics Parameter Settings

The initial customization screen shown in Figure 4-1 is used to define the basic parameters for the core, including the component name, reference clock frequency, and silicon type.

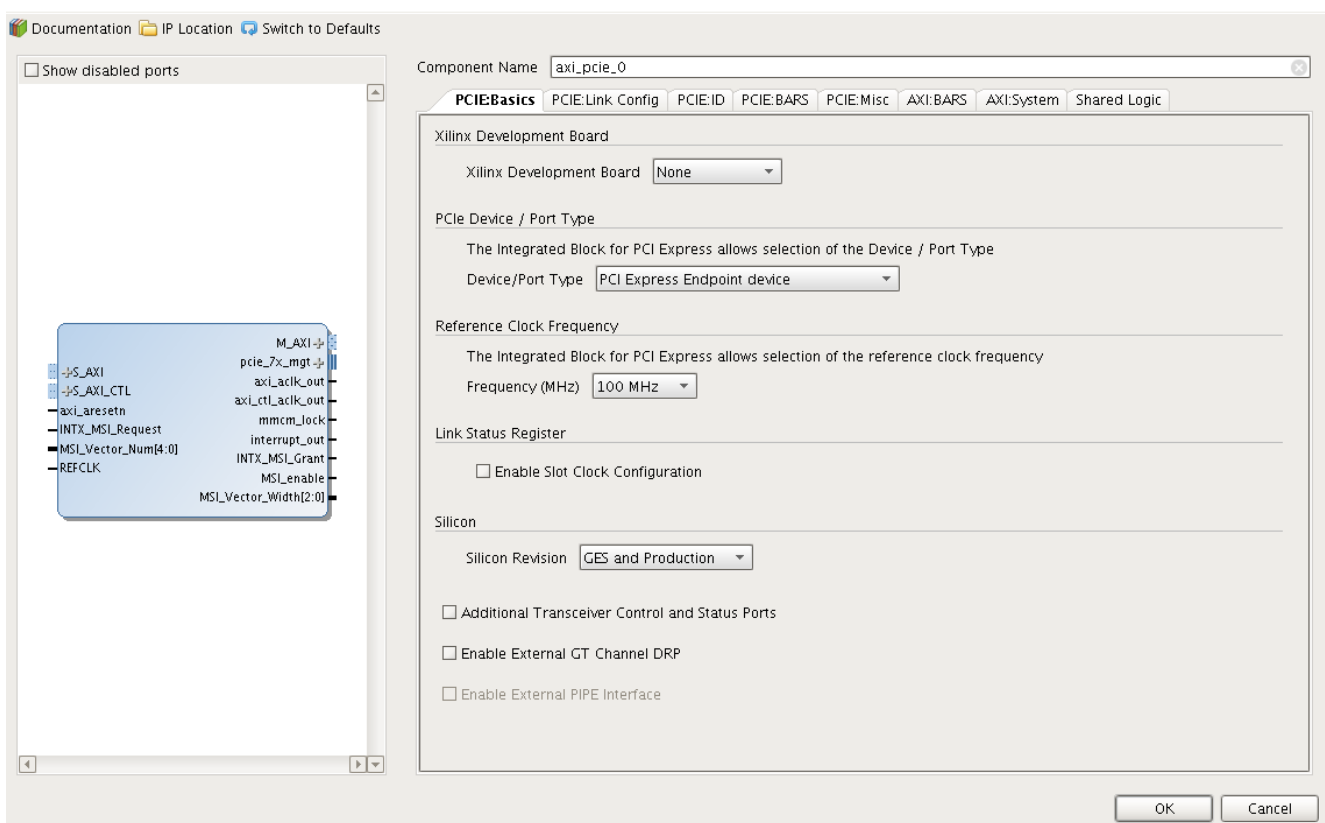


Figure 4-1: Basics Parameter Settings

Component Name

Base name of the output files generated for the core. The name must begin with a letter and can be composed of these characters: a to z, 0 to 9, and "_."



IMPORTANT: *The name cannot be the same as a core module name; for example, "axi_pcie" is a reserved name.*

PCIe Device/Port Type

Indicates the PCI Express logical device type.

Reference Clock Frequency

Selects the frequency of the reference clock provided on the `refclk` input.

Slot Clock Configuration

Enables the Slot Clock Configuration bit in the Link Status register. Selecting this option means the link is synchronously clocked. See [Clocking, page 40](#) for more information on clocking options.

Silicon Type

Selects the silicon type.

Additional Transceiver Control and Status Ports

When this option is checked, the transceiver debug ports are enabled. See [Table B-1](#) for a list of these ports.

Enable External GT Channel DRP

The external GT channel DRP ports are pulled out to the core top.

- `ext_ch_gt_drpdo[15:0]`
- `ext_ch_gt_drpd[15:0]`
- `ext_ch_gt_drpen[0:0]`
- `ext_ch_gt_drwe[0:0]`
- `ext_ch_gt_drprdy[:0]`
- `ext_ch_gt_drpaddr[8:0]`

`gt_ch_drp_rdy` indicates that the external GT channel DRP is ready to use, and that it is not in use by internal logic.

Enable External PIPE Interface

When selected, this option enables an external third-party bus functional Model (BFM) to connect to the PIPE interface of the AXI PCIe core. Note that this option is disabled by default and is enabled when the **Include Shared Logic (Clocking) in example design**

option is selected in the Shared Logic page of the Vivado IDE. This is applicable for both Endpoint and Root Port configurations.

PCIe Link Configuration

The PCIe Link Config page is shown in [Figure 4-2](#).

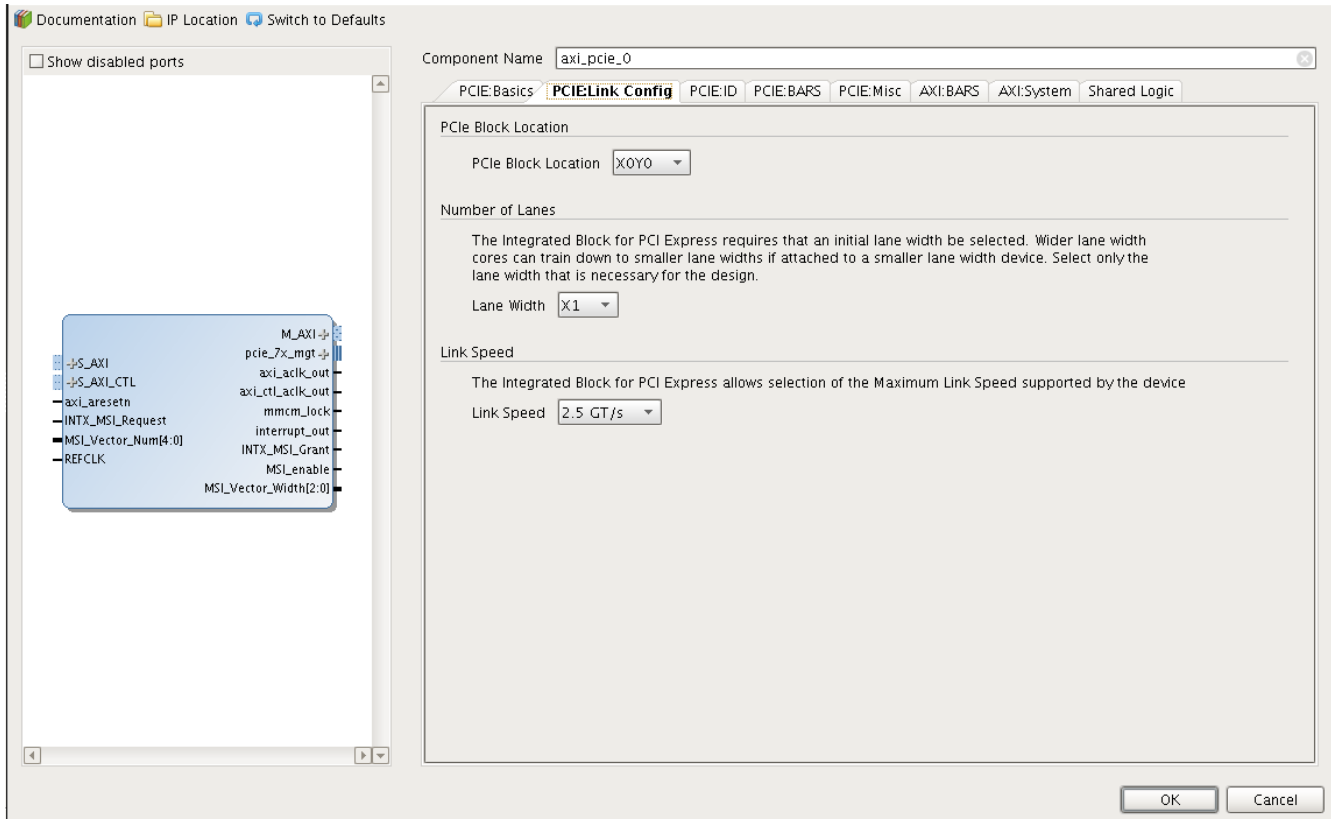


Figure 4-2: PCIe Link Configuration

Number of Lanes

The AXI Memory Mapped to PCI Express[®] core requires the selection of the initial lane width. [Table 4-1](#) defines the available widths and associated generated core. Wider lane width cores can train down to smaller lane widths if attached to a smaller lane-width device.

Table 4-1: Lane Width and Product Generated

Lane Width	Product Generated
x1	1-Lane
x2	2-Lane
x4	4-Lane
x8	8-Lane

Link Speed

The AXI Memory Mapped to PCI Express core allows you to select the maximum link speed supported by the device. [Table 4-2](#) defines the lane widths and link speeds supported by the device. Higher link speed cores are capable of training to a lower link speed if connected to a lower link speed capable device.

Table 4-2: Lane Width and Link Speed

Lane Width	Link Speed
x1	2.5 Gb/s, 5 Gb/s
x2	2.5 Gb/s, 5 Gb/s
x4	2.5 Gb/s, 5 Gb/s
x8	2.5 Gb/s

PCIe Block Location

The AXI Memory Mapped to PCI Express core allows you to select the PCI Express Hard Block within AMD FPGAs.

PCIe ID Settings

The Identity Settings pages are shown in [Figure 4-3](#). These settings customize the IP initial values and device class code.

Figure 4-3: PCIe ID Settings

ID Initial Values

- **Vendor ID:** Identifies the manufacturer of the device or application. Valid identifiers are assigned by the PCI™ Special Interest Group to guarantee that each identifier is unique. The default value, 10EEh, is the Vendor ID for AMD. Enter your vendor identification number here. FFFFh is reserved.
- **Device ID:** A unique identifier for the application; the default value, which depends on the configuration selected, is 70<link speed><link width>h. This field can be any value; change this value for the application.



IMPORTANT: Since this Device ID parameter is evaluated based on the device family (7 in case of 7 series IPs), EP, or RP mode, link width and link speed, it is expected to change whenever any of these parameters changes. For example, if the value of the Device ID is changed first and then any one of the link speed, link width, EP, or RP is changed, then the Device ID value is re-evaluated and replaces the previous value set by you. Hence, AMD recommends changing the link width, link speed, and Device Port type first and then the Device ID value.

- **Revision ID:** Indicates the revision of the device or application; an extension of the Device ID. The default value is 00h; enter values appropriate for the application.

- **Subsystem Vendor ID:** Further qualifies the manufacturer of the device or application. Enter a Subsystem Vendor ID here; the default value is 10EEh. Typically, this value is the same as Vendor ID. Setting the value to 0000h can cause compliance testing issues.
- **Subsystem ID:** Further qualifies the manufacturer of the device or application. This value is typically the same as the Device ID; the default value depends on the lane width and link speed selected. Setting the value to 0000h can cause compliance testing issues.

Class Code

The Class Code identifies the general function of a device, and is divided into three byte-size fields. The Vivado IDE allows you to either enter the 24-bit value manually (default) by either selecting the **Enter Class Code Manually** checkbox or using the Class Code lookup assistant to populate the field. De-select the checkbox to enable the Class Code assistant.

- **Base Class:** Broadly identifies the type of function performed by the device.
- **Sub-Class:** More specifically identifies the device function.
- **Interface:** Defines a specific register-level programming interface, if any, allowing device-independent software to interface with the device.

Class code encoding can be found in the PCI-SIG® specifications [\[Ref 8\]](#).

Class Code Look-up Assistant

The Class Code Look-up Assistant provides the Base Class, Sub-Class and Interface values for a selected general function of a device. This Look-up Assistant tool only displays the three values for a selected function. You must enter the values in Class Code for these values to be translated into device settings.

PCIe Base Address Registers

The PCIe Base Address Registers (BARs) screen shown in [Figure 4-4](#) set the base address register space for the Endpoint configuration. Each BAR (0 through 5) configures the BAR Aperture Size and Control attributes of the Physical Function, as described in [Table 4-3](#).

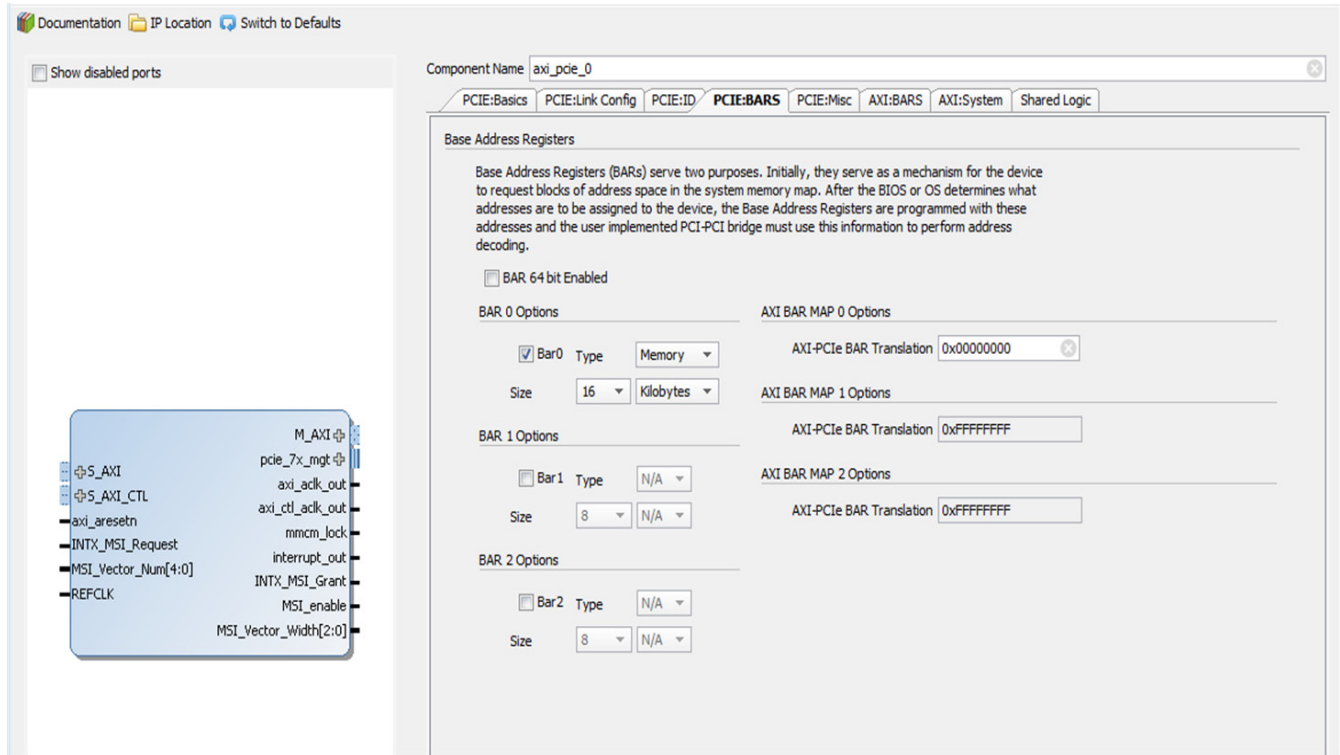


Figure 4-4: PCIe Base Address Register

Base Address Register Overview

The AXI Memory Mapped to PCI Express core in Endpoint configuration supports up to three 32-bit BARs or three 64-bit BARs. The AXI Memory Mapped to PCI Express in Root Port configuration supports one 32-bit BARs or one 64-bit BAR.

All BAR registers share these options:

- **Checkbox:** Click the checkbox to enable BAR. Deselect the checkbox to disable BAR.
- **Type:** BARs can be **Memory** apertures only. Memory BARs can be either 64-bit or 32-bit. Prefetch is enabled for 64-bit and not enabled for 32-bit. When a BAR is set as 64 bits, it uses the next BAR for the extended address space, making it inaccessible.
- **Size:** The available Size range depends on the PCIe® Device/Port Type and the Type of BAR selected. [Table 4-3](#) lists the available BAR size ranges.

Table 4-3: BAR Size Ranges for Device Configuration

PCIe Device/Port Type	BAR Type	BAR Size Range
PCI Express Endpoint	32-bit Memory	8 KB - 2 GB
	64-bit Memory	8 KB - 2 GB
Root Port of PCI Express Root Complex	32-bit Memory	8 KB - 2 GB
	64-bit Memory	8 KB - 4 GB

Interrupt Pin

Indicates the usage of Legacy interrupts. The AXI Memory Mapped to PCI Express core implements INTA only.

Intercept and Decode Incoming MSI Packet

Indicates that a received Memory Write packet with an address equal to the value of Root Port MSI Base register will be decoded and reported as an MSI interrupt in the Root Port Interrupt FIFO register.

MSI Vectors Requested

Indicates the number of MSI vectors requested by the core.

Completion Timeout Configuration

Indicates the completion timeout value for incoming completions due to outstanding memory read requests.

Dynamic Slave Bridge Address Translation

Enables the address translation vectors within the AXI Memory Mapped to PCI Express bridge logic to be changed dynamically through the AXI4-Lite interface.

AXI Base Address Registers

The AXI Base Address Registers (BARs) screen shown in [Figure 4-6](#) sets the AXI base address registers and the translation between AXI Memory space and PCI Express Memory space. Each BAR has a Base Address, High Address, and translation field which can be configured through the Vivado IDE.

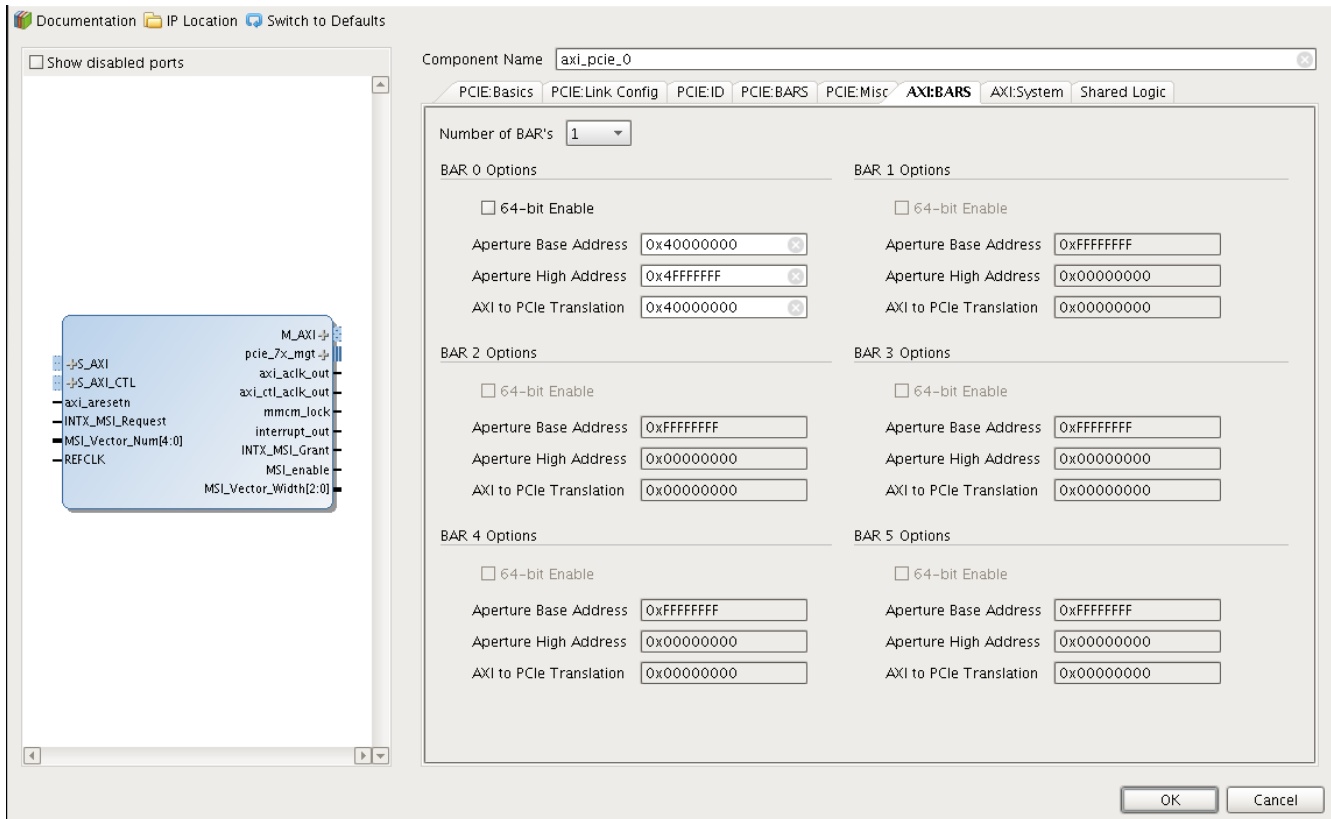


Figure 4-6: AXI Base Address Registers

Number of BARs

Indicates the number of AXI BARs enabled. The BARs are enabled sequentially.

64-bit Enable

Allows 32-bit AXI address space to be mapped into 64-bit PCIe address space.

Aperture Base Address

Sets the base address for the address range of AXI BAR. You should edit this parameter to fit design requirements.

Aperture High Address

Sets the upper address threshold for the address range of AXI BAR. You should edit this parameter to fit design requirements.

AXI to PCIe Translation

Configures the translation mapping between AXI and PCI Express address space. You should edit this parameter to fit design requirements.

AXI System

The AXI System screen shown in [Figure 4-7](#) sets the AXI Addressing and AXI Interconnect parameters.

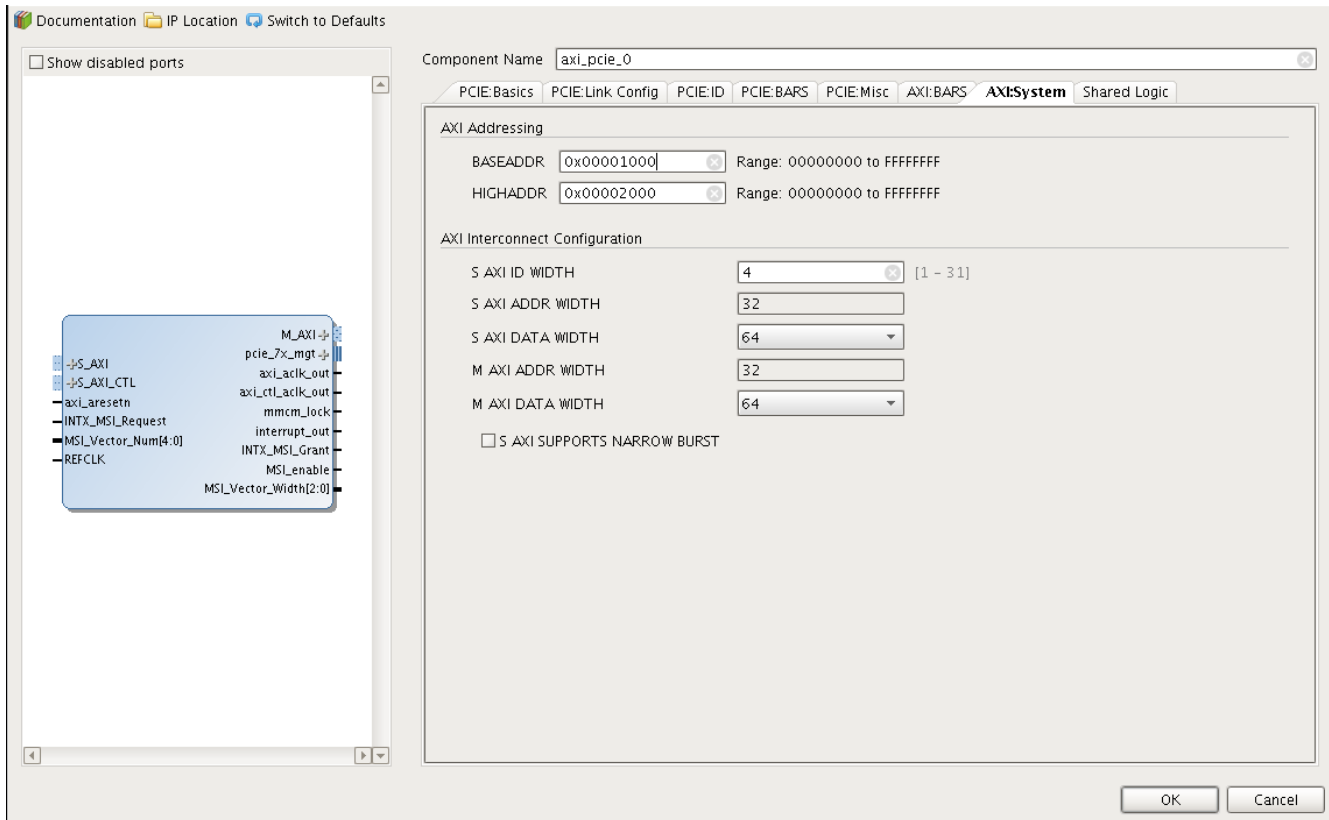


Figure 4-7: AXI System Settings

BASEADDR

Sets the AXI Base Address for the device. You should edit this parameter to fit design requirements.

HIGHADDR

Sets the AXI High Address threshold for the device. You should edit this parameter to fit design requirements.

S AXI ID WIDTH

Sets the ID width for the AXI Slave Interface.

Note: Multiple IDs are not supported for AXI Master Interface. Therefore, all signals concerned with ID are not available at AXI Master Interface.

S AXI ADDR WIDTH

AXI supports 32-bit addressing so this field is always set to 32.

S AXI DATA WIDTH

Sets the data bus width for the AXI Slave interface. This can be 64-bit or 128-bit based on your requirements. For X4G2 and X8G1, the core supports only 128-bit to achieve maximum performance.

M AXI ADDR WIDTH

AXI supports 32-bit addressing so this field is always set to 32.

M AXI DATA WIDTH

Sets the data bus width for the AXI Master interface. This can be 64-bit or 128-bit based on your requirement. For X4G2 and X8G1, the core supports only 128-bit to achieve maximum performance.

S AXI SUPPORTS NARROW BURST

Configures the IP to accept narrow burst transactions. When not enabled, the IP is optimized accordingly.

Shared Logic

Enables you to share common blocks across multiple instantiations by selecting one or more of the options on this page. For a details description of the shared logic feature, see [Shared Logic in Chapter 3](#).

Add. Debug Options

[Figure 4-8](#) consists of an option to enable the JTAG debugger for debugging purpose.

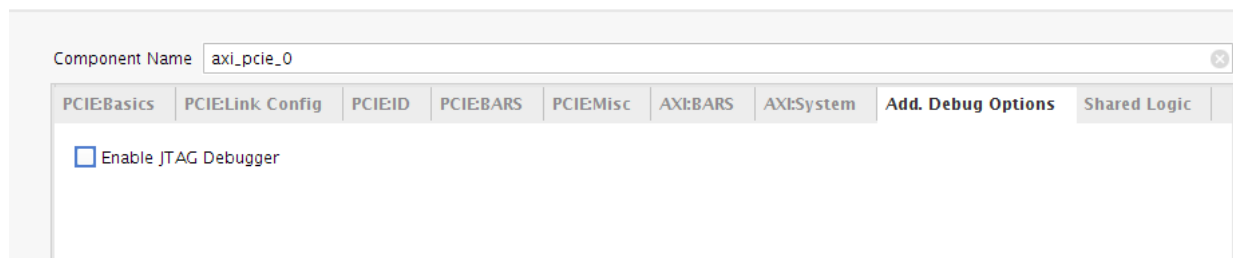


Figure 4-8: Add. Debug Options

The feature provides debug usability for the following:

- Option to integrate PCIe debug options in the IP core top which is needed to be included in the user design.

- Assumes an external clock to be used.
- **LTSSM State Transitions:** This shows all the LTSSM state transitions that have been made starting from link-up.
- **PHY Reset FSM Transitions:** This shows the PHY reset FSM (internal state machine that is used by the PCIe solution IP).
- **Receiver Detect:** This shows all the lanes that have completed Receiver Detect successfully.

The following are the steps:

1. Open a new Vivado Design Suite and connect to the board.
2. The `hw_axi_1` should be present.

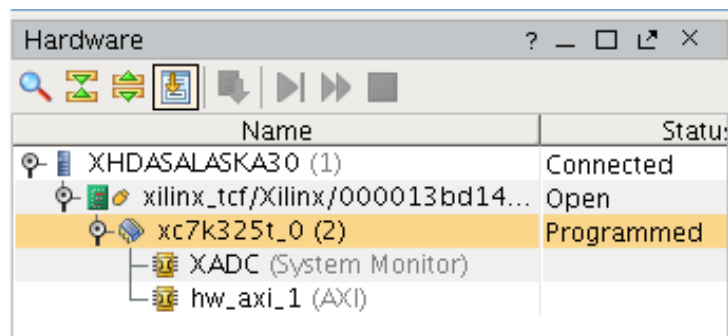


Figure 4-9: `hw_axi_1` Location

3. Type `source test_rd.tcl` in the Vivado Tcl Console.
4. For post-processing, double-click the following:
 - `draw_ltssm.tcl` (Windows) or `wish draw_ltssm.tcl`
 - `draw_reset.tcl` (Windows) or `wish draw_reset.tcl`
 - `draw_rxdet.tcl` (Windows) or `wish draw_rxdet.tcl`

This displays the pictorial representation of the LTSSM state transitions.

Output Generation

For details, see the *Vivado Design Suite User Guide: Designing with IP* (UG896) [Ref 12].

For information regarding the example design, see [Example Design Output Structure in Chapter 5](#).

Constraining the Core

This section contains information about constraining the core in the Vivado Design Suite.

Required Constraints

The AXI Memory Mapped to PCI Express core requires a clock period constraint for the REFCLK input that agrees with the C_REF_CLK_FREQ parameter setting. In addition, pin-placement (LOC) constraints are needed that are board/part/package specific.

See [Placement Constraints](#) for more details on the constraint paths for FPGA architectures.

Additional information on clocking can be found in the AMD Solution Center for PCI Express (see [Solution Center for PCI Express](#)).

System Integration

Some additional components to this system in the Vivado IP integrator can include the need to connect the MicroBlaze™ processor or Zynq™ device Arm® processor peripheral to communicate with PCI Express (in addition to the AXI4-Lite register port on the PCIe bridge). The AXI Interconnect provides this capability and performs the necessarily conversions for the various AXI ports that might be connected to the AXI Interconnect IP (described in [\[Ref 9\]](#)).

The AXI Memory Mapped to PCI Express core can be configured with each port connection for an AXI Vivado IP integrator system topology. When instantiating the core, ensure the following bus interface tags are defined.

```
BUS_INTERFACE M_AXI
BUS_INTERFACE S_AXI
BUS_INTERFACE S_AXI_CTL
```

PCIe Clock Integration

The PCIe differential clock input in the system might need to use a differential input buffer (that is instantiated separately) from the AXI Memory Mapped to PCI Express core. The Vivado IP integrator automatically inserts the appropriate clock buffer.

Placement Constraints

The AXI Memory Mapped to PCI Express core provides a Xilinx design constraint (XDC) file for all supported PCIe, Part, and Package permutations. You can find the generated XDC file in the Sources tab of the Vivado IDE after generating the IP in the Customize IP dialog box.

For design platforms, it might be necessary to manually place and constrain the underlying blocks of the AXI Memory Mapped to PCI Express core. The modules to assign a LOC constraint include:

- the embedded integrated block for PCIe itself
- the GTX transceivers (for each channel)
- the PCIe differential clock input (if utilized)

The following subsection describes the example location constraints.

Location Constraints for Virtex 7 and Kintex 7 FPGAs

This section highlights the LOC constraints to be specified in the XDC file for the AXI Memory Mapped to PCI Express core for design implementations.

For placement/path information on the integrated block for PCIe itself, the following constraint can be utilized:

```
set_property LOC PCIE_X*Y* [get_cells {U0/comp_axi_enhanced_pcie/
comp_enhanced_core_top_wrap/axi_pcie_enhanced_core_top_i/pcie_7x_v2_0_inst/
pcie_top_i/pcie_7x_i/pcie_block_i}]
```

For placement/path information of the GTX transceivers, the following constraint can be utilized:

```
set_property LOC GTXE2_CHANNEL_X*Y* [get_cells {U0/comp_axi_enhanced_pcie/
comp_enhanced_core_top_wrap/axi_pcie_enhanced_core_top_i/pcie_7x_v2_0_inst/
gt_ges.gt_top_i/pipe_wrapper_i/pipe_lane[0].gt_wrapper_i/
gtx_channel.gtxe2_channel_i}]
```

For placement/path constraints of the input PCIe differential clock source (using the example provided in [System Integration](#)), the following can be utilized:

```
set_property LOC IBUFDS_GTE2_X*Y* [get_cells {*/PCIE_Diff_Clk_I/
USE_IBUFDS_GTE2.GEN_IBUFDS_GTE2[0].IBUFDS_GTE2_I}]
```

Location Constraints for Artix 7 FPGAs

Special consideration must be given to Artix™ 7 device implementations. The same IP block constraint can be used as described previously (see [Location Constraints for Virtex 7 and Kintex 7 FPGAs, page 80](#)). However, the PCIe serial transceiver wrapper instance is different in the IP. Use the following LOC constraint for the GTP transceivers in Artix 7 devices.

```
set_property LOC GTPE2_CHANNEL_X*Y* [get_cells {U0/comp_axi_enhanced_pcie/
comp_enhanced_core_top_wrap/axi_pcie_enhanced_core_top_i/pcie_7x_v2_0_inst/
gt_ges.gt_top_i/pipe_wrapper_i/pipe_lane[0].gt_wrapper_i/
gtp_channel.gtpe2_channel_i}]
```


Also for Artix 7 devices, the GTP_COMMON must be constrained to a location. The following LOC constraint can be utilized.

```
set_property LOC GTPE2_COMMON_X*Y* [get_cells {U0/comp_axi_enhanced_pcie/  
comp_enhanced_core_top_wrap/axi_pcie_enhanced_core_top_i/pcie_7x_v2_0_inst/  
gt_ges.gt_top_i/pipe_wrapper_i/pipe_lane[0].pipe_quad.pipe_common.qpll_wrapper_i/  
gtp_common.gtpe2_common_i}]
```

Clock Frequencies

The AXI Memory Mapped to PCI Express Bridge supports reference clock frequencies of 100 MHz, and 250 MHz and is configurable within the Vivado IDE.

Clock Management

For details, see [Clocking in Chapter 3](#).

Clock Placement

For details, see [Placement Constraints](#).

Banking

This section is not applicable for this IP core.

Transceiver Placement

The Transceiver primitives adjacent to the PCIe hard block should be used to aid in the place and route of the solution IP. The adjacent Transceiver banks one above or one below the desired PCIe hard block can also be used. Transceivers outside this range are not likely to meet the timing requirements for the PCI Express Solution IP and should not be used.

Simulation

- For comprehensive information about Vivado simulation components, as well as information about using supported third party tools, see the *Vivado Design Suite User Guide: Logic Simulation* (UG900) [\[Ref 13\]](#).
- For information regarding simulating the example design, see [Simulation Design Overview in Chapter 5](#).



IMPORTANT: For cores targeting 7 series or Zynq 7000 SoC devices, UNIFAST libraries are not supported. AMD IP is tested and qualified with UNISIM libraries only.

PIPE Mode Simulations

The PIPE Simulation mode allows you to run the simulations without serial transceiver block to speed up simulations. To run the simulations using the PIPE interface to speed up the simulation, generate the core after selecting the **Enable External PIPE Interface** option in the Basic tab of the Customize IP dialog box.

For details, see [Enable External PIPE Interface, page 67](#). For third-party bus functional model support, see the *PIPE Mode Simulation Using Integrated Endpoint PCI Express Block in Gen2 x8 Configurations Application Note* (XAPP1184) [Ref 16].

Third-party simulation support pulls out the following ports when Enable External PIPE interface ports is selected. The following tables describe the ports that add to the boundary of the AXI PCIe core when this option is selected.

For PIPE ports to and from the pcie_top, each lane has independent input and output bus signals.

[Table 4-4](#) and [Table 4-5](#) describe the PIPE bus signals available at the top level of the core and their corresponding mapping inside the EP core (pcie_top) PIPE signals.



IMPORTANT: A new file, *xil_sig2pipe.v*, is delivered in the simulation directory, and the file replaces *phy_sig_gen.v*. BFM/VIPs should interface with the *xil_sig2pipe* instance in *board.v*.

Table 4-4: Common Input/Output Commands with Endpoint PIPE Signals Mapping

In Commands	Endpoint PIPE Signals Mapping	Out Commands	Endpoint PIPE Signals Mapping
common_commands_in[3:0]	not used ⁽¹⁾	common_commands_out[0]	pipe_clk ⁽²⁾
		common_commands_out[1]	pipe_tx_rate_gt ⁽³⁾
		common_commands_out[2]	pipe_tx_rcvr_det_gt
		common_commands_out[3]	pipe_tx_deemph_gt
		common_commands_out[6:4]	pipe_tx_margin_gt
		common_commands_out[11:7]	not used ⁽¹⁾

Notes:

1. These ports functionality has been deprecated and can be left unconnected.
2. pipe_clk is an output clock based on the core configuration. For Gen1 rate, pipe_clk is 125 MHz. For Gen2, pipe_clk is 250 MHz.
3. pipe_tx_rate_gt indicates the pipe rate: 1'b0 for Gen1, and 1'b1 for Gen2.

Table 4-5: Input/Output Bus with Endpoint PIPE Signals Mapping

Input Bus	Endpoint PIPE Signals Mapping	Output Bus	Endpoint PIPE Signals Mapping
pipe_rx_0_sigs[15:0]	pipe_rx0_data_gt	pipe_tx_0_sigs[15: 0]	pipe_tx0_data_gt
pipe_rx_0_sigs[17:16]	pipe_rx0_char_is_k_gt	pipe_tx_0_sigs[17:16]	pipe_tx0_char_is_k_gt
pipe_rx_0_sigs[18]	pipe_rx0_elec_idle_gt	pipe_tx_0_sigs[18]	pipe_tx0_elec_idle_gt

Table 4-5: Input/Output Bus with Endpoint PIPE Signals Mapping

Input Bus	Endpoint PIPE Signals Mapping	Output Bus	Endpoint PIPE Signals Mapping
pipe_rx_0_sigs[24:19]	not used ⁽¹⁾	pipe_tx_0_sigs[19]	pipe_tx0_compliance_gt
		pipe_tx_0_sigs[20]	pipe_rx0_polarity_gt
		pipe_tx_0_sigs[22:21]	pipe_tx0_powerdown_gt

Notes:

- 1. This ports functionality has been deprecated and can be left unconnected.
- 2. Lanes 1 to 7 use similar signal definitions.

Example Design

This chapter contains information about the example design provided in the Vivado™ Design Suite.

Overview

The example simulation design for the Endpoint configuration of the AXI-PCIe block consists of two parts.

- **Root Port Model:** a test bench that generates, consumes, and checks PCI Express bus traffic
- **AXI Block RAM Controller**

Simulation Design Overview

For the simulation design, transactions are sent from the Root Port Model to the AXI Memory Mapped to PCI Express core configured as an Endpoint and processed inside the AXI Block RAM controller design.

[Figure 5-1](#) illustrates the simulation design provided with the AXI Memory Mapped to PCI Express core.

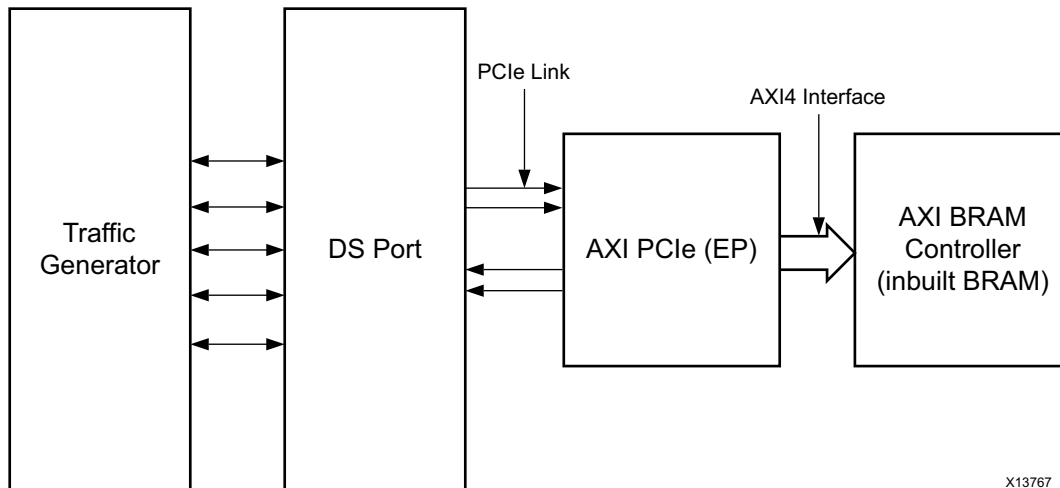


Figure 5-1: Example Design Block Diagram

Note: The example design supports Verilog as the target language.

Customizing and Generating the Example Design

In Customize IP dialog box, make the following selections for the example design.

1. In the PCIe:Basics page, the example design supports only an Endpoint (EP) device.
2. The PCIe:ID defaults are supported.
3. The PCIe:BARS defaults are supported.
4. The PCIe:Misc page defaults are supported.
5. In the AXI:BARS page, default values are assigned to the Base Address, High Address, and AXI to PCIe Translation values.
6. The AXI:System page default values are supported.

Note: After customizing the core, right-click the component name, and select **Open IP Example Design**. This opens a separate example design. Simulate the core by following the steps in the next section.

Simulating the Example Design

The example design can be run in any configuration using:

- Vivado simulator
- Cadence IES Simulator
- Mentor Graphics Questa Advanced Simulator
- VCS Simulator

Vivado Simulator

By default, the simulator is set to Vivado simulator. To run a simulation, click **Run Behavioral Simulation** in the Flow Navigator.

Cadence IES

For a Cadence IES simulation, the following steps are required:

1. In Vivado IDE, change the simulation settings as follows:
 - Target simulator: **Incisive Enterprise Simulator (IES)**
2. On the simulator tab, select **Run Simulation > Run behavioral simulation**.

Mentor Graphics Questa Advanced Simulator

For a Questa Advanced Simulator simulation, the following steps are required:

1. In the Vivado IDE, change the simulation settings as follows:
 - Target simulator: **Questa Advanced Simulator/ModelSim**
2. On the Simulator tab, select **Run Simulation > Run behavioral simulation**.

VCS Simulator

For a VCS simulation, the following steps are required:

1. In Vivado IDE, change the simulation settings as follows:
 - Target simulator: **Verilog Compiler Simulator (VCS)**
2. On the simulator tab, select **Run Simulation > Run behavioral simulation**.



IMPORTANT: *Simulation is not supported for configurations with the Silicon Revision option set to IES. Only implementation is supported.*

Implementation Design Overview

For implementation design, the AXI Block RAM controller can be used as a scratch pad memory to write and read to Block RAM locations.

Example Design Elements

The core wrapper includes:

- An example Verilog HDL or VHDL wrapper (instantiates the cores and example design).
- A customizable demonstration test bench to simulate the example design.

Example Design Output Structure

Figure 5-2 provides the output structure of the example design.

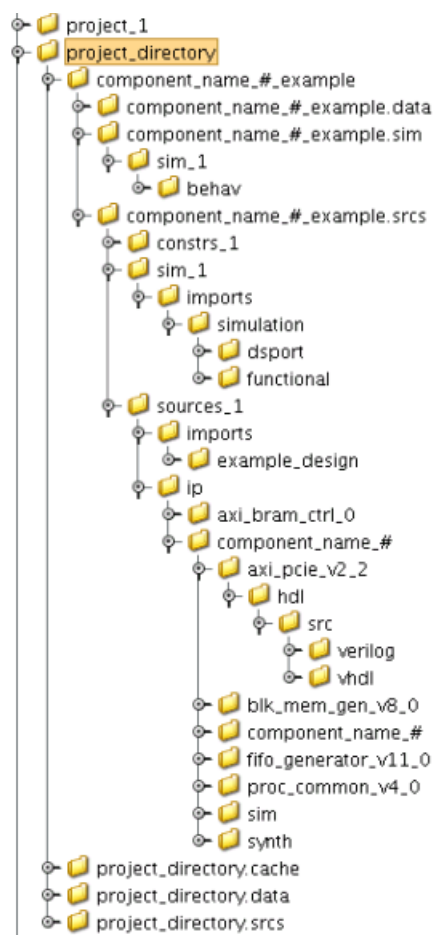


Figure 5-2: Example Design Output Structure

Table 5-1 provides a descriptions of the contents of the example design directories.

Table 5-1: Example Design Structure

Directory	Description
project_1/axi_pcie_0_example	Contains all example design files.
project_1/axi_pcie_0_example/ axi_pcie_0_example.srcs/sources_1/imports/ example_design/	Contains the top module for the example design, xilinx_axi_pcie_ep.v.
project_1/axi_pcie_0_example/ axi_pcie_0_example.srcs/sources_1/ip/axi_pcie_0	Contains the XDC file based on device selected, all design files and subcores used in axi_pcie, and the top modules for simulation and synthesis.
project_1/axi_pcie_0_example/ axi_pcie_0_example.srcs/sources_1/ip/ axi_bram_ctrl_0	Contains block RAM controller files used in example design.
project_1/axi_pcie_0_example/ axi_pcie_0_example.srcs/sim_1/imports/ simulation/dsprot	Contains all RP files, cgator and PIO files.
project_1/axi_pcie_0_example/ axi_pcie_0_example.srcs/sim_1/imports/ simulation/functional	Contains the test bench file.
project_1/axi_pcie_0_example/ axi_pcie_0_example.srcs/constrs_1/imports/ example_design	Contains the example design XDC file.

Test Bench

This chapter contains information about the test benches provided in the Vivado™ Design Suite environment.

Root Port Model Test Bench for Endpoint

The PCI Express® Root Port Model is a robust test bench environment that provides a test program interface that can be used with the provided PIO design or with a user design. The purpose of the Root Port Model is to provide a source mechanism for generating downstream PCI Express TLP traffic to stimulate your design, and a destination mechanism for receiving upstream PCI Express TLP traffic from your design in a simulation environment.

Source code for the Root Port Model is included to provide the model for a starting point for your test bench. All the significant work for initializing configuration space, creating TLP transactions, generating TLP logs, and providing an interface for creating and verifying tests are complete, allowing you to dedicate efforts to verifying the correct functionality of the design rather than spending time developing an Endpoint core test bench infrastructure.

The Root Port Model consists of:

- Test Programming Interface (TPI), which allows you to stimulate the Endpoint device for the model.
- Example tests that illustrate how to use the test program TPI.

Figure 1 illustrates the Root Port Model coupled with the PIO design.

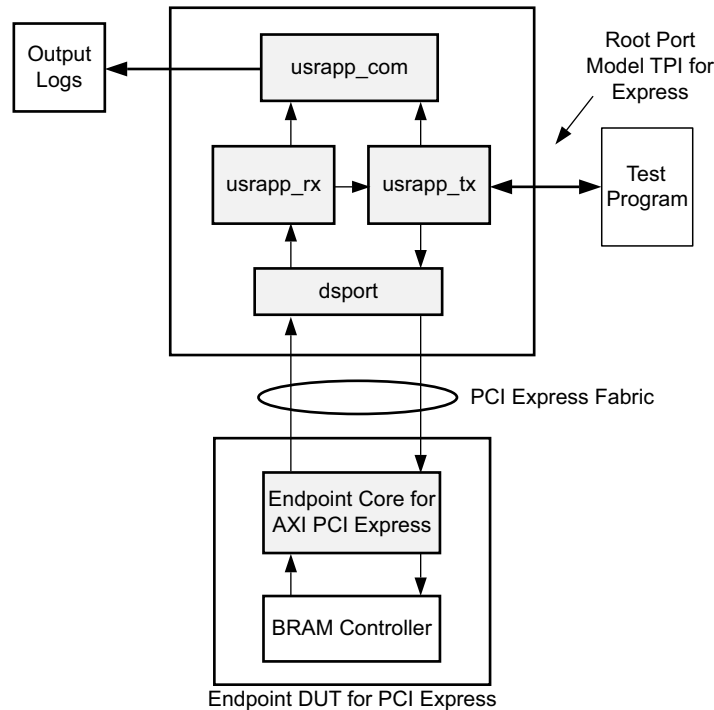


Figure 6-1: Root Port Model for AXI_PCIE Endpoint

Architecture

The Root Port Model consists of these blocks, illustrated in [Figure 6-1](#):

- dsport (Root Port)
- usrapp_tx
- usrapp_rx
- usrapp_com (Verilog only)

The usrapp_tx and usrapp_rx blocks interface with the dsport block for transmission and reception of TLPs to/from the Endpoint Design Under Test (DUT). The Endpoint DUT consists of the Endpoint for AXI-PCIe and the Block RAM controller design (displayed) or customer design.

The usrapp_tx block sends TLPs to the dsport block for transmission across the PCI Express Link to the Endpoint DUT. In turn, the Endpoint DUT device transmits TLPs across the PCI Express Link to the dsport block, which are subsequently passed to the usrapp_rx block. The dsport and core are responsible for the data link layer and physical link layer processing when communicating across the PCI Express logic. Both usrapp_tx and usrapp_rx utilize the usrapp_com block for shared functions, for example, TLP processing and log file outputting.

Transaction sequences or test programs are initiated by the usrapp_tx block to stimulate the logic interface of the Endpoint device. TLP responses from the Endpoint device are received

by the `usrapp_rx` block. Communication between the `usrapp_tx` and `usrapp_rx` blocks allow the `usrapp_tx` block to verify correct behavior and act accordingly when the `usrapp_rx` block has received TLPs from the Endpoint device.

Simulating the Example Design

To simulate the design, see [Chapter 5, Example Design](#).

Endpoint Model Test Bench for Root Port

The Endpoint model test bench for the AXI Memory Mapped to PCI Express core in Root Port configuration is a simple example test bench that connects the Configurator example design and the PCI Express Endpoint model allowing the two to operate like two devices in a physical system. Because the Configurator example design consists of logic that initializes itself and generates and consumes bus traffic, the example test bench only implements logic to monitor the operation of the system and terminate the simulation.

The Endpoint model test bench consists of:

- Verilog source code for all Endpoint model components
- PIO slave design

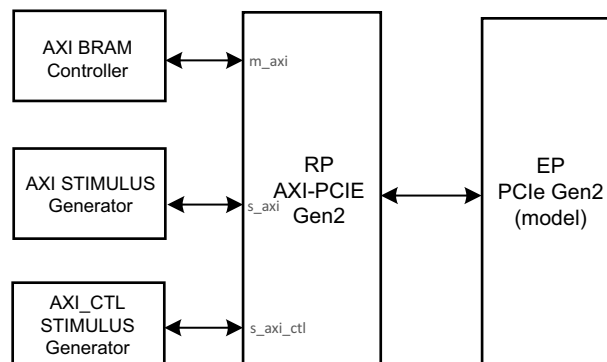


Figure 6-2: Endpoint Model for AXI_PCIE Root Port

Architecture

The Endpoint model consists of these blocks:

- PCI Express Endpoint (AXI Memory Mapped to PCI Express in Endpoint configuration) model.
- PIO slave design, consisting of:
 - `pio_rx_engine`

- `pio_tx_engine`
- `pio_ep_mem`
- `pio_to_ctrl`

The `pio_rx_engine` and `pio_tx_engine` blocks interface with the `ep` block for reception and transmission of TLPs from/to the Root Port Design Under Test (DUT). The Root Port DUT consists of the core configured as a Root Port and the Block RAM controller along with `s_axi` and `s_axi_ctrl` models to drive traffic on `s_axi` and `s_axi_ctrl`.

Simulating the Example Design

To simulate the design, see [Chapter 5, Example Design](#).

Upgrading

This appendix contains information about migrating a design from ISE® Design Suite to the Vivado™ Design Suite, and for upgrading to a more recent version of the IP core. For customers upgrading in the Vivado Design Suite, important details (where applicable) about any port changes and other impact to user logic are included.

Migrating to the Vivado Design Suite

For information on migrating to the Vivado Design Suite, see *ISE to Vivado Design Suite Migration Methodology Guide* (UG911) [\[Ref 10\]](#).

Upgrading in the Vivado Design Suite

This section provides information about any changes to the user logic or port designations that take place when you upgrade to a more current version of this core in the Vivado Design Suite.

Parameter Changes

[Table A-1](#) shows the changes to parameters in the current version of the core.

Table A-1: Parameter Changes

User Parameter Name	Display Name	New/ Change/ Removed	Details	Default Value
PCIE_CAP_SLOT_IMPLEMENTED	Enable Slot Clock Configuration	Removed	Removed this parameter and added a new one.	NA
SLOT_CLOCK_CONFIG	Enable Slot Clock Configuration	Added	Added this parameter in place of PCIE_CAP_SLOT_IMPLEMENTED.	true (Checked)
BASE_CLASS_MENU	Base Class Menu	Changed	Changed the default value.	Memory controller
SUB_CLASS_INTERFACE_MENU	Sub Class Interface Menu	Changed	Changed the default value.	Other memory controller

Table A-1: Parameter Changes (Cont'd)

User Parameter Name	Display Name	New/ Change/ Removed	Details	Default Value
en_jtag_debug	Enable JTAG Debugger	New	Added the parameter to enable the JTAG Debug option Unchecked.	Unchecked
REDUCE_OOB_FREQ	Reduce OOB Frequency	Added	Added the parameter to reduce OOB frequency	false

Port Changes

Table A-2 shows the changes to ports in the current version of the core.

Table A-2: Port Changes

Signal Name	Direction	Width
user_link_up	Out	1

Debugging

This appendix provides information for using the resources available on the AMD Support website, debug tools, and other step-by-step processes for debugging designs that use the AXI Memory Mapped to PCI Express core.

Finding Help with AMD Adaptive Computing Solutions

To help in the design and debug process when using the AXI Memory Mapped to PCI Express, the [Support web page](#) contains key resources such as product documentation, release notes, answer records, information about known issues, and links for obtaining further product support. The [Community Forums](#) are also available where members can learn, participate, share, and ask questions about AMD Adaptive Computing solutions.

Documentation

This product guide is the main document associated with the AXI Memory Mapped to PCI Express. This guide, along with documentation related to all products that aid in the design process, can be found on the [Support web page](#) or by using the AMD Adaptive Computing Documentation Navigator. Download the Documentation Navigator from the [Downloads page](#). For more information about this tool and the features available, open the online help after installation.

Answer Records

Answer Records include information about commonly encountered problems, helpful information on how to resolve these problems, and any known issues with an AMD Adaptive Computing product. Answer Records are created and maintained daily ensuring that users have access to the most accurate information available.

Answer Records for this core can be located by using the Search Support box on the main [Support web page](#). To maximize your search results, use proper keywords such as:

- Product name
- Tool message(s)

- Summary of the issue encountered

A filter search is available after results are returned to further target the results.

Master Answer Record for the AXI Memory Mapped to PCI Express

AR: [54646](#)

Extensive debugging collateral is available in AR: [56802](#).

Technical Support

AMD Adaptive Computing provides technical support at the [Community Forums](#) for this AMD LogiCORE™ IP product when used as described in the product documentation. AMD cannot guarantee timing, functionality, or support if you do any of the following:

- Implement the solution in devices that are not defined in the documentation.
- Customize the solution beyond that allowed in the product documentation.
- Change any section of the design labeled DO NOT MODIFY.

To ask questions, navigate to the [Community Forums](#).

Debug Tools

There are many tools available to address AXI Memory Mapped to PCI Express design issues. It is important to know which tools are useful for debugging various situations.

Vivado Design Suite Debug Feature

The AMD Vivado™ Design Suite debug feature inserts logic analyzer and virtual I/O cores directly into your design. The debug feature also allows you to set trigger conditions to capture application and integrated block port signals in hardware. Captured signals can then be analyzed. This feature in the Vivado IDE is used for logic debugging and validation of a design running in AMD devices.

The Vivado logic analyzer is used with the logic debug IP cores, including:

- ILA 2.0 (and later versions)
- VIO 2.0 (and later versions)

See *Vivado Design Suite User Guide: Programming and Debugging* (UG908) [[Ref 14](#)].

Reference Boards

Various AMD development boards support AXI Memory Mapped to PCI Express. These boards can be used to prototype designs and establish that the core can communicate with the system.

- 7 series evaluation boards
 - KC705
 - VC707
 - ZC706

Third-Party Tools

This section describes third-party software tools that can be useful in debugging.

LSPCI (Linux)

LSPCI is available on Linux platforms and allows you to view the PCI Express device configuration space. LSPCI is usually found in the `/sbin` directory. LSPCI displays a list of devices on the PCI buses in the system. See the LSPCI manual for all command options. Some useful commands for debugging include:

- `lspci -x -d [<vendor>]: [<device>]`

This displays the first 64 bytes of configuration space in hexadecimal form for the device with vendor and device ID specified (omit the `-d` option to display information for all devices). The default Vendor/Device ID for AMD cores is 10EE:6012. Here is a sample of a read of the configuration space of a AMD device:

```
> lspci -x -d 10EE:6012
81:00.0 Memory controller: Xilinx Corporation: Unknown device 6012
00: ee 10 12 60 07 00 10 00 00 00 80 05 10 00 00 00
10: 00 00 80 fa 00 00 00 00 00 00 00 00 00 00 00 00
20: 00 00 00 00 00 00 00 00 00 00 00 00 00 ee 10 6f 50
30: 00 00 00 00 40 00 00 00 00 00 00 00 00 05 01 00 00
```

Included in this section of the configuration space are the Device ID, Vendor ID, Class Code, Status and Command, and Base Address Registers.

- `lspci -xxxx -d [<vendor>]: [<device>]`

This displays the extended configuration space of the device. It can be useful to read the extended configuration space on the root and look for the Advanced Error Reporting (AER) registers. These registers provide more information on why the device has flagged an error (for example, it might show that a correctable error was issued because of a replay timer timeout).

- `lspci -k`

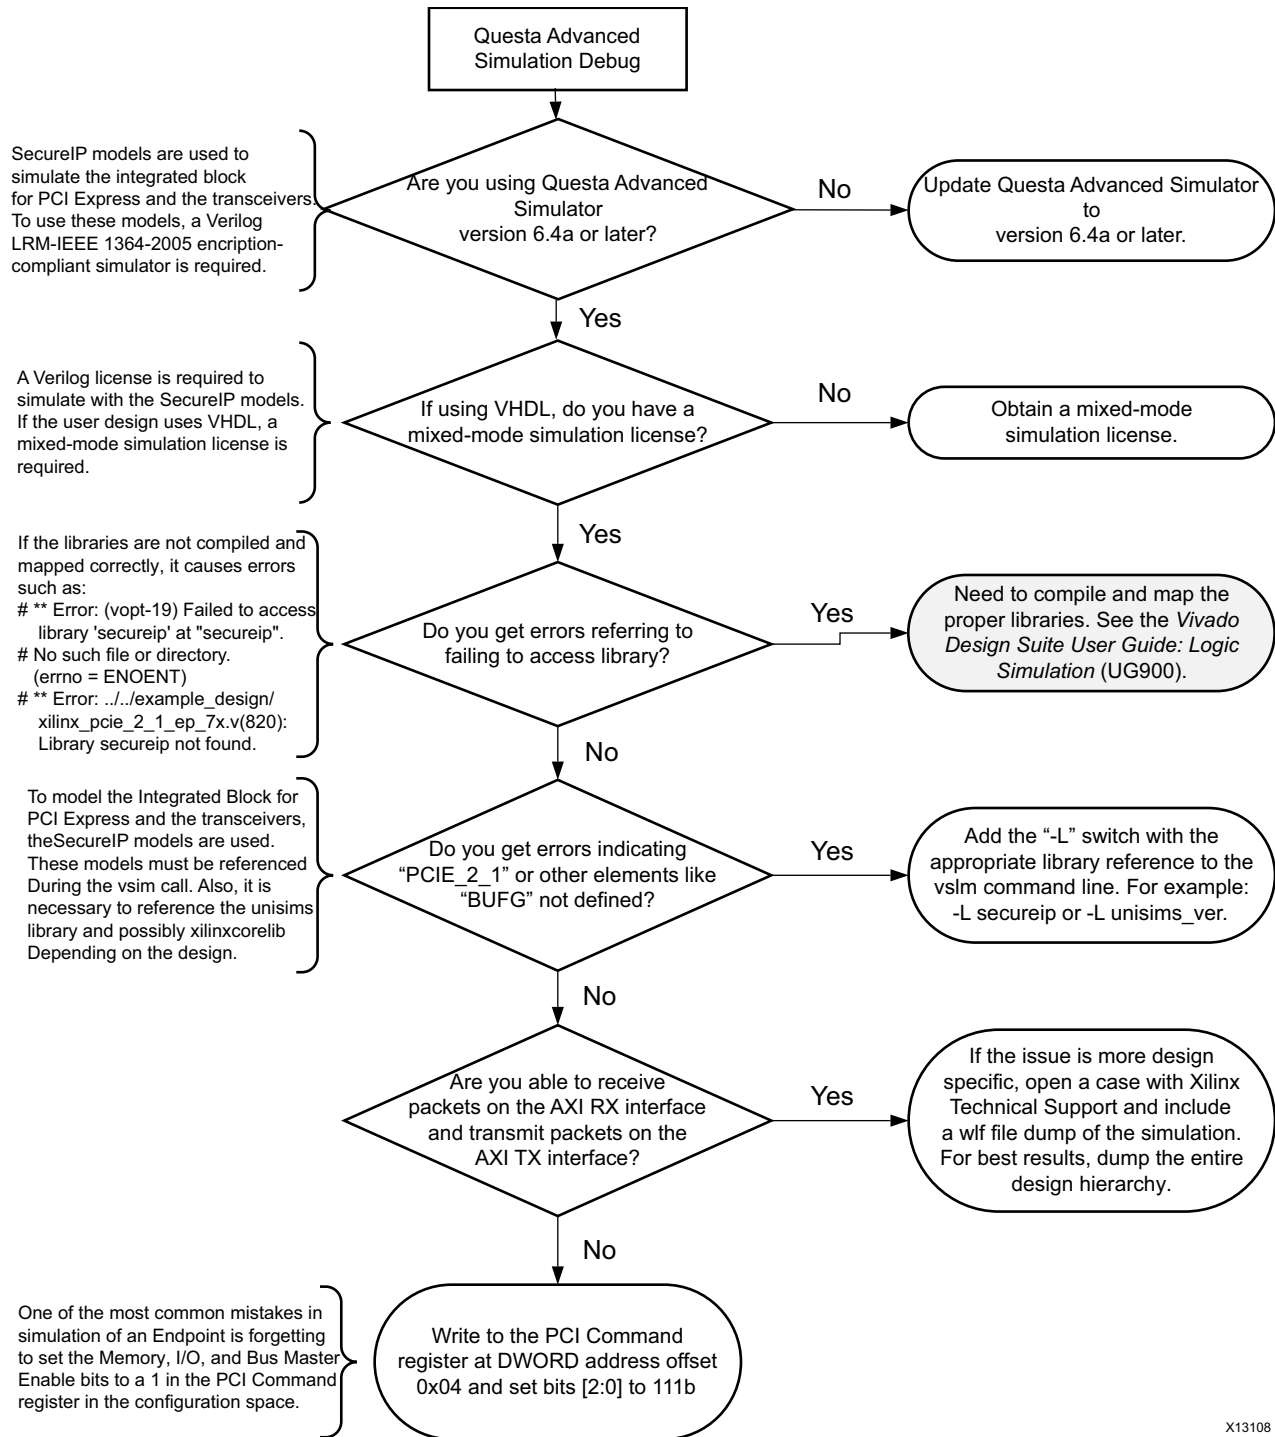
Shows kernel drivers handling each device and kernel modules capable of handling it (works with kernel 2.6 or later).

PCI-SIG Software Suites

PCI-SIG[®] software suites such as PCIE-CV can be used to test compliance with the specification. This software can be downloaded at www.pcisig.com.

Simulation Debug

The simulation debug flow for Mentor Graphics Questa Advanced Simulator is illustrated in [Figure B-1](#). A similar approach can be used with other simulators.

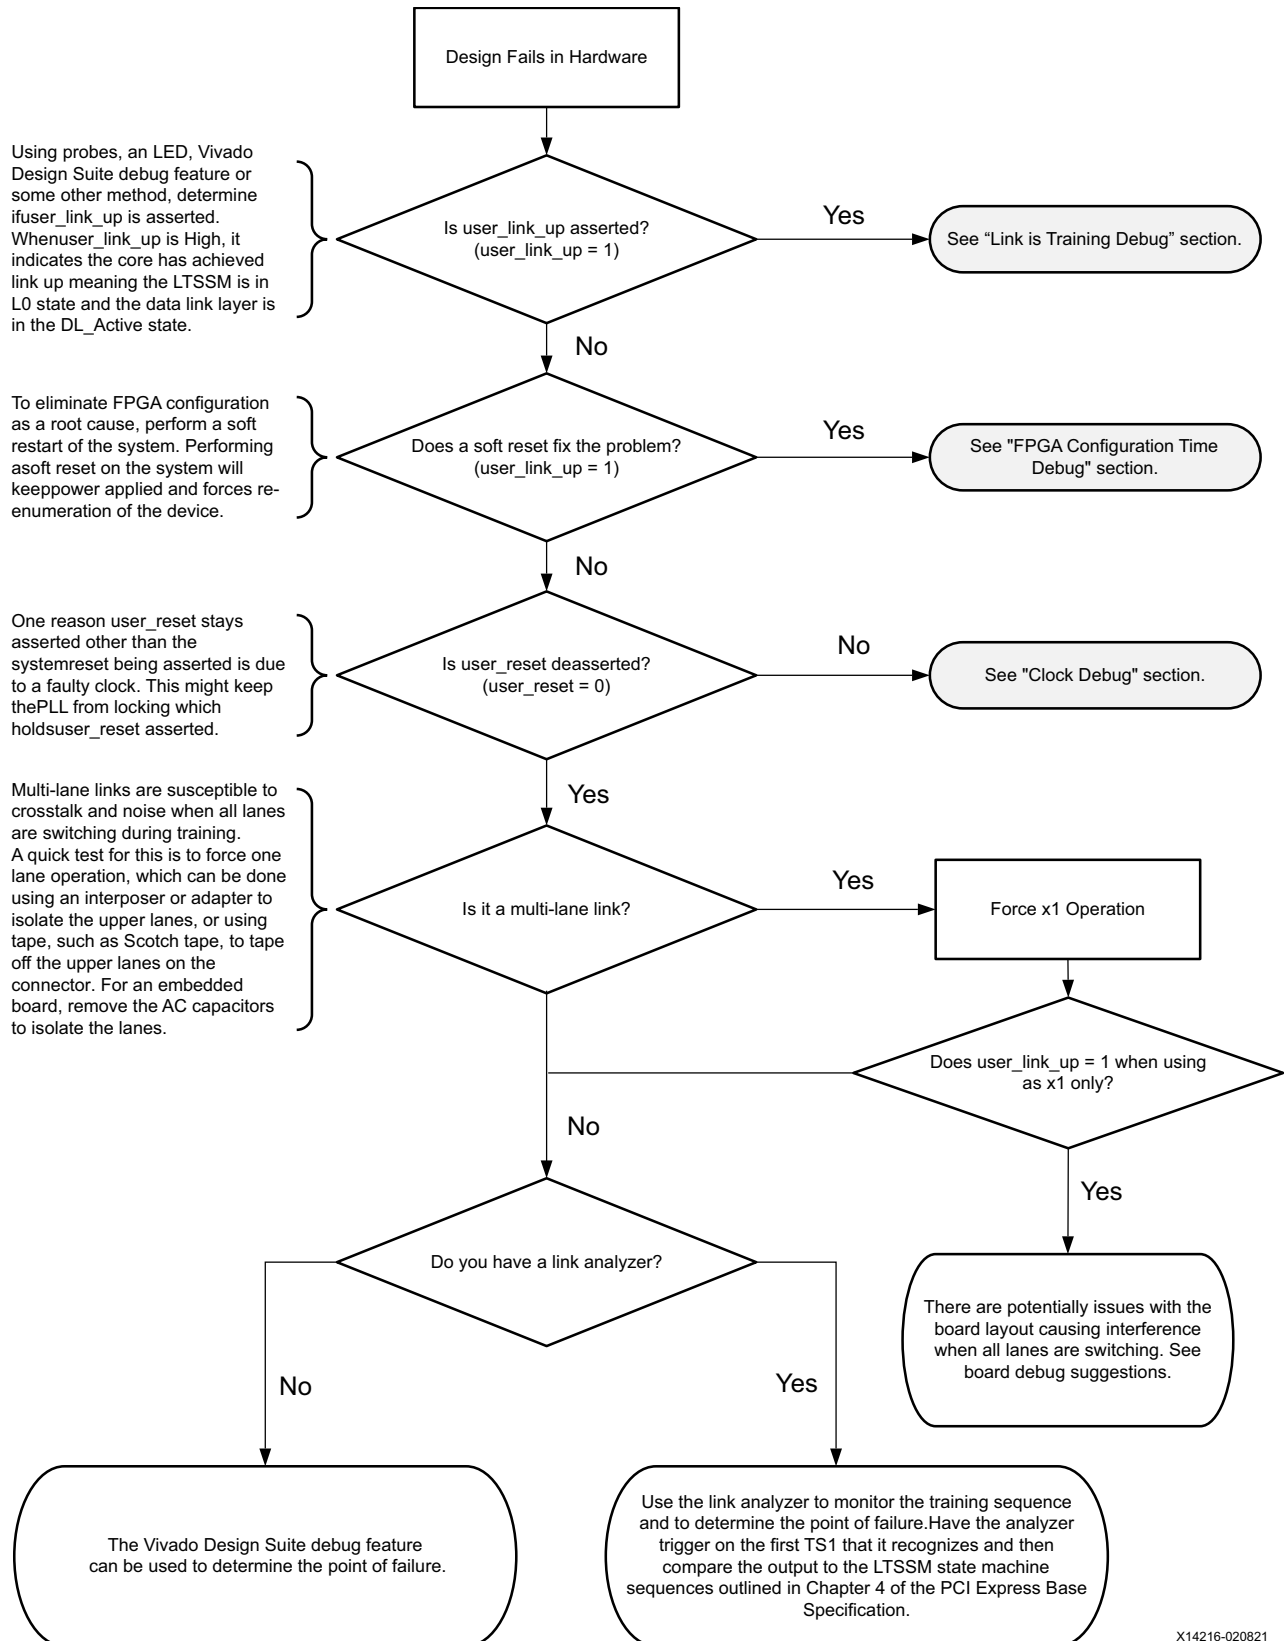


X13108

Figure B-1: Questa Advanced Simulator Simulation Debug Flow

Hardware Debug

Hardware issues can range from device recognition issues to problems seen after hours of testing. This section provides debug flow diagrams for some of the most common issues. Endpoints that are shaded gray indicate that more information is found in sections after [Figure B-2](#).



X14216-020821

Figure B-2: Design Fails in Hardware Debug Flow Diagram

FPGA Configuration Time Debug

Device initialization and configuration issues can be caused by not having the FPGA configured fast enough to enter link training and be recognized by the system. Section 6.6 of *PCI Express Base Specification, rev. 2.1* [Ref 8] states two rules that might be impacted by FPGA Configuration Time:

- A component must enter the LTSSM Detect state within 20 ms of the end of the Fundamental reset.
- A system must guarantee that all components intended to be software visible at boot time are ready to receive Configuration Requests within 100 ms of the end of Conventional Reset at the Root Complex.

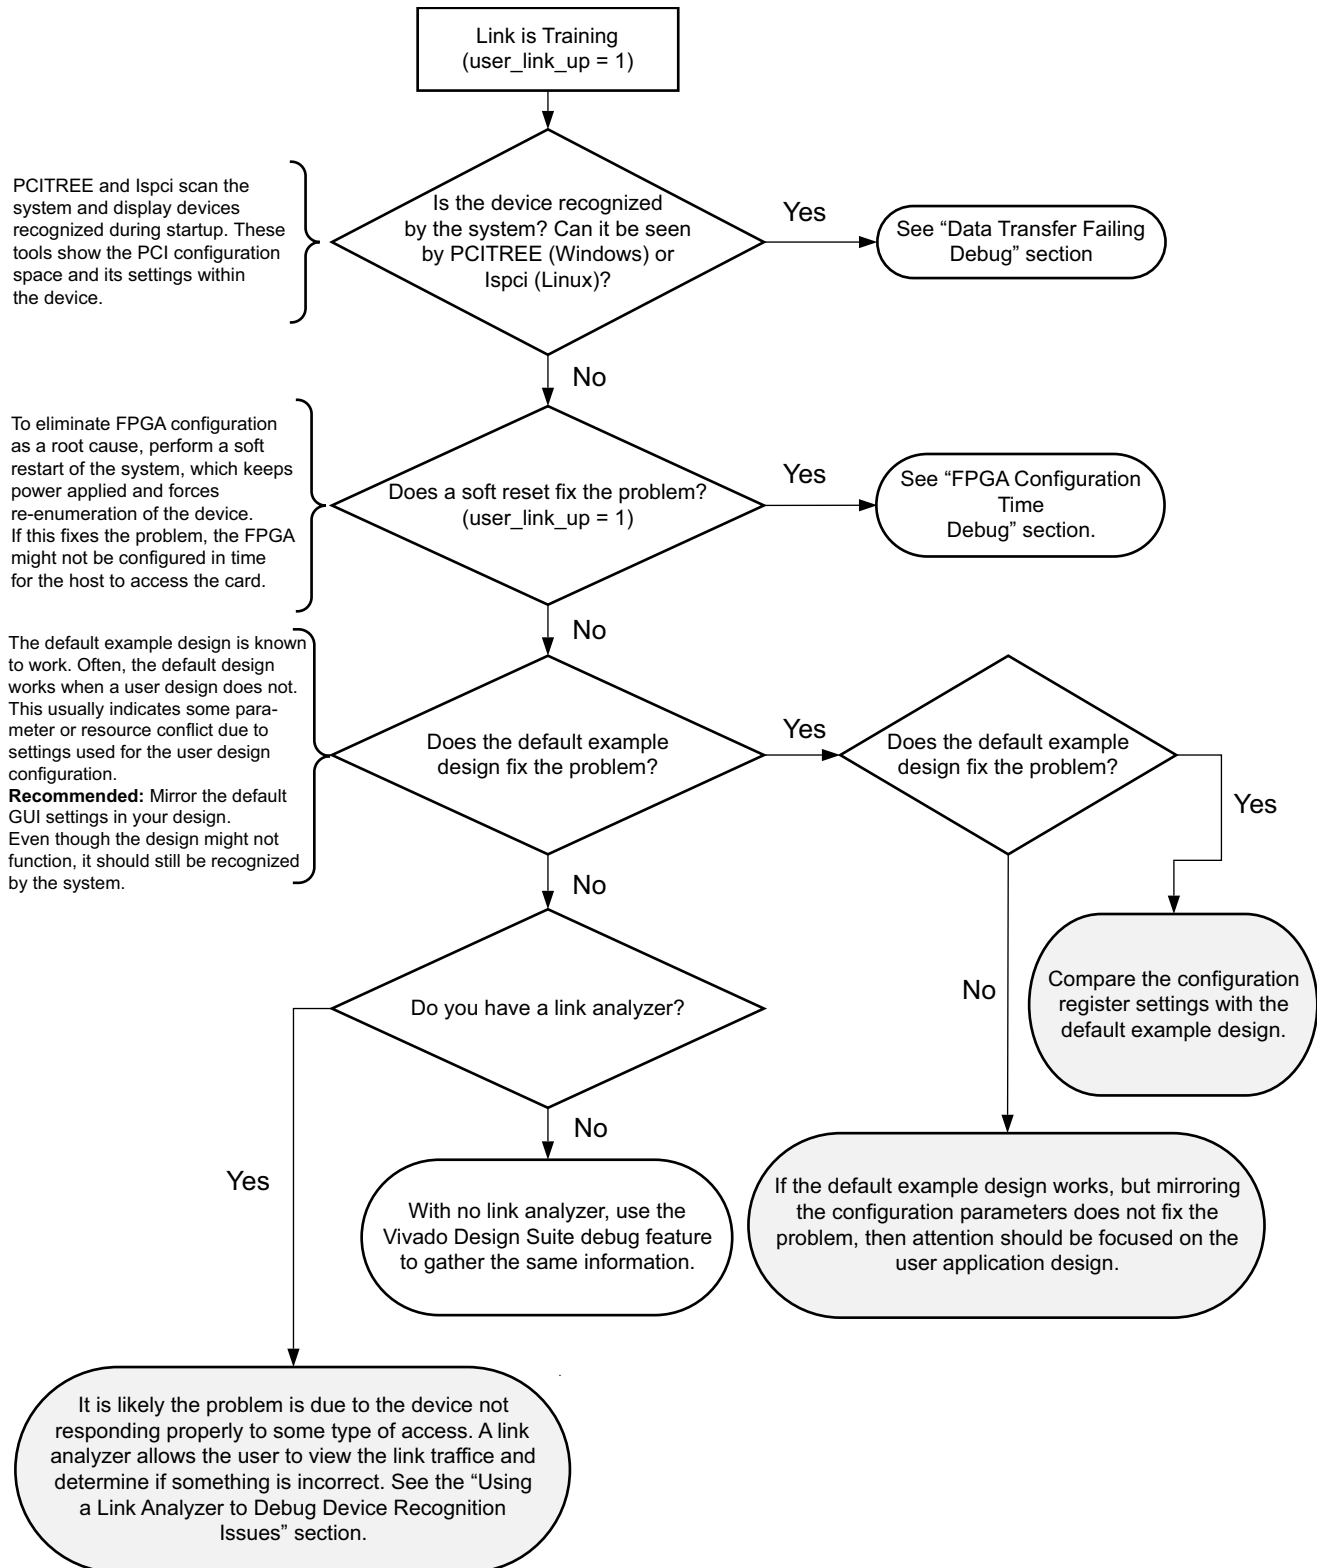
These statements basically mean the FPGA must be configured within a certain finite time, and not meeting these requirements could cause problems with link training and device recognition.

Configuration can be accomplished using an onboard PROM or dynamically using JTAG. When using JTAG to configure the device, configuration typically occurs after the Chipset has enumerated each peripheral. After configuring the FPGA, a soft reset is required to restart enumeration and configuration of the device. A soft reset on a Windows-based PC is performed by going to **Start > Shut Down** and then selecting **Restart**.

To eliminate FPGA configuration as a root cause, you should perform a soft restart of the system. Performing a soft reset on the system keeps power applied and forces re-enumeration of the device. If the device links up and is recognized after a soft reset is performed, the FPGA configuration is most likely the issue. Most typical systems use ATX power supplies which provide some margin on this 100 ms window as the power supply is normally valid before the 100 ms window starts.

Link is Training Debug

Figure B-3 shows the flowchart for link trained debug.



X13110-020821

Figure B-3: Link Trained Debug Flow Diagram

FPGA Configuration Time Debug

Device initialization and configuration issues can be caused by not having the FPGA configured fast enough to enter link training and be recognized by the system. Section 6.6 of *PCI Express Base Specification, rev. 2.1* [Ref 8] states two rules that might be impacted by FPGA Configuration Time:

- A component must enter the LTSSM Detect state within 20 ms of the end of the Fundamental reset.
- A system must guarantee that all components intended to be software visible at boot time are ready to receive Configuration Requests within 100 ms of the end of Conventional Reset at the Root Complex.

These statements basically mean the FPGA must be configured within a certain finite time, and not meeting these requirements could cause problems with link training and device recognition.

Configuration can be accomplished using an onboard PROM or dynamically using JTAG. When using JTAG to configure the device, configuration typically occurs after the Chipset has enumerated each peripheral. After configuring the FPGA, a soft reset is required to restart enumeration and configuration of the device. A soft reset on a Windows based PC is performed by going to **Start > Shut Down** and then selecting **Restart**.

To eliminate FPGA configuration as a root cause, you should perform a soft restart of the system. Performing a soft reset on the system keeps power applied and forces re-enumeration of the device. If the device links up and is recognized after a soft reset is performed, then FPGA configuration is most likely the issue. Most typical systems use ATX power supplies which provides some margin on this 100 ms window as the power supply is normally valid before the 100 ms window starts.

Clock Debug

One reason to not deassert the `user_reset_out` signal is that the FPGA PLL (MMCM) and Transceiver PLL have not locked to the incoming clock. To verify lock, monitor the transceiver RXPLLLKDET output and the MMCM LOCK output. If the PLLs do not lock as expected, it is necessary to ensure the incoming reference clock meets the requirements in *7 Series FPGAs GTX/GTH Transceivers User Guide* [Ref 3]. The REFCLK signal should be routed to the dedicated reference clock input pins on the serial transceiver, and the design should instantiate the IBUFDS_GTE2 primitive in the design. See the *7 Series FPGAs GTX/GTH Transceivers User Guide* for more information on PCB layout requirements, including reference clock requirements.

Reference clock jitter can potentially close both the TX and RX eyes, depending on the frequency content of the phase jitter. Therefore, maintain as clean a reference clock as possible. Reduce crosstalk on REFCLK by isolating the clock signal from nearby high-speed traces. Maintain a separation of at least 25 mils from the nearest aggressor signals.

The PCI Special Interest Group website provides other tools for ensuring the reference clocks are compliant to the requirements of the *PCI Express Specification*: www.pcisig.com/specifications/pciexpress/compliance/compliance_library.

Debugging PCI Configuration Space Parameters

Often, a user application fails to be recognized by the system, but the AMD PIO Example design works. In these cases, the user application is often using a PCI configuration space setting that is interfering with the system's ability to recognize and allocate resources to the card.

The AMD solutions for PCI Express handle all configuration transactions internally and generate the correct responses to incoming configuration requests. Chipsets have limits to the amount of system resources they can allocate and the core must be configured to adhere to these limitations.

The resources requested by the Endpoint are identified by the BAR settings within the Endpoint configuration space. You should verify that the resources requested in each BAR can be allocated by the chipset. Generate a core that implements a small amount of memory (approximately 2 KB) to identify if this is the root cause.

The Class Code setting selected in the Vivado IDE can also affect configuration. The Class Code informs the Chipset as to what type of device the Endpoint is. Chipsets might expect a certain type of device to be plugged into the PCI Express slot and configuration might fail if it reads an unexpected Class Code. The BIOS could be configurable to work around this issue.

Using a link analyzer, it is possible to monitor the link traffic and possibly determine when during the enumeration and configuration process problems occur.

Using a Link Analyzer to Debug Device Recognition Issues

In cases where the link is up (`user_link_up = 1`), but the device is not recognized by the system, a link analyzer can help solve the issue. It is likely the FPGA is not responding properly to some type of access. The link view can be used to analyze the traffic and see if anything looks out of place.

To focus on the issue, it might be necessary to try different triggers. Here are some trigger examples:

- Trigger on the first `INIT_FC1` and/or `UPDATE_FC` in either direction. This allows the analyzer to begin capture after link up.
- The first TLP normally transmitted to an Endpoint is the Set Slot Power Limit Message. This usually occurs before Configuration traffic begins. This might be a good trigger point.

- Trigger on Configuration TLPs.
- Trigger on Memory Read or Memory Write TLPs.

Data Transfer Failing Debug

[Figure B-4](#) shows the flowchart for data transfer debug.

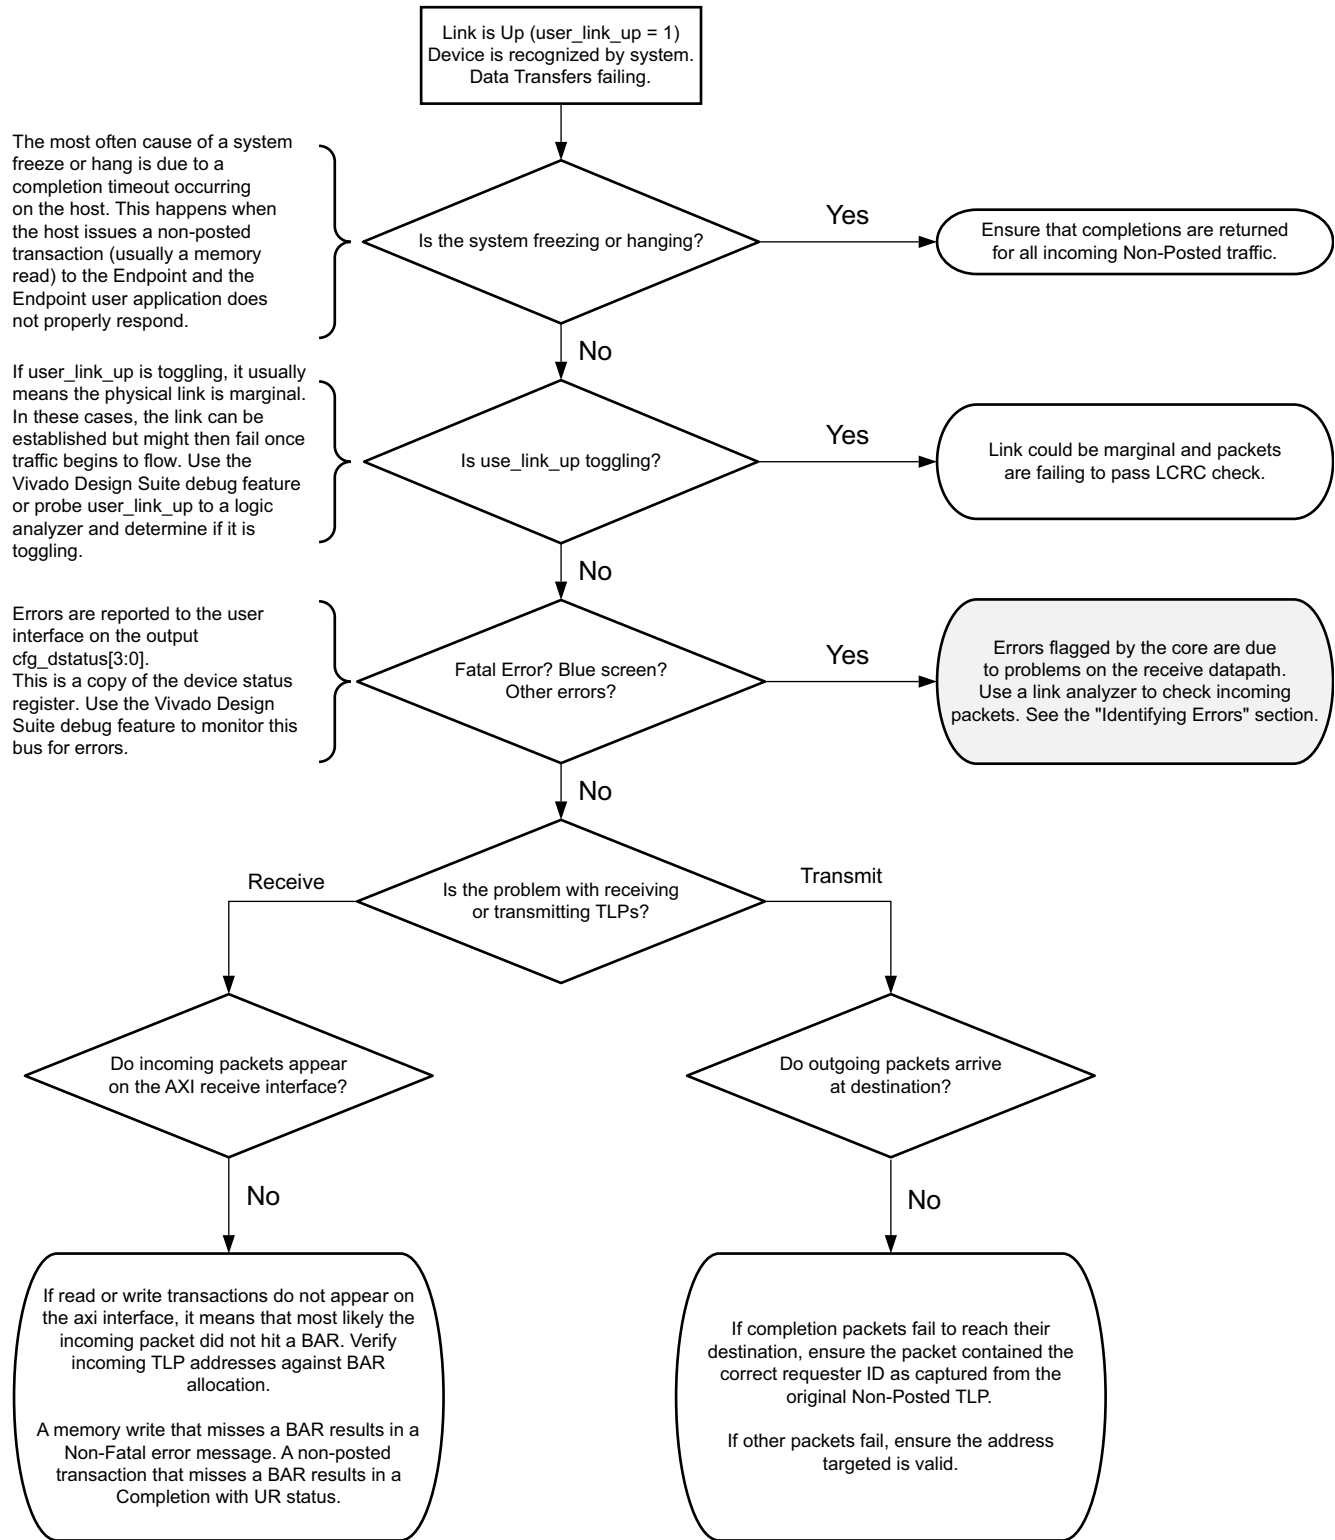


Figure B-4: Data Transfer Debug Flow Diagram

Identifying Errors

See [Abnormal Conditions in Chapter 3](#) for information about how the Slave side and Master side of the AXI Memory Mapped to PCI Express handle abnormal conditions.

Next Steps

If the debug suggestions listed previously do not resolve the issue, open a support case or visit the AMD User Community forums to have the appropriate AMD expert assist with the issue.

To can create a technical support case, go to the [AMD Service Portal](#), and include the following items when opening a case:

- Detailed description of the issue and results of the steps listed above.
- Vivado Design Suite debug feature captures taken in the steps above.

To discuss possible solutions, use the [AMD User Community Forums](#) page.

Transceiver Debug

[Table B-1](#) describes the ports used to debug transceiver related issues.



IMPORTANT: *Debugging transceiver issues is recommended for advanced users only.*

Table B-1: Ports Used for Transceiver Debug

Port	Direction	Width	Description
pipe_txprbssel	I	3	PRBS input.
pipe_rxprbssel	I	3	PRBS input.
pipe_rxprbsforceerr	I	1	PRBS input.
pipe_rxprbscntreset	I	1	PRBS input.
pipe_loopback	I	1	PIPE loopback.
pipe_rxprbserr	O	1	PRBS output.
pipe_rst_fsm	O		Should be examined if <code>pipe_rst_idle</code> is stuck at 0.
pipe_qrst_fsm	O		Should be examined if <code>pipe_rst_idle</code> is stuck at 0.
pipe_sync_fsm_tx	O		Should be examined if <code>pipe_rst_fsm</code> stuck at 11'b10000000000, or <code>pipe_rate_fsm</code> stuck at 24'b0001000000000000000000000000.
pipe_sync_fsm_rx	O		Deprecated.
pipe_drp_fsm	O		Should be examined if <code>pipe_rate_fsm</code> is stuck at 100000000.

Table B-1: Ports Used for Transceiver Debug (Cont'd)

Port	Direction	Width	Description
pipe_rst_idle	O		Wrapper is in IDLE state if pipe_rst_idle is High.
pipe_qrst_idle	O		Wrapper is in IDLE state if pipe_qrst_idle is High.
pipe_rate_idle	O		Wrapper is in IDLE state if pipe_rate_idle is High.
PIPE_DEBUG_0/gt_txresetdone	O		Generic debug ports to assist debug. These are generic debug ports to bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_1/gt_rxresetdone	O		Generic debug ports to assist debug. These are generic debug ports to bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_2/gt_phystatus	O		Generic debug ports to assist debug. These are generic debug ports to bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_3/gt_rxvalid	O		Generic debug ports to assist debug. These are generic debug ports to bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_4/gt_txphaligndone	O		Generic debug ports to assist debug. These generic debug ports bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_5/gt_rxphaligndone	O		Generic debug ports to assist debug. These generic debug ports bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_6/gt_rxcommadet	O		Generic debug ports to assist debug. These generic debug ports bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.

Table B-1: Ports Used for Transceiver Debug (Cont'd)

Port	Direction	Width	Description
PIPE_DEBUG_7/gt_rdy	O		Generic debug ports to assist debug. These generic debug ports bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_8/user_rx_converge	O		Generic debug ports to assist debug. These generic debug ports bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
PIPE_DEBUG_9/PIPE_TXELECIDLE	O		Generic debug ports to assist debug. These generic debug ports bring out internal PIPE Wrapper signals, such as raw GT signals. DEBUG_0 to DEBUG_9 are intended for per lane signals. The bus width of these generic debug ports depends on the number of lanes configured in the wrapper.
pipe_txinhibit	I	1	Connects to TXINHIBIT on transceiver channel primitives.

Interface Debug

AXI4-Lite Interfaces

Read from a register that does not have all 0s as a default to verify that the interface is functional. Output `s_axi_arready` asserts when the read address is valid and output `s_axi_rvalid` asserts when the read data/response is valid. If the interface is unresponsive ensure that the following conditions are met.

- The `axi_ctl_aclk` and `axi_ctl_aclk_out` clock inputs are connected and toggling.
- The interface is not being held in reset, and `axi_aresetn` is an active-Low reset.
- Ensure that the main core clocks are toggling and that the enables are also asserted.
- Has a simulation been run? Verify in simulation and/or a Vivado Design Suite debug feature capture that the waveform is correct for accessing the AXI4-Lite interface.

Additional Resources and Legal Notices

Finding Additional Documentation

Documentation Portal

The AMD Adaptive Computing Documentation Portal is an online tool that provides robust search and navigation for documentation using your web browser. To access the Documentation Portal, go to <https://docs.xilinx.com>.

Documentation Navigator

Documentation Navigator (DocNav) is an installed tool that provides access to AMD Adaptive Computing documents, videos, and support resources, which you can filter and search to find information. To open DocNav:

- From the AMD Vivado™ IDE, select **Help > Documentation and Tutorials**.
- On Windows, click the **Start** button and select **Design Tools > DocNav**.
- At the Linux command prompt, enter `docnav`.

Note: For more information on DocNav, refer to the *Documentation Navigator User Guide* ([UG968](#)).

Design Hubs

AMD Design Hubs provide links to documentation organized by design tasks and other topics, which you can use to learn key concepts and address frequently asked questions. To access the Design Hubs:

- In DocNav, click the **Design Hubs View** tab.
- Go to the [Design Hubs](#) web page.

Support Resources

Note: For support resources such as Answers, Documentation, Downloads, and Forums, see [Support](#).

References

This section provides links to supplemental material useful to this document:

1. *Vivado Design Suite AXI Reference Guide* ([UG1037](#))
2. *7 Series FPGAs Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG054](#))
3. *7 Series FPGAs GTX/GTH Transceivers User Guide* ([UG476](#))
4. *Virtex 7 FPGAs Gen3 Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG023](#))
5. *UltraScale Architecture Integrated Block for PCI Express LogiCORE IP Product Guide* ([PG156](#))
6. *AXI Bridge for PCI Express Gen3 Product Guide* ([PG194](#))
7. [AMBA AXI Protocol Specification](#)
8. [PCI-SIG® Specifications](#)
9. *AXI to AXI Connector Data Sheet* ([DS803](#))
10. *ISE to Vivado Design Suite Migration Methodology Guide* ([UG911](#))
11. *Vivado Design Suite User Guide: Getting Started* ([UG910](#))
12. *Vivado Design Suite User Guide: Designing with IP* ([UG896](#))
13. *Vivado Design Suite User Guide: Logic Simulation* ([UG900](#))
14. *Vivado Design Suite User Guide: Programming and Debugging* ([UG908](#))
15. *Vivado Design Suite User Guide: Designing IP Subsystems using IP Integrator* ([UG994](#))
16. *PIPE Mode Simulation Using Integrated Endpoint PCI Express Block in Gen2 x8 Configurations Application Note* ([XAPP1184](#))

Revision History

The following table shows the revision history for this document.

Date	Version	Revision
11/24/2023	2.9	Updated Feature Summary .
02/22/2021	2.9	<ul style="list-style-type: none"> Added user_link_up to Port Descriptions. Updated link speed for x8 late width in Table 4-2. Editorial updates in IP Facts. Updates for core v2.9 in Upgrading in the Vivado Design Suite. Editorial updates in Hardware Debug.
04/04/2018	2.8	<ul style="list-style-type: none"> Included specific list of Artix 7 devices that are not supported (IP Facts). Added more detail to interrupt_out port description (in the Top-Level Interface Signals table in the Product Specification chapter). Changed C_PCIE2AXIBAR_n to C_PCIEBAR2AXIBAR_n (in the BAR and Address Translation section in the Designing with the Core chapter).
10/04/2017	2.8	<ul style="list-style-type: none"> Updated axi_aresetn description in Top-Level Interface Signals table. Updated Bit[3] description in Interrupt Decode Register table. Added Important note for Device ID in Design Flow Steps chapter.
04/05/2017	2.8	<ul style="list-style-type: none"> Updated to latest 2017.1 release. Added Additional Debug Options section in Design Flow Steps chapter. Added en_jtag_debug in Parameter Changes table in Upgrading appendix.
11/30/2016	2.8	Updated to latest 2016.4 release.
10/05/2016	2.8	<ul style="list-style-type: none"> Updated the Root Port MSI Base Register 2 (Offset 0x150) section (Parameter Dependencies) in Chapter 2.
06/08/2016	2.8	<ul style="list-style-type: none"> Removed MSI-X support in document. Updated Fig. 4-4: PCIe Base Address Register.
04/06/2016	2.8	<ul style="list-style-type: none"> Small editorial update to the Root Port Error FIFO Read Register (Offset 0x154) section. Added the Intercept and Decode Incoming MSI Packet customization option (PCIe Miscellaneous tab). Updated the description for the following customization options: <ul style="list-style-type: none"> Base Address Register Overview (AXI BARs tab) 64-bit Enable (AXI BARs tab) AXI to PCIe Translation (AXI BARs tab) Changed text because I/O BARs are no longer supported. Added parameter and port changes to the Migrating and Upgrading appendix.

Date	Version	Revision
09/30/2015	2.7	<ul style="list-style-type: none"> Added MSI-X and multiple vector address (only single MSI is supported) to Unsupported Features list. Added the C_MSI_DECODE_ENABLE parameter. Clarified that the component name (in Vivado IDE) cannot be same as reserved module names from core. Updated PIPE signal mapping information with new tables: Common In/Out Commands and Endpoint PIPE Signals Mappings, and Input/Output Bus with Endpoint PIPE Signals Mapping.
06/24/2015	2.6	<ul style="list-style-type: none"> Moved performance and resource utilization data to www.xilinx.com. Updated footnote details for top-level interface signal m_axi_arsixe[2:0]. Corrected the documented clocking interface signal names. Updated the documented sequence to clear correctable, non-fatal, and fatal bits of the Interrupt Decode Register. Added the Root Port BAR section. Added the BAR Addressing section. Added the Enable External PIPE Interface option to Vivado IDE options, and added more information about PIPE Mode simulation support. Updated Vivado Lab Edition to Vivado Design Suite Debug Feature. Added new parameter migration information.
04/01/2015	2.6	<ul style="list-style-type: none"> Changed document title to match core name in Vivado IP catalog. Added further clarifying footnotes to Line Rate for PCIe Support for Gen1/Gen2 table. Added the pipe_txinhibit signal. Added Test Bench design details. Updated Vivado lab tools to Vivado Lab Edition.
11/19/2014	2.5	<ul style="list-style-type: none"> Correction made to 64-bit BAR size range upper limit. Updated the simulation procedures for Cadence Incisive Enterprise Simulator (IES), and Verilog Compiler Simulator (VCS). Added Important Note regarding the recommended version of Mentor Graphics Questa SIM to use to avoid simulation failure. Minor change made to debug diagrams.
10/01/2014	2.5	Updated for core v2.5.
06/02/2014	2.4	<ul style="list-style-type: none"> Updated for core v2.4. Added new device support. Removed the axi_aclk and axi_ctl_aclk input ports.
04/02/2014	2.3	<ul style="list-style-type: none"> Updated simulation information. Minor changes and updates.
12/18/2013	2.3	<ul style="list-style-type: none"> Updated for core v2.3. Updated Example Design chapter. Updated parameter changes and port changes in Migrating and Updating chapter.

Date	Version	Revision
10/02/2013	2.2	<ul style="list-style-type: none"> Updated for core v2.2. Updated resource utilization numbers. Added information about the Shared Logic feature. Added example design information. Added port and parameter upgrade information. Added transceiver debug information.
03/20/2013	2.0	Updated for core v2.0, and for Vivado Design Suite-only support.
12/18/2012	1.2	<ul style="list-style-type: none"> Updated core v1.06a, ISE Design Suite 14.4, and Vivado Design Suite 2012.4. Updated Chapter 4, Customizing and Generating the Core. Added Appendix B, Debugging.
10/16/2012	1.1	<ul style="list-style-type: none"> Updated core v1.05a, ISE Design Suite 14.3, and Vivado Design Suite 2012.3. Added Chapter 4, Customizing and Generating the Core. Major updates to PCIe Clock Integration. Added Unsupported Request to Upstream Traffic and Clock Frequencies.
07/25/2012	1.0	Initial Xilinx release. This release is for core version 1.04.a with ISE Design Suite 14.2 and Vivado Design Suite 2012.2. This document replaces DS820, <i>LogiCORE IP AXI Bridge for PCI Express Data Sheet</i> .

Please Read: Important Legal Notices

The information presented in this document is for informational purposes only and may contain technical inaccuracies, omissions, and typographical errors. The information contained herein is subject to change and may be rendered inaccurate for many reasons, including but not limited to product and roadmap changes, component and motherboard version changes, new model and/or product releases, product differences between differing manufacturers, software changes, BIOS flashes, firmware upgrades, or the like. Any computer system has risks of security vulnerabilities that cannot be completely prevented or mitigated. AMD assumes no obligation to update or otherwise correct or revise this information. However, AMD reserves the right to revise this information and to make changes from time to time to the content hereof without obligation of AMD to notify any person of such revisions or changes. THIS INFORMATION IS PROVIDED "AS IS." AMD MAKES NO REPRESENTATIONS OR WARRANTIES WITH RESPECT TO THE CONTENTS HEREOF AND ASSUMES NO RESPONSIBILITY FOR ANY INACCURACIES, ERRORS, OR OMISSIONS THAT MAY APPEAR IN THIS INFORMATION. AMD SPECIFICALLY DISCLAIMS ANY IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, OR FITNESS FOR ANY PARTICULAR PURPOSE. IN NO EVENT WILL AMD BE LIABLE TO ANY PERSON FOR ANY RELIANCE, DIRECT, INDIRECT, SPECIAL, OR OTHER CONSEQUENTIAL DAMAGES ARISING FROM THE USE OF ANY INFORMATION CONTAINED HEREIN, EVEN IF AMD IS EXPRESSLY ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

AUTOMOTIVE APPLICATIONS DISCLAIMER

AUTOMOTIVE PRODUCTS (IDENTIFIED AS "XA" IN THE PART NUMBER) ARE NOT WARRANTED FOR USE IN THE DEPLOYMENT OF AIRBAGS OR FOR USE IN APPLICATIONS THAT AFFECT CONTROL OF A VEHICLE ("SAFETY APPLICATION") UNLESS THERE IS A SAFETY CONCEPT OR REDUNDANCY FEATURE CONSISTENT WITH THE ISO 26262 AUTOMOTIVE SAFETY STANDARD ("SAFETY DESIGN"). CUSTOMER SHALL, PRIOR TO USING OR DISTRIBUTING ANY SYSTEMS THAT INCORPORATE PRODUCTS, THOROUGHLY TEST SUCH SYSTEMS FOR SAFETY PURPOSES. USE OF PRODUCTS IN A SAFETY APPLICATION WITHOUT A SAFETY DESIGN IS FULLY AT THE RISK OF CUSTOMER, SUBJECT ONLY TO APPLICABLE LAWS AND REGULATIONS GOVERNING LIMITATIONS ON PRODUCT LIABILITY.

© Copyright 2012–2023 Advanced Micro Devices, Inc. AMD, the AMD Arrow logo, Artix, Kintex, Spartan, Virtex, Vitis, Vivado, Zynq, and combinations thereof are trademarks of Advanced Micro Devices, Inc. AMBA, AMBA Designer, Arm, ARM1176JZ-S, CoreSight, Cortex, PrimeCell, Mali, and MPCore are trademarks of Arm Limited in the EU and other countries. PCI, PCIe, and PCI Express are trademarks of PCI-SIG and used under license. Other product names used in this publication are for identification purposes only and may be trademarks of their respective companies.